

Government Audits of Municipal Corruption and Belief Updating: Experimental Evidence

Raymond Duch

Felipe Torres-Raposo

Nuffield College, University of Oxford

University College London

September 23, 2026

Abstract

Government audits are widely used to expose public malfeasance, but it is unclear whether such information changes citizens' beliefs or political behavior. We conducted a field survey experiment in Chile in which over 5,000 citizens were randomly assigned to receive short WhatsApp videos summarizing their municipality's most recent audit results. The audit information caused significant and persistent updating of beliefs about local government corruption, with similar effects for copartisans and non-copartisans of the incumbent. Updating was heuristic rather than calibrated: belief changes did not vary with the magnitude of reported malfeasance, though citizens reacted more strongly to deteriorating than to improving audit results. Downstream, the treatment reduced stated intention to vote for the incumbent mayor by 4.5 percentage points but had no detectable effect on donations to municipal public goods. Audit disclosures thus shift both beliefs and electoral evaluations, while leaving broader civic behavior unchanged.

Keywords: Corruption, Belief Updating, Audits

Audit reports are widely endorsed (Berliner & Wehner, 2022; OECD, 2016) and evidence suggests they mitigate corruption (Denly, 2022; Lagunes, 2021). They can also have a causal effect on the corruption beliefs of average citizens. In partnership with the NGO Chile Transparente and the Comptroller General Office (CGO), we assigned video treatments, via WhatsApp, summarizing municipality audit results released by the CGO just prior to the experiment. Treated subjects observed a video reporting 1) the level of malfeasance in pesos; 2) the cost of malfeasance in terms of foregone flu vaccines; and 3) whether levels have improved or deteriorated since the previous audit. The placebo group saw a video unrelated to corruption or malfeasance. Corruption beliefs changed in response to this audit information and this change persisted one week and one month post-treatment.

The design of our field survey experiment builds on the meta study conducted by Dunning, Grossman, Humphreys, Hyde, et al. (2019). They find no evidence that voter information campaigns shape voter behavior. This echoes other mixed evidence on whether, and how, citizens respond to information about corruption. Field experiments that publicize audit findings sometimes document electoral punishment of incumbents (e.g., Ferraz and Finan 2008), but other studies find demobilization, backlash, or null behavioral effects (Cheeseman & Peiffer, 2021; A. Chong et al., 2015).

To explain downstream outcomes such as incumbent vote we argue for well-identified evidence linking information content to directly measured belief updating. Our experiment demonstrates that credible audit disclosures causally and persistently increase citizens' perceptions of local corruption. Moreover, there is evidence of information downstream effects on voting intentions although no evidence that this is mediated by belief updating.

We also argue for a better understanding of the mechanisms underlying belief updating. Messaging content is complex and potentially highly dimensional; the public is diverse and causal effects are likely to be heterogeneous. Our results clearly attest to this. First, while citizens respond strongly to the presence of malfeasance information, their updating is largely insensitive to the magnitude of irregularities reported. Beliefs shift in the direction implied by the signal, but

they do not scale proportionally with the quantitative content of the audits, nor do they consistently converge toward ground truth. Second, temporal benchmarking matters asymmetrically: citizens react more strongly to worsening audit results than to improving ones, suggesting that trends serve as a heuristics in corruption evaluation. Third, while copartisans of incumbent mayors hold more favorable prior beliefs about corruption, with respect to our main outcome measures, they update in parallel with non-copartisans when presented with audit evidence. Credible institutional information can overcome partisan resistance at the level of belief formation, but citizens process.

A strength of our study is that we learn that information content can have a significant impact on belief updating by average citizens. But our attempt to engage subjects with the nuances of municipal audit reports for the most part failed. Hence a limitation of the study is that we can say little about how to design the content and delivery of corruption information so as to optimize the public's engagement.

The patterns we observe (strong directional updating but insensitivity to the reported magnitude, and asymmetric weighting of trend cues) echo recent findings in economic experiments on belief formation. Factual information shifts beliefs about macroeconomic conditions and policy, but uptake of the numeric content is typically partial (Coibion et al., 2023; D'Acunto & Weber, 2024; Roth & Wohlfart, 2020); a parallel literature on misinformation correction documents similarly incomplete revision (Costello et al., 2024). Our contribution is to characterize this incomplete updating in the corruption domain: belief change is uncorrelated with the magnitude of audited malfeasance, and, consistent with work on negativity asymmetry, worsening trends move beliefs more than improvements (Barron, 2021).

Our results also speak to a long-standing debate over partisan motivated reasoning. Classic accounts predict that partisans resist or discount unwelcome information (Kunda, 1990; Taber & Lodge, 2006), yet copartisans and non-copartisans of the incumbent updated in the same direction and to comparable degrees. This parallel response is consistent with recent evidence that political persuasion is largely symmetric across partisan groups (Coppock, 2023) and with rational-updating models in political economy (DellaVigna & Gentzkow, 2010): credible, non-partisan audit infor-

mation can shift beliefs even in a polarized setting.

Our paper proceeds as follows. We first present the theory and then the experimental design. The third section presents the results of the experiment. Finally, we discuss the implications of our findings and suggest directions for future studies of the public’s beliefs about corruption.

Corruption Belief Updating

Our pre-registered conjecture is that the irregularities reported in the Contraloria audits cause updating of corruption beliefs. We review our strategy for measuring audit irregularities and corruption beliefs. Finally, we discuss models of belief updating that inform our experimental design and analysis of the results.

Audit Reports. A small but influential field-experimental literature asks whether audit disclosures shape voter behavior, with mixed results. Ferraz and Finan (2008) report that audited corruption violations reduced Brazilian incumbents’ reelection likelihood by roughly 7 percentage points. Bobonis et al. (2016) find only a small and statistically insignificant 2 percentage point effect of Puerto Rican pre-election audit disclosure on reelection rates. Analogous designs in Mexico have produced weaker or null effects: Arias et al. (2022) report belief updating in response to detailed information on federal transfers and unauthorized spending, while A. Chong et al. (2015), using the same audit data, find no significant change in corruption perceptions or incumbent support and a reduction in turnout.¹ This literature has limitations. First, information is typically delivered through pamphlets or door-to-door visits, leaving take-up and exposure uncertain. Second, downstream electoral or behavioral effects are commonly taken as evidence of belief updating without direct measurement of beliefs themselves.

¹Jablonski et al. (2021) compiled audit-based service-quality indices for Ugandan villages and informed voters of relative performance within their district; the treatments did not shift beliefs about incumbent integrity or effort, nor voting outcomes.

Corruption and Audits in Chile. The salience of corruption and credibility of the auditing institution in Chile make it well-suited for testing audit information effects. Chile ranks 27th of 180 on Transparency International’s Corruption Perceptions Index (Transparency International, 2023), yet in the years preceding our study a series of high-profile scandals kept corruption near the top of the public agenda. Chileans ranked corruption as the second most important national issue in 2017, fourth in 2019 and tenth in 2021 (crowded out by COVID-19 pandemic issues) (LAPOP, 2021) (see appendix Figure A.3).

The audit reports are produced by the Comptroller General Office (CGO) who have extensive access to records on national and municipal government activity. The CGO is ranked in the top tercile globally by the *Supreme Audit Institutions Independence Index* (World Bank, 2021). It is consistently ranked among the most trusted public institutions in annual surveys by Chile’s Council for Transparency (Consejo Nacional para la Transparencia, 2019, 2020).²

Over the last decade the CGO has audited roughly 180 of Chile’s 345 municipalities each year, with audited municipalities selected on a scoring system that weights budget size, transfers to the private sector, and prior audit findings; the weights change from one audit plan to the next. Each audit typically focuses on a single activity (procurement, hiring, or municipal finances) and takes around six months. The 2020 cycle ran from January to December and covered the 116 municipalities used in our study; results were released several months before the May 2021 municipal election.

Corruption Beliefs. Our outcome variable is corruption beliefs. Most measures of corruption beliefs employ a qualitative indicator. These measurements can include individual perceptions of levels or changes in the amount of “corruption” (Peiffer, 2020; Reinikka & Svensson, 2006), bribery (Corbacho et al., 2016; Letki et al., 2023), fraud (Zhou & Oostendorp, 2014), nepotism (Gagliarducci & Manacorda, 2020), waste (Bandiera et al., 2009), or vote-buying (Gonzalez-Ocantos et al., 2014; Hicken et al., 2015). We adopt a qualitative measure consistent with the measurement approach used by (Arias et al., 2022; Enríquez et al., 2024).

²Further details on perceptions of trust in the CGO are available in Section A.3 of the appendix.

Audit reports typically provide a quantitative assessment of government malfeasance. Thus, two of our measures capture the quantitative dimension of these reports. First, we elicit respondents' perceptions of local government malfeasance using a distributional question that assigns probabilities to pre-specified bins of malfeasance levels, following the approach of Arias et al. (2022) and Arias et al. (2019). Second, we include a belief measure that mirrors those used in the inflation expectations (Roth & Wohlfart, 2020) and wage expectations literature (Jäger et al., 2024), where respondents are asked to indicate their best guess from a wide set of discrete ordinal options.

Citizens rarely read audit reports. Their beliefs about local corruption arrive instead through a mix of formal media reporting, interpersonal networks, direct experience, and occasional signals from anti-corruption institutions (Pan et al., 2023). Because this channel mix varies across citizens, priors are heterogeneous and only weakly tethered to objective corruption levels: they encode impressions and anecdotes rather than quantities. Second, when a credible quantitative signal does arrive, citizens process it under motivational and cognitive constraints. The three patterns below develop these points: the first concerns where priors come from; the second and third concern why updating deviates from Bayesian benchmarking.

Channel mix and heterogeneous priors. Credible audit reports can in principle fill informational gaps and anchor beliefs to ground truth. In practice, corruption information mostly reaches citizens through informal sources (friends, family, social media) that often inflates corruption beliefs, while mainstream media fill some gaps but under-report malfeasance (Zhu et al., 2013). The resulting priors reflect citizens' information environments more than the underlying incidence of corruption, and they vary accordingly: in our pre-treatment data, priors differ systematically with the channels through which respondents report learning about politics (Section A.5, Table A.4 and Figure A.6 in the appendix). Because such priors are anchored in impressions rather than magnitudes, we expect citizens to process an arriving audit signal heuristically: registering its presence and direction more readily than its quantitative content, the pattern we develop below.

Partisanship and motivated reasoning. Even when accurate information is available, partisans

may selectively recruit and evaluate evidence to align with attachments to the incumbent, underweighting negative information when it threatens a preferred identity (Bisgaard, 2019; Kunda, 1990; Taber & Lodge, 2006). Economic models formalize this as belief-based utility: citizens derive non-instrumental utility from holding identity-consistent beliefs even at the cost of accuracy (Bénabou & Tirole, 2002, 2016). This mechanism predicts asymmetric updating between copartisans and non-copartisans of the incumbent, a prediction we test directly.

Cognitive limits and base-rate neglect. Independent of partisan motives, citizens face acquisition and processing limits that can prevent beliefs from converging on numerical ground truth. Salient anecdotes crowd out base rates, and bounded rationality constrains the encoding of quantitative content (Bastianello & Imas, 2026). The implication for our setting is that updating may be reliably directional yet remain insensitive to the magnitude of reported irregularities. Consistent with this, Figures A.7 and A.8 in the appendix show that municipal priors are uncorrelated with the frequency of past audits or the severity of past irregularities, suggesting that repeated exposure to anti-corruption information does not, on its own, calibrate beliefs to objective levels.

From Mechanisms to Hypotheses. The preceding discussion yields two premises: priors about local corruption are heterogeneous and weakly calibrated to ground truth, and citizens process new quantitative signals under motivational and cognitive constraints. They imply a structured set of seven hypotheses. A first set concerns *heuristic updating*: if subjects allocate limited attention to the quantitative content of the audit signal, we expect a directional shift in corruption beliefs (H1) but little sensitivity to the reported magnitude of malfeasance (H3), and little narrowing of the gap between subjective beliefs and audit ground truth (H4). A second set concerns *priors and reference points*: benchmarking against past audits in the same municipality should anchor the direction and size of updating (H2), while subjects at the extremes of the heterogeneous prior distribution should update conservatively, attenuating the treatment effect (H5). A third set concerns *motivated reasoning and its behavioral consequences*: copartisans and non copartisans may process the same signal differently even under incentivized elicitation (H6), and any belief shift should

plausibly translate into willingness to contribute to municipal public goods (H7a) and support for the incumbent (H7b).

The pre-analysis plan registered before data collection committed us to five of these predictions: the directional treatment effect (H1), the temporal-benchmarking contrast (H2), the information-gap narrowing test (H4), and the two behavioral outcomes (H7a, H7b). The remaining three hypotheses, namely magnitude-correlation prediction (H3), the extreme-priors heterogeneity (H5), and the partisanship contrast (H6), are exploratory extensions developed during analysis. Each hypothesis below is tagged accordingly. Table A.3, in the appendix, reports the predicted direction, observed direction, and standardized magnitude for each hypothesis.

Cognitive Shortcut. The randomly assigned information treatment in our experiment is a quantitative narrative about local governmental corruption. Our pre-registered expectation was that subjects would update conditional on how the information content interacted with their priors. A much more basic, non-registered, conjecture is that subjects allocate limited attention to the information content we provided them. With limited cognitive investment in the information content, the signal is that public officials commit malfeasance. Subjects update their beliefs about levels of corruption when they learn about local municipal government malfeasance.

H1: Being informed of municipal government malfeasance increases beliefs about levels of corruption.

We are able to correlate features of the corruption messaging with updating behavior. Respondents learn about magnitudes of malfeasance and whether performance has improved or gotten worse. While respondents are randomly assigned to observe the information treatment, they are not randomly assigned to the municipality for which a high or low malfeasance audit outcome is realized. Nevertheless, the response of subjects to these quantitative details provides suggestive insights into belief updating. There are features of the quantitative information treatment that we expect to correlate with the direction and magnitude of belief updating and that speak to a number of conjectures regarding corruption beliefs.

Benchmarking. A potential mechanism underpinning belief updating about corruption is bench-

marking informative signals about government malfeasance. Benchmarking is a foundational assumption of many models of responsibility attribution (Barro, 1973; Ferejohn, 1986; Fiorina, 1981). There is evidence of benchmarking incumbent government performance; certainly concerning economic performance (Arel-Bundock et al., 2021; Aytaç, 2018; Kayser & Peress, 2012); but also in other policy areas such health (Becher et al., 2023; R. M. Duch et al., 2025; Rodríguez et al., 2025) and education (Charbonneau & Van Ryzin, 2015). Studies though suggest that benchmarking is challenging for average citizens (Arel-Bundock et al., 2021; Healy & Lenz, 2014).

If citizens benchmark, what do they benchmark against? The literature suggests two reference points: Spatial benchmarking, for example, juxtaposes incumbent malfeasance information with comparisons from other similar jurisdictions (Banerjee et al., 2014; Bhandari et al., 2023; Enríquez et al., 2024). Temporal benchmarking, on the other hand, compares incumbent performance with their historical performance or with previous outcomes under a challenger (Arias et al., 2025; Avenburg, 2019; Botero et al., 2015; Breitenstein, 2019). We explore the temporal benchmarking mechanism by providing respondents with a comparison of their municipality’s current audit performance with previous audit results. The expectation is that these benchmarks inform corruption beliefs by providing a reference point in assessing levels of malfeasance but the evidence of their effectiveness is mixed.

H2: Beliefs regarding corruption levels will be higher if reported malfeasance has increased compared to the previous audit and will be lower if reported malfeasance has declined.

Magnitude Encoding. Empirical studies of corruption suggest that reported malfeasance magnitudes correlate with belief updating (Arias et al., 2022; Boas et al., 2019; A. Chong et al., 2015; Enríquez et al., 2024; Ferraz & Finan, 2011; Larreguy et al., 2020). A concern though is that average citizens struggle to correctly assess these quantitative policy outcome metrics. We see this, for example, with respect to crime statistics (Esberg & Mummolo, 2018; Larsen & Olsen, 2020); immigration (Blinder & Schaffner, 2019; Hopkins et al., 2019), government debt (Roth et al., 2022), and trade balances (Mutz, 2021; Rodríguez Chatruc et al., 2021).

An outstanding puzzle is how to frame malfeasance magnitudes to ensure engagement and

comprehension. The Enríquez et al. (2024) information treatments, for example, consisted of the percentage of audited funds subject to irregularities. Similarly, Boas et al. (2019) informed voters of the rate of funds rejected by the local audit agency. Based on pre-testing of alternative framing strategies, the information treatment in our trial reports the levels of malfeasance in Chilean pesos – and these vary significantly across municipalities.

H3: Belief updating will correlate with the reported magnitude of malfeasance levels in the municipality.

Information Gap. Citizens’ priors regarding levels of malfeasance in their community can be inaccurate. Our pre-registered hypothesis was that acquiring accurate information about corruption would reduce the gap between subjective beliefs and objective reality (Cavallo et al., 2017; D’Acunto & Weber, 2024). There is some evidence that, under certain conditions, average citizens receive ground-truth information and subsequently update their beliefs in a direction that narrows this gap (Jäger et al., 2024; Tappin et al., 2020). However, other studies offer less encouraging findings: corrective information can sometimes backfire, exacerbating the discrepancy between beliefs and reality (Carey et al., 2025; Nyhan & Reifler, 2010; Zappalà, 2024). Our interest is in whether credible information regarding malfeasance causes citizens to form more accurate beliefs about corruption levels (Liang, 2025; Sanna & Lagnado, 2025).

Our experiment measures this gap between beliefs and ground truth. Ground truth is derived from audit agency data and we elicit beliefs about this ground truth with incentivized survey questions. We measure these incentivized corruption beliefs, for control and treated subjects, both pre-treatment (measured four months before the intervention) and post-treatment. Subjects have a financial incentive to use this information to narrow this information gap. But even given the financial incentives a host of factors may compromise informed updating.³

H4: Providing respondents with audit agency data about their municipality will reduce the gap between their subjective beliefs and objective corruption metrics.

Anchoring. Belief updating may differ for individuals with extreme priors – either high or

³Both quantitative outcomes were incentivized: respondents could earn a 20% bonus on their survey participation payment for answering these questions correctly.

low (Fan et al., 2026). This might result from anchoring and conservative updating at the extremes (Coutts, 2019; Eil & Rao, 2011; Kovach, 2021) or expectations reference points (Coibion & Gorodnichenko, 2015). Several corruption studies have provided evidence that belief updating is conditional on citizens' priors relative to what is reported in the signal. A case in point is Arias et al. (2022), who report that Mexican voters, with high corruption priors, revised their beliefs more favorably after receiving corruption audits about the incumbent, while voters with low corruption priors updated beliefs negatively. Belief updating in this case was conditional on whether voters had high or low corruption priors.

There may be diminishing returns to information for people with very high or low corruption priors. Individuals with relatively negative priors regarding malfeasance in their local government respond to negative versus positive corruption signals differently than those with relatively positive priors. The effect of providing negative information about the malfeasance of a municipal government will be decreasing in individuals' prior beliefs that the government is malfeasant. Similarly, the effect of providing positive information would be decreasing in the individuals' prior beliefs that the government is not malfeasant. Those with negative (positive) malfeasance priors may already anticipate audit results that report negative (positive) information about irregularities in their municipal government, and hence, there is minimal updating reflected in their posteriors.

H5: Subjects with relatively high corruption priors will be unresponsive to negative reports of malfeasance. Subjects with relatively low corruption priors will be unresponsive to positive reports of malfeasance.

Motivated Reasoning. Partisan motivated reasoning may color how individuals process policy-related information (Kunda, 1990; Mian et al., 2021; Peterson & Iyengar, 2021). Thaler (2024) presents experimental results suggesting that partisanship conditions Bayesian updating with respect to a range of quantitative policy metrics. Recent research also suggests that partisanship affects how individuals respond politically to information about government corruption (de Figueiredo et al., 2022; Elia & Schwindt-Bayer, 2022; Solaz et al., 2019). This implies that partisanship conditions belief updating when individuals receive new quantitative information about malfeasance

by their local government (Anduiza et al., 2013; Cornejo, 2022; Cubel et al., 2024).

An alternative conjecture is that while partisans' priors about levels of corruption in their municipality differ, both copartisans and non-copartisans will update their beliefs similarly in response to factual information about corruption. Recent experimental findings give some credence to this argument: Tappin et al. (2023) demonstrate that partisans update when they receive persuasive messages and that there was no evidence that party elite cues meaningfully diminished partisans' receptivity to the messages; the Kim (2025) experimental results suggest that in the absence of affective triggers, partisans were persuaded by both congenial and uncongenial information treatments; and Guess and Coppock (2020) find no support for the notion that counter-attitudinal evidence or arguments reinforce partisan pre-existing opinions and beliefs, a phenomenon sometimes known as "backlash". In the case of our online field experiment, this "parallel persuasion" argument (Coppock, 2023) suggests, first, that copartisans of the governing municipal party will perceive the levels of municipal government malfeasance to be much lower than the levels perceived by non-copartisans. Second, it implies that our malfeasance treatment will have similar effects on copartisans and non-copartisans.

H6a (Motivated reasoning): The effect of audit information on corruption beliefs is smaller for copartisans of the incumbent than for non-copartisans.

H6b (Parallel persuasion): Copartisans and non-copartisans differ in their priors but update by similar amounts, so the partisan gap in beliefs persists after treatment.

Belief Behavior Link. Ultimately, we are interested in corruption belief updating because these changing beliefs affect political behavior. Again, while these updated beliefs are not randomly assigned their correlation with behaviors can be informative. We will focus on two political behaviors: public goods contributions and incumbent vote. Our pre-registered hypothesis is that levels of government malfeasance are negatively correlated with the public's enthusiasm for contributing to public goods provision (Beekman et al., 2014; Campos-Vazquez & Mejia, 2016; Jahnke & Weisser, 2019; Peyton, 2020). There is also an extensive literature suggesting a link between governmental malfeasance and vote for incumbent candidates (Arias et al., 2022; Banerjee et al., 2014;

Boas et al., 2019; Buntaine et al., 2018; A. Chong et al., 2015; Dunning, Grossman, Humphreys, Hyde, et al., 2019; Enríquez et al., 2024; Ferraz & Finan, 2008). Our expectation is that the belief updating we observe would be correlated with incumbent vote preferences.

H7a: Reported malfeasance is negatively correlated with contributions to public.

H7b: Reported malfeasance is negatively correlated with vote for incumbent governing parties.

Experimental Design

Recruitment In anticipation of the 2020 annual municipal audits, we recruited approximately 49,883 subjects, primarily through social media (Facebook, Instagram, Twitter) and standard panel-recruitment methods such as posters, in-person recruitment, and local media advertisements.⁴ Recruitment aimed to approximate the national census distribution by gender, age, education, and region while ensuring adequate coverage of the audited communes. To limit self-selection, messaging conveyed only basic logistical details (including compensation) and made no reference to the survey’s topic or research questions.

Sample. In January 2020, the CGO selected 116 municipalities for the 2020 audit program; our pool contained 46,723 residents across these municipalities. Between January and March 2021, we conducted a pre-treatment survey with 5,528 pool members resident in these municipalities.⁵ Table A.1 in the appendix benchmarks the final sample against the Chilean census: the sample matches the census age profile but is more female and significantly better educated.

Table A.2 benchmarks our sample against actual voting outcomes and electorate demographics in the 116 municipalities: It is more left-leaning, more likely to report being non-partisan, and reports higher turnout. Reported turnout in the 2016 municipal election was 64.67%, versus 31.50% actual turnout in these municipalities, a 2.05-fold difference. Partisan attachments are skewed:

⁴Full details, including the ads, are in appendix Section A2.

⁵Subjects had a window of more than four weeks to respond to the invitation. Of 6,050 eligible respondents who completed the baseline survey, 5,528 participated in the final study; we excluded 10 municipalities whose audit reports were released late or too close to the election. The effective response rate was about 25 percent.

roughly 73% of respondents identify with left or center-left parties (which received a 31% average vote share), and 38% report no affiliation (independents received about 16%).⁶ The sample also skews female and young among reported 2016 voters: 64.7% of self-reported voters were women (versus a 55% average), and roughly 43% were aged 18–29 (versus about 15% of the voting population). Our estimates are therefore best read as applying to a comparatively politically engaged, left-leaning, and young population, an external-validity caveat to bear on the results that follow.

Field Survey Experiment. Our design occupies an intermediate position along the spectrum from survey experiments (Mutz, 2011) to classic field experiments (A. Chong et al., 2015; Dunning, Grossman, Humphreys, Hyde, et al., 2019). Two features classify it as a field experiment (Gerber & Green, 2012): first, the treatment consists of real information about actual government malfeasance, not a hypothetical vignette; its content is drawn directly from official CGO audit reports and varies with actual audit findings across the 116 municipalities. Second, we controlled random assignment, and the treatment was delivered via WhatsApp as a standalone message, through a medium respondents already use for political communication and in a context that closely resembles a naturally occurring setting. The design deviates from a classic field experiment in that treatment is not delivered in a fully naturalistic setting: subjects were recruited beforehand, and because WhatsApp also served as the survey channel, those who consented may have anticipated contact. In this respect our study resembles a field experiment with survey outcomes (Kalla et al., 2021). It also differs from natural-experiment studies of audit disclosure, which exploit the public release of audit findings and their diffusion through media and campaigns (Ferraz & Finan, 2008): there, exposure to the audit signal is uncontrolled and diffuse; here, it is individually randomized and directly delivered.

The Malfeasance Signal. Treated subjects are informed of their municipality’s audit outcomes through a short 50-second video delivered via WhatsApp. The medium is well suited to the task:

⁶We compare partisan attachment rather than recalled vote choice, since we only asked whether respondents voted in 2016. We did not elicit which party or candidate they voted for, given that the election preceded our survey by several years and we were concerned about recall accuracy.

85% of our respondents reported using WhatsApp multiple times a day and 75% said they receive political information through it, well above Facebook (44%), Instagram (49%), and TikTok (22%).


How best to signal malfeasance, however, is not obvious. Magnitude can be conveyed as an overall level, a share of the budget (Buntaine et al., 2018), or foregone expenditure (Arias et al., 2022; Larreguy et al., 2020), and Bhandari et al. (2023) show that temporal benchmarking, comparing current malfeasance to the district’s own past, moves beliefs more than spatial benchmarking against other districts. To choose among these options empirically rather than by assumption, we fielded a separate pilot before producing the treatment videos: a WhatsApp chatbot experiment ($n = 3,996$) testing 24 messaging strategies. Each respondent was randomly shown six video messages and rated each on whether it was convincing, trustworthy, reliable, and precise (items drawn from prior work on message effectiveness (O’Keefe, 2021; Thomas et al., 2019) alongside comprehension and attentiveness checks gauging whether respondents grasped the quantitative content (the benchmark and the magnitude of the amounts). The pilot showed that framing malfeasance as foregone resources and using a temporal benchmark was most effective (R. Duch & Torres, 2020). We therefore built both frames into the full-scale field survey experiment, which is why the treatment described above conveys the audited amount, its cost in foregone flu vaccines, and whether corruption improved or worsened relative to the prior audit.

Figure 1 presents the three frames of the treatment video – it illustrates both a worsening (top row) and improving (bottom row) audit result.⁷ The videos are customized to reflect the CGO audit results for the participant’s municipal government. First, respondents received information about the total amount of malfeasance found in the 2020 audit report. A second frame compares current reported malfeasance levels with those reported in the municipality’s previous audit. The third frame expresses malfeasance magnitude as the equivalent number of foregone influenza vaccines, since citizens struggle to interpret large raw sums without a reference point or sense of the relevant base rate. Anchoring abstract quantities to familiar objects follows the localized-analogies literature (Spatharioti et al., 2024). We chose the influenza vaccine for several reasons. Chilean

⁷Screenshots of the placebo video can be found in appendix Figure A.1. The videos for each of the 116 municipalities and the data used in these videos are available at: <https://link-omitted-to-preserve-anonymity>.

municipalities, though not the program’s funder, are central to its delivery (disseminating information, coordinating logistics, and mobilizing residents), making it a natural local benchmark. Its annual rollout, reinforced by the salience of health policy after COVID-19, makes it more recognizable than a one-time policy. And because vaccine prices are public and the program runs nationally, the foregone benefit can be quantified precisely and reported for every municipality in our sample.

Figure 1: Screenshots Depicting the Video Treatment.



Note: These are screenshots of the video information treatments that subjects received. The top row has screenshots of respondents in municipalities with rising malfeasance. The first part of the video reports the amount of malfeasance found in their municipality. The second screenshot is a temporal comparison – corruption has increased over time in this municipality. The third screenshot on the top row shows “foregone losses” expressed as the number of influenza vaccines. The three images on the bottom row show the same information, although for a municipality with decreasing levels of irregularities over time (we label this “positive” information).


Treatment Assignment. Figure 2 provides a summary of the experiment design. Treatment assignment was conducted within each of the 116 municipalities. Our initial pilot results indicated that respondents’ prior beliefs are strongly correlated with their posteriors. Subject’s malfeasance priors (both qualitative and quantitative) were measured in the pre-treatment survey. We constructed *High* and *Low* blocks using these subjective malfeasance scales (from 1 to 10). Within each municipality, we blocked on high versus low malfeasance priors; subjects within each block were randomly assigned to either their municipal malfeasance treatment video or the placebo condition video. This design element allowed us to create homogeneous sets of respondents based on their priors and improve the precision of our estimates (Gerber & Green, 2012).

The videos were sent to these 5,528 participants via WhatsApp; 3,696 participants assigned to treatment and 1,832 assigned to the placebo condition.⁸ A post-treatment survey was administered to participants 3-6 days after they were treated with the videos. 3,592 subjects responded to the survey: 2,249 out of 3,696 treated (61%) and 1,343 out of the 1,832 placebo subjects (73%). Finally, a follow-up survey was administered one month after the video treatments, with 1,198 out of 3,696 treated (32%) responding and 1,057 out of the 1,832 placebo subjects (58%) responding.⁹

⁸Immediately after respondents watched the video, they were asked a series of compliance and attention check questions. Out of the 1,832 respondents in the placebo group, 1,055 answered correctly the follow-up Chatbot questionnaire (57%). Of the 3,696 respondents in treatment, 1,961 answered correctly the Chatbot attention check questions (53%).

⁹The initial design anticipated two treatment arms plus a placebo arm. The final implementation consisted of only one treatment arm and the placebo without adjusting the sample sizes of the final treatment and placebo arms.

Figure 2: Flow Chart of Experimental Design.



Outcomes. Table 1 presents the treatment and placebo group mean values and standard errors for our three outcome variables measured at each survey wave.¹⁰ The Qualitative outcome measurement question is adapted from Enríquez et al. (2024): “How much of the municipal government expenditures are corrupt?” (10-point scale from none to all). It’s average value for treated subjects rises from 4.49 in pre-treatment to 5.34 post-treatment and then declines somewhat to 5.15.

Building on recent findings regarding incentives and accuracy of beliefs (Rathje et al., 2023; Zimmermann, 2020), we incentivized the two quantitative outcomes: *Distribution* and *Resources*. The *Distribution* variable is subjects’ reported ranking of their municipality in terms of levels of malfeasance: Low (1); Moderate (2); High (3); and Very High (4). They are presented with the actual distribution of Chilean municipalities in four categories. Participants earn 300 CLP (Chilean Pesos) if they correctly categorize their municipality.¹¹ For treated subjects, in pre-treatment, the average Distribution score is 1.91, just below moderate malfeasance, rising to 2.29 in the initial post-treatment survey and then declining to 2.19 in the follow-up.

For the incentivized *Resource* measure, we ask participants to indicate how much out of every 1000 CLP municipality expenditures is associated with corrupt activities. Participants earn an additional 300 CLP if they get the correct answer. This Resource measure has been converted to ordinal scores ranging from 1 (“Less than 100 CLP”) to 2 (“Between 100 and 200 CLP”) until 10 (“More than 900 out of 1000 CLP”). For treated subjects, in pre-treatment, the average Resource score is 3.85; in post-treatment, the average increases to 4.57, but it declines in the follow-up to 4.22.

The use of incentives does not appear to bias estimates of the treatment effect. The incentive structure is identical in the pre- and post-treatment waves, so any level shift it induces is common to both and differences out; what we estimate is the change in beliefs following treatment, holding the incentive constant. If anything, incentives render our test conservative: they raise attentiveness and cognitive effort on these demanding factual questions (Camerer & Hogarth, 1999; Krosnick,

¹⁰Several recent studies have highlighted the importance of robust measurement strategies to avoid experimental measurement error (Gillen et al., 2019; Haaland et al., 2023).

¹¹This incentive represented about 20% increase in their subject payment, if they answered correctly both incentivized questions.

1991), so magnitude-insensitive updating cannot easily be attributed to inattention. Nor is there strong reason to expect incentive-induced experimenter demand: the available evidence suggests such effects are minimal (de Quidt et al., 2018; Mummolo & Peterson, 2019). A valid concern is that that treatment effects can vary with financial incentives (Baillon et al., 2022; Bullock et al., 2015; Prior et al., 2015). By raising the stakes, incentives may subtly redefine the task relative to an unincentivized setting, which has implications for how far the treatment effects generalize. We also include a non-incentivized measure of respondents' assessment of municipal malfeasance, allowing us to verify that the treatment moves beliefs even absent the incentive. In Table 1 we observe that the average values for all three outcome variables rise in the post-treatment wave. Again, for all three outcome variables, the average values of subjects in the placebo condition are similar across all three waves.

Respondents are asked to donate 500 CLP (about 15 percent of their earnings) to the provision of a local public good. In the initial wave, about one-third of treated respondents donated 500 CLP – this drops to 15 percent in the post-treatment survey for treated subjects and then drops to 10 percent in the one-month follow-up.

Table 1 summarizes covariate values. Our sample is disproportionately female. Most of our sample had completed primary and secondary education and some form of tertiary education. The mean partisanship across all waves is Left-leaning with a score of 4.4 measured on a Left-Right scale ranging from 1 to 10. Approximately one-quarter of the respondents indicated they would vote for the incumbent mayor (*Copartisan*). Average individual monthly income is in the range of \$648-\$678 in pre-treatment and declines somewhat over the subsequent waves.

Balance. Respondents are observed at three waves: baseline, when treatment is assigned and implemented; one week post-treatment; and one month post-treatment. We assess covariate *balance* in appendix Figure A.9, which compares standardized mean differences across the treatment and placebo arms at each wave, using an indicative tolerance of 0.1 (Stuart et al., 2013). With few exceptions, the unadjusted standardized mean differences fall within this threshold. Appendix Table

Table 1: Descriptive Mean Values for Malfeasance Belief Outcome Measures and Covariates

	Survey Wave					
	Pre-Treat		Post-Treat		Follow-up	
	Placebo	Treatment	Placebo	Treatment	Placebo	Treatment
Qualitative Measures						
Subjective malfeasance scale (1-10)	4.47 (2.44)	4.49 (2.44)	4.60 (2.35)	5.34 (2.34)	4.39 (2.30)	5.15 (2.29)
Certainty subjective malfeasance scale (1-10)	5.44 (2.72)	5.39 (2.70)	5.28 (2.72)	5.74 (2.59)	5.46 (2.72)	5.52 (2.59)
Quantitative Measures - Incentivized						
Resources (1-10)	3.83 (2.36)	3.85 (2.39)	3.78 (2.36)	4.57 (2.60)	3.52 (2.22)	4.22 (2.48)
Distribution (1-4)	1.90 (0.78)	1.91 (0.79)	1.93 (0.79)	2.29 (0.87)	1.89 (0.78)	2.19 (0.84)
Behavioral Measures						
500 CLP Municipal Donation	32% (1.09%)	33% (0.77%)	17% (1.27%)	15% (0.78%)	11% (1.42%)	10% (1.34%)
Covariates						
Female	64.79% (1.11%)	64.69% (0.78%)	64.78% (1.3%)	64.61% (1%)	64.81% (1.46%)	66.61% (1.35%)
Education (Yrs)	15.09 (0.06)	15.05 (0.04)	15.10 (0.08)	15.04 (0.06)	15.06 (0.08)	15.03 (0.08)
Age	36.42 (0.31)	35.61 (0.21)	36.41 (0.37)	35.37 (0.28)	36.45 (0.43)	35.17 (0.38)
Partisanship (1 to 10)	4.55 (0.05)	4.43 (0.03)	4.58 (0.06)	4.38 (0.04)	4.58 (0.06)	4.31 (0.06)
Co-partisanship	28% (1.05%)	27% (0.73%)	28% (1.22%)	26% (0.93%)	28% (1.38%)	26% (1.27%)
Income (USD)	678 (30)	648 (18)	656 (28)	621 (24)	604 (29)	556 (21)
Sample	1832	3696	1343	2249	1057	1198

Note: This table reports means for all outcomes and covariates by wave and treatment group. *Behavioral* outcomes show the share donated to respondents' municipalities. *Female* is the proportion of women; *Education* in years; *Partisanship* ranges 0 ("left") to 10 ("right"); *Co-partisanship* is the share intending to vote for the incumbent mayor; *Income* uses self-reported bracket midpoints. Standard errors in parentheses.

A.5 reports regressions of treatment assignment on pre-treatment covariates, confirming balance across arms.

Non-random treatment direction. Conditional on the actual audit reports, the information treatment is relatively positive or negative, and this direction is *not* randomly assigned; it reflects the reported audit results. Appendix Table A.13 compares respondents in municipalities receiving negative versus positive reports. The two samples are broadly similar, though those in negative municipalities had higher political knowledge, were more likely to be government copartisans, and were more likely to receive municipal social benefits; there were no significant differences in outcome-variable priors. Appendix Table A.14 extends this comparison to municipal-level malfeasance and

political outcomes. Negative and positive municipalities are audited at similar rates and exhibit comparable levels of malfeasance, both in the total number of irregularities and in the incidence of *severe* irregularities. Across electoral competition, participation, candidate composition, ideological polarization, and ethnic composition of the electorate, we observe no substantive differences, although positive-outcome municipalities show somewhat higher turnout and more candidates and electoral coalitions.

Attrition. We next examine how differential *attrition* between baseline and the two post-treatment waves affects the composition of the arms. Appendix Table A.7 reports logit regressions of post-treatment non-response on the treatment-arm dummy. The treatment-arm coefficients are positive and precisely estimated, indicating higher attrition in the treatment arm, and attrition correlates positively with income, education, and civic knowledge. Comparing baseline priors for attriters versus those who completed both follow-up waves Appendix Table A.6 compares attriters with those who completed both follow-up waves: Attriters report higher baseline perceptions of malfeasance on the Qualitative outcome but do not differ significantly on the quantitative outcomes; the two groups are broadly similar in gender, education, age, and partisanship, though they differ in income. Despite this differential attrition, overall covariate balance remains intact (Figure A.9), and the study’s differential attrition rate of 13% sits within the upper range observed in comparable field experiments (Ghanem et al., 2023). We assess robustness to attrition in two complementary ways, with full detail and the plausibility of each method’s identifying assumptions discussed in Section A.6. First, we estimate inverse probability weights (IPW) from a first-stage model predicting attrition status; In the appendix, Table A.10 presents the IPW-weighted treatment effects alongside the unweighted estimates from Table 2, and the two are consistent.¹² Second, we estimate Lee bounds (Lee, 2009) and Generalized Lee bounds (Semenova, 2025), incorporating baseline covariates to tighten the bounds. Across both approaches the treatment effects are stable, providing reassurance that attrition does not meaningfully distort the main findings.

¹²McFadden’s pseudo- R^2 values for the propensity-score model are reported at the foot of Table A.10; the model’s classification performance (confusion matrix) is reported in appendix Figure A.10.

Compliance. A post-treatment survey was administered 3–6 days after treatment. Asked whether they recalled viewing the videos, 846 respondents (63% of post-treatment respondents) recalled the placebo video and 1,670 (74%) recalled the treatment videos. In appendix Table A.19 we define Post-Treatment Compliers as the roughly 45% of subjects, across both conditions, who answered the post-treatment survey and reported viewing the videos. Appendix Table A.20 compares baseline measures for compliers and non-compliers and finds little evidence that they differ. In the main text we report intention-to-treat (ITT) estimates.

Results

Our results support the pre-registered hypothesis that corruption beliefs, measured pre-treatment (Corrupt_{ti}) and post-treatment ($\text{Corrupt}_{(t+1)i}$), respond to a short randomly assigned video information intervention.¹³ Variation in the quantitative details that subjects observe allows us to gain some understanding of how different elements of the malfeasance message affects belief updating. While these features of the message are not randomly assigned they provide insights into how average citizens engage with malfeasance information reported in audits.

Information Treatment.

When subjects in our experiment are treated with audit agency information about local municipal government their assessment of corruption levels in their community increases quite dramatically. In our experiment, treatment, T_i , is randomly assigned to subjects. Treated subjects, $T_i = 1$, receive CGO information about actual corruption levels in the respondent’s municipal government, i.e., $\text{Corrupt}_{ti}^{\text{True}}$. In the upper left panel of Figure 3 we observe a large significant effect for all three outcome metrics – posterior corruption beliefs for treated respondents are significantly higher, on average, than their priors while there is no significant change, on average, for participants

¹³We pre-registered our experimental design and analysis plan at the AEA RCT Registry [link excluded to preserve anonymity].

in the placebo arm.¹⁴ The difference between the control and treatment groups corresponds to approximately 0.31 standard deviations in the Qualitative outcome, 0.29 in Resources, and 0.12 in the Distribution outcome.¹⁵ Regression coefficients for this treatment variable in Table 2 confirm this strong treatment effect. Controlling for malfeasance priors, citizens update when they have access to audit information.

Magnitudes. The bottom panel in Figure 3, presents average updating broken down by four bins of our standardized measure of malfeasance (on the x-axis). Figure 3 distinguishes between respondents living in municipalities where malfeasance increase (Negative) versus those where it declined (Positive). For the majority of subjects who observed an increase in reported malfeasance, we see that subjects in the lowest malfeasance quartile bin exhibited significantly lower increases in malfeasance beliefs. However, across all the other three standardized malfeasance groups, there is no statistically significant difference in average belief updating. For respondents in municipalities with improving audit reports, there is no statistically significant difference in updating across all four malfeasance magnitude bins.

For our three outcome variables, we specify a *Malfeasance Magnitude* model that includes the standardized measure of malfeasance and its interaction with the treatment variable. Table 2 reports the regression results for three specifications of each outcome: the information treatment alone, the treatment interacted with standardized malfeasance magnitude, and the treatment interacted with the negative benchmark. Across all three outcome variables, the interaction between malfeasance magnitude and the treatment is statistically insignificant. Treated subjects believe corruption levels are higher post-intervention, but this updating is essentially uncorrelated with the magnitude of the reported malfeasance: the additional effect is 0.07 SD for the Qualitative outcome and 0.01 SD for Resources. The point estimate for Distribution is larger (0.39 SD) but imprecisely estimated and statistically indistinguishable from zero.

¹⁴The “on average” is important here because as we see below, respondents with high versus low priors, within both these treatment arms, respond quite differently.

¹⁵The outcomes were standardized using the full-sample standard deviation.

Benchmarking. In the second frame of the treatment video, subjects received a temporal comparison of their 2020 municipal audit results with the municipality’s previous audit. Of the 116 municipalities in the study, 94 registered worse malfeasance in 2020 than in their previous audit (the *Negative* condition); in 22, malfeasance was unchanged or declined (the *Positive* condition). If beliefs respond to whether things have gotten worse or better, updating should be larger in the Negative condition than in the Positive condition.

The upper-right frame of Figure 3 reveals a clear benchmarking asymmetry. Treated subjects in both conditions hold higher corruption posteriors than their priors, but the response to negative information is substantially larger than the response to positive information. Among subjects informed that malfeasance increased (Negative), belief updating rises by approximately 0.32 standard deviations relative to the placebo group for the Qualitative and Resource outcomes, and by 0.34 standard deviations for the Distribution outcome. Among subjects informed that malfeasance declined (Positive), the corresponding differences fall to roughly 0.23 standard deviations for the Qualitative outcome, 0.12 for the Resource outcome, and 0.08 for the Distribution outcome. For the two incentivized outcomes, updating in the Positive condition is statistically indistinguishable from the placebo group: positive benchmarking does not move incentivized beliefs, whereas negative benchmarking does.


In Table 2 we test this asymmetry by interacting the treatment indicator with a *Negative* benchmark dummy (Positive is the reference category). The interaction term is positive and most precisely estimated for the incentivized outcomes, consistent with Figure 3. For the Qualitative outcome the asymmetry is modest, because subjects revise upward under positive benchmarking as well; for the Distribution and Resource outcomes the asymmetry is pronounced, since only negative benchmarking produces detectable movement. Citizens thus treat the audit message as a directional signal: a report that conditions have worsened produces robust upward revision, including on the incentivized measures, whereas a report that conditions have improved produces at most a weak qualitative response and no movement on the incentivized measures.¹⁶

¹⁶The positive-condition estimates should be interpreted with caution. Only 22 municipalities recorded audits that improved compared to their previous CGO audit, resulting in slightly fewer than 500 subjects classified in the

This pattern of asymmetric benchmarking (negative weighing more heavily than positive information) is consistent with other evidence (Barron, 2021; Zhou & Oostendorp, 2014) and reinforces our broader interpretation that belief updating here is heuristic rather than calibrated: subjects register the *sign* of the temporal comparison far more readily than its magnitude.

positive category. Power calculations are provided in appendix Section A.13.

Figure 3: Audit Treatment Effects.



Note: This figure reports average belief updating (Posterior – Prior) by treatment, benchmarking condition, and malfeasance magnitude. *Information* (top left) shows overall updating by group. *Benchmarking* (top right) compares *Negative* and *Positive* conditions with their placebos. *Magnitudes* (bottom) shows updating by treatment, split by standardized malfeasance levels.

The malfeasance video treatment effect, which increases corruption belief posteriors, is significantly higher across our three belief measures for subjects who are informed that their municipality’s audit results are worse than previous audits. For subjects informed that audit performance has improved, only the qualitative metric registers a significant difference from the placebo. Thus

Table 2: Regression Results of Malfeasance Beliefs - All Outcomes

	Qualitative			Resources			Distribution		
	Information	Magnitude	Benchmarking	Information	Magnitude	Benchmarking	Information	Magnitude	Benchmarking
Intercept	2.734*** (0.401)	2.717*** (0.402)	2.471*** (0.446)	2.302*** (0.662)	2.274*** (0.648)	2.198*** (0.642)	1.122*** (0.153)	1.125*** (0.152)	1.101*** (0.166)
Treat	0.749*** (0.079)	0.736*** (0.091)	0.575** (0.208)	0.740*** (0.088)	0.739*** (0.094)	0.321+ (0.167)	0.347*** (0.028)	0.327*** (0.035)	0.110 (0.067)
Stand - Malfeasance		0.490 (0.489)			0.590 (0.432)			0.217 (0.233)	
Treat:Stand - Malfeasance		0.179 (0.882)			0.012 (0.597)			0.302 (0.353)	
Negative			0.311 (0.212)			0.132 (0.174)			0.031 (0.084)
Treat:Negative			0.202 (0.227)			0.489* (0.190)			0.278*** (0.073)
Prior	0.467*** (0.020)	0.464*** (0.020)	0.466*** (0.020)	0.446*** (0.018)	0.444*** (0.018)	0.445*** (0.019)	0.421*** (0.024)	0.415*** (0.024)	0.420*** (0.026)
Covariates	✓	✓	✓	✓	✓	✓	✓	✓	✓
Num.Obs.	3439	3439	3439	3439	3439	3439	3439	3439	3439
R2	0.263	0.264	0.267	0.220	0.221	0.225	0.203	0.207	0.213
R2 Adj.	0.257	0.258	0.261	0.214	0.214	0.218	0.196	0.200	0.206

Note: This table reports regression estimates across three outcomes and specifications. *Stand-Malfeasance* is the min-max standardized malfeasance reported in the video (0–1). The *Benchmarking* column reports main and moderating effects for those receiving negative information (increasing malfeasance). Standard errors clustered at the commune level. * $p < .05$, ** $p < .01$, *** $p < .001$.

all subjects, regardless of whether audits have improved or worsened, hold on average higher corruption posteriors, although these are significantly higher in municipalities with worsening audit performance.

Information Gap. The content of the video treatment is hypothesized to reduce gaps between ground truth and subjects’ corruption beliefs ($\text{Corrupt}_{ti}^{\text{True}} - \text{Corrupt}_{ti}$). We label this gap *Corrup Diff* and generate a distinct term for both Resources and Distribution. A limitation is that pre- and post-treatment belief measures (Distribution and Resources) do not precisely reflect the malfeasance information reported in the video treatment, which consist of: 1.) the value of malfeasance expenditures and 2.) the foregone resources associated with them.¹⁷ This variation between the signal and the updating question could attenuate our estimates of the impact of information treatments on information gaps learning (Fan et al., 2026).


There is little evidence that belief updating reduces the subjects’ information gap. The upper panel in Figure 4 presents the average updating in the y-axis ($(\text{Corrupt}_{t+1i} - \text{Corrupt}_{ti})$) for the

¹⁷The information treatments were based on the ChatBot experiments that identified the optimal messaging strategy. These were conducted after the original pre-treatment surveys (that occurred four months before the intervention) and did not precisely match the pre-treatment belief measures.

placebo treatment arm and for those in the worsening versus improving malfeasance versions of the audit video treatments. Across all treatment arms, the gap between our corruption signal and priors on the x-axis ($(\text{Corrupt}_{ti}^{\text{True}} - \text{Corrupt}_{ti})$) is correlated with belief updating. These correlations suggest that the gap between prior beliefs and our corruption signal is not causing belief updating. One indication is the positive correlation between our variable, measuring the gap between priors and the malfeasance signal, and belief updating for subjects in the placebo treatment arm. They never receive the signal, but their updating pattern is identical to that of treated subjects who do receive the signal. The correlations we observed in Figure 4 are likely driven by regression to the mean. Those with relatively high (low) malfeasance priors register decreasing (rising) beliefs in post-treatment.

The histogram in the lower panel of Figure 4 indicates that Distribution priors are typically lower than actual corruption levels. Hence, the average for *Corrup Diff* is positive. The Resources histogram in Figure 4 suggests the opposite – they are typically higher than actual corruption levels. Hence, the average value for *Corrup Diff* is negative. On average, though, both Resource and Distribution measured malfeasance beliefs increase. As a result, for the Distribution measure, treated subjects, on average, have more accurate posterior beliefs – from an average 0.36 discrepancy pre-treatment to essentially zero in post-treatment. On the other hand, results for the Resources outcome variable suggests that subjects, on average, have less accurate beliefs about corruption in post-treatment – their accuracy goes from -1.61 to -2.22. Given that these two metrics have differently signed net belief accuracies in pre-treatment, our information gap model of updating predicts that, on average, Resource and Distribution beliefs would move in the opposite direction. In fact, they move in the same direction, and as a result, the Resource posteriors are less accurate than priors. Subjects, on average, respond to treatment (regardless of quantitative content) by worsening their assessment of corruption levels. Subjects do not appear to be using the quantitative content of the message to reduce their information gap.

Figure 4: Belief Updating by Gap Between Actual Malfeasance and Pre-Treatment Beliefs.



Note: This figure on Top-left: average belief updating (Posterior - Prior) against the prior-truth gap ($\text{Corrupt}_{ti}^{\text{True}} - \text{Corrupt}_{ti}$), ranging -9 to +9 for *Resources* and -3 to +3 for *Distribution*; negative values indicate overestimation, positive underestimation. Treated respondents split into *Negative* (told corruption increased) and *Positive* (told it decreased). Bottom panels show the prior-truth gap distribution for *Resources* (left) and *Distribution* (right). Regression estimates are reported in Table A.21.

Regression estimates in Table A.21 in the appendix confirm this conclusion. For the *Distribution* outcome, the estimated coefficient for the interaction term, $\text{Treat} \times \text{Corrup Diff}$, is positive and precisely estimated. This could be evidence of updating that narrows an information gap. But the results for the *Resources* outcome suggest regression to the mean. In the *Resources* model, all subjects update conditional on the *Corrup Diff* term (including placebo subjects who receive no information signal). Note that the estimated coefficient for the interaction term, $\text{Treat} \times \text{Corrup Diff}$ in the *Resource* model, is essentially zero. The difference between actual corruption levels

and subjects' priors is clearly correlated with belief updating in the Distribution model – but the overall evidence does not suggest updating aimed at reducing an information gap.

In contrast to our results, Arias et al. (2022) find that malfeasance revelations move beliefs in proportion to the prior–signal gap, closing the distance between priors and the audit signal. Two features of our design likely explain the different findings. First, we elicit priors directly from each subject four months before the intervention, whereas Arias et al. infer treatment-group priors from the control group's posteriors. Direct elicitation permits individual-level tests of the prior–signal interaction but adds measurement error from the four-month gap; the inferred-priors approach is more efficient under prior stability but renders prior-contingent updating partly model-dependent. Second, the treatments differ in form and content: Arias et al. delivered a single percentage-of-funds-irregular figure through coordinated leafleting and group meetings ahead of an election, whereas we deliver a 50-second WhatsApp video bundling a peso-denominated audit total, a foregone-vaccines equivalence, and a temporal benchmark. Crucially, both differences cut in the same direction: our noisier prior measure (directly elicited, four months before treatment) and our denser, lower-engagement delivery channel (a 50-second video bundling three quantities, against a single figure delivered through in-person campaigns) each make precise, prior-calibrated Bayesian updating harder to detect and leave more room for the coarse, heuristic updating we document. The contrast with Arias et al. is therefore consistent with, rather than a challenge to, our central claim that citizens register a general signal rather than calibrating to the specific magnitude.

More broadly, the divergence illustrates a challenge for the literature: when standalone trials differ simultaneously in prior elicitation, delivery channel, treatment content, electoral salience, and subject-pool composition, it is hard to know whether divergent findings reflect substantive heterogeneity in citizen behavior or differences in what is being measured and how intensively subjects engage with it. Coordinated, multi-site designs in the tradition of the Metaketa initiative (Dunning, Grossman, Humphreys, Hyde, et al., 2019) (especially those that hold channel and content fixed while randomly varying engagement effort, prior measurement, and treatment

composition) offer the most credible route to cumulative evidence on when audit information is processed in a Bayesian fashion and when, as here, it is processed heuristically.

High/Low Malfeasance Priors. For each outcome we split subjects at the mean prior within their municipal sample; a high score indicates a belief that malfeasance is high. Two forces move beliefs in Figure 5, and separating them is the key to reading it. The first is regression to the mean, visible in the placebo panels: absent any signal, high-prior subjects drift toward less perceived malfeasance and low-prior subjects toward more, by roughly symmetric amounts. The second is the treatment itself, which must be measured against that drift.


Read this way, the raw movements in the *Treatment* frame are less puzzling than they first appear. Low-prior subjects revise sharply toward more malfeasance, while high-prior subjects continue their placebo-like drift toward less, only more slowly. Measured against placebo, however, the treatment pushes *both* groups toward higher perceived malfeasance: the *Treat* coefficient is positive in both subsamples. The visual asymmetry arises because the uniform push lands on groups drifting in opposite directions, and because low-prior subjects have more headroom to move. The *Benchmarking* and *Magnitude* panels repeat the pattern: raw movement toward more malfeasance is concentrated among low-prior subjects in every condition and at every level of reported malfeasance, while high-prior movement stays muted throughout.

Table A.22 confirms that this asymmetry in raw updating is not a heterogeneous response to the signal. Against placebo, the core *Treat* effect is positive, significant, and nearly identical across subsamples (.34 and .35 on the incentivized distribution measure, both $p < .001$), and every *Treat* $\times$ *Malfeasance* interaction is indistinguishable from zero in both. The one margin on which priors condition the response is *Benchmarking*: the pooled effect of learning that an audit worsened is driven almost entirely by low-prior subjects ($Treat \times Negative = 0.380$, $p < .001$, versus 0.003 for high priors on the distribution measure; 0.887, $p < .01$, versus -0.116 on resources). For subjects who already believe corruption is high, news of deterioration is merely confirming.

The implication for Arias et al. (2022) follows directly. Rather than calibrating to the prior–

signal gap, the audit registers a common “corruption exists here” impression whose size is largely independent of where the prior sat: the treatment effect neither varies systematically with the prior nor scales with the reported amount. What looks like prior-contingent updating in the raw data is placebo drift plus a roughly uniform shift. The benchmarking response among low-prior subjects is the lone trace of gap-based updating in our data.

Figure 5: Audit Treatment Updating Effects by Low versus High Priors.



Note: This figure shows average belief updating (Posterior – Prior) by treatment status, benchmarking, and standardized malfeasance levels. Within each group, respondents are split by prior beliefs: *Low* (below-mean within municipality) and *High* (above-mean). Top left (Information): updating for treated and placebo groups by prior level. Top right (Benchmarking): treated respondents only, split by benchmark type (*Negative* or *Positive*) and prior level. Bottom (Magnitude): average updating for treated respondents by standardized malfeasance levels and prior group. Dashed grey lines mark placebo drift: the average movement of placebo-group beliefs in the same condition and prior group. Placebo drift does not differ significantly by benchmark condition ($p = .118$) or by reported magnitude ($p = .508$).

Partisanship


We now assess the two different predictions regarding how partisanship conditions belief updating. Motivated reasoning (*H6a*) implies that copartisans of the incumbent discount negative audit information, so the audit’s effect should be attenuated – or absent – among copartisans, producing a treatment $\times$ partisanship interaction. Parallel persuasion (*H6b*) implies the opposite: copartisans and non-copartisans begin from different priors but revise their beliefs by a *common* amount, leaving the pre-treatment partisan gap intact after treatment (Coppock, 2023). The data favor parallel persuasion *H6b* and provide no support for *H6a*. Across all three outcomes the audit moves copartisans and non-copartisans by statistically indistinguishable amounts, while the partisan gap in belief *levels* present before treatment persists afterward.

Figure 6 presents the evidence: Panel A is a schematic, following Coppock (2023): each partisan group is represented by a line running from its average prior to its average posterior. Two features carry the interpretation. The *slope* of each line is the average causal effect of the audit; if the two lines are *parallel*, the treatment moves both groups by the same amount. The *vertical distance* between the lines is the descriptive partisan gap in beliefs; if that distance is preserved from prior to posterior, the audit informs both groups without closing the gap between them. Parallel persuasion is, literally, the case in which the two lines are parallel and offset.

The remaining panels confirm this parallel persuasion schematic. Across all three outcome measures, copartisans (who hold systematically lower corruption beliefs) and non-copartisans (who hold higher beliefs) trace near-parallel trajectories from prior to posterior: the audit shifts both groups upward by similar magnitudes, and the partisan gap at baseline is reproduced almost exactly at the posterior. The gap is widest for the non-incentivized *Qualitative* measure and narrower for the two incentivized measures (*Resources* and *Distribution*), consistent with evidence that incentivizing factual questions compresses partisan divergence (Bullock et al., 2015; Prior et al., 2015; Rathje et al., 2023). Crucially, that compression operates on the *level* of beliefs (the intercepts), not on the audit’s *effect* (the slopes), which remain similar across partisan groups whether

or not the outcome is incentivized.

Figure 6: Parallel Persuasion: Corruption Priors and Posteriors for Copartisans and Non-Copartisans.



Note: Panel A shows the parallel-persuasion schematic of Coppock (2023): slope is the average causal effect; intercept differences reflect descriptive differences in beliefs between groups. Remaining panels show prior and posterior beliefs by treatment and partisanship across the three outcomes, split into copartisans (support incumbent) and non-copartisans. *Negative/Positive* indicate increasing/declining malfeasance; *No Information* is baseline priors.

The model in Table 3 interacts the treatment indicator with copartisanship (and, in the three-way term, with standardized malfeasance), so the interaction coefficients isolate whether the audit's effect is larger for one partisan group than the other (Olden & Møen, 2022). The evidence for parallel persuasion is clear on both of its predictions. First, the treatment effect is common to both groups: the $\text{Treat} \times \text{Copartisan}$ interaction is small and statistically insignificant for all three outcomes (0.03, 0.14, -0.02), so we cannot reject that the audit moves copartisans and non-copartisans equally; this is the basis for rejecting motivated reasoning (H6a). Second, the partisan gap in belief levels is large and significant (the Copartisan main effect is -0.84 , -0.52 , and -0.26 for the Qualitative, Resources, and Distribution measures) and, as Figure 6 shows, survives treatment. It survives even for the incentivized outcomes that tend to shrink partisan differences on factual questions (Bullock et al., 2015; Prior et al., 2015; Rathje et al., 2023). That the gap survives incentives is itself informative: if the residual difference were merely partisan cheerleading, incentives should have eliminated it; instead it reflects a durable difference in what copartisans and non-copartisans believe to be true about their local government.

Table 3: Triple Interaction - Malfeasance

	Qualitative	Resources	Distribution
Intercept	3.062*** (0.387)	2.360*** (0.638)	1.215*** (0.150)
Treat	0.710*** (0.095)	0.694*** (0.117)	0.326*** (0.041)
Co-partisan	-0.837*** (0.168)	-0.524** (0.173)	-0.259*** (0.060)
Stand - Malfeasance	0.155 (0.420)	0.093 (0.309)	0.076 (0.275)
Treat:Stand - Malfeasance	0.337 (0.671)	0.436 (0.473)	0.409 (0.318)
Treat:Co-partisan	0.029 (0.181)	0.140 (0.191)	-0.019 (0.074)
Treat:Co-partisan:Stand - Malfeasance	-0.771 (1.778)	-1.914 (2.474)	-0.562 (0.488)
Prior	0.419*** (0.020)	0.434*** (0.018)	0.381*** (0.024)
Covariates	✓	✓	✓
Num.Obs.	3414	3414	3414
R2	0.283	0.228	0.224
R2 Adj.	0.276	0.220	0.216

Note: This table presents models including a three-way interaction between treatment condition, standardized malfeasance, and copartisan status. Standardized malfeasance measures the malfeasance level described in the treatment. Copartisans are respondents who intended to vote for the incumbent; non-copartisans include those undecided or supporting a different candidate. Standard errors are clustered at the commune level. * $p < .05$, ** $p < .01$, *** $p < .001$.

The benchmarking results require one qualification. Stratifying the sample by partisanship (Appendix Table A.15, Panel B), the negative benchmark is statistically significant among non-copartisans for all three outcomes ($\hat{\beta}^{Treat:Negative} = 0.43, 0.59, 0.32$) but not among copartisans. Further analysis, though, that estimates a three-way $Treat \times Negative \times Copartisan$ interaction (reported in appendix Table A.16) indicates that the difference between the groups is indistinguishable from zero. The contrast in the stratified estimates arises mechanically because the copartisan subsample is small ($N=929$ vs. $2,510$) and its benchmarking coefficients, though insignificant, are positive and of comparable magnitude. The data are therefore consistent with a common benchmarking response that is simply estimated less precisely among copartisans, and we read the benchmarking evidence as suggestive of greater non-copartisan sensitivity rather than as statistically distinguishable partisan heterogeneity. The same conclusion holds for the magnitude and information-gap specifications (Appendix Table A.15, Panels C–D): individual subgroup coefficients occasionally diverge in sign, but the corresponding three-way interactions are again insignificant.

Two conclusions follow. First, the audit’s effect on corruption beliefs is statistically common to copartisans and non-copartisans – the central prediction of parallel persuasion and the basis for rejecting motivated reasoning (H6a). Second, parallel persuasion also predicts a persistent partisan gap in belief levels, which we observe in both Figure 6 and Table A.15, and which survives even for the incentivized outcomes that tend to shrink partisan differences on factual questions. That the gap survives incentives is itself informative: if the residual partisan difference were merely partisan cheerleading, incentives should have eliminated it; instead it reflects a durable difference in what copartisans and non-copartisans believe to be true about their local government.

Short-term or Persistent Updating

Belief updating was measured one week and one month after the information treatment, and the effect persists.¹⁸ Dividing one-month by one-week estimates, overall persistence ratios range from

¹⁸Recent efforts to measure the persistence of information treatments typically take additional measurements four to eight weeks out (Cavallo et al., 2017; Haaland & Roth, 2021; Kuziemko et al., 2015).

0.40 to 0.71, and from 0.50 to 0.73 under the negative benchmark (Figure A.17). This compares favorably with Coppock (2017), in which roughly half of an effect survives ten days: 14 of our 21 estimates persist at least 30 days.¹⁹ Differential attrition does not account for the follow-up estimates: treated attriters are, if anything, more pessimistic than retained treated respondents, so our persistence estimates are plausibly lower bounds. The exception is the positive benchmark, where persistence ratios fall to between 0 and 0.35 and average beliefs for the Resources outcome reverse sign (Figures A.17 and A.18). Two factors plausibly drive this attenuation: information acquired through the treatment fades from memory over a month (Coppock, 2017), and exposure to counterarguments can offset it in a politically charged environment (D. Chong & Druckman, 2010); the latter bears especially on the positive arm, since our intervention coincided with municipal election campaigns in which ambient corruption messaging contradicts good news while reinforcing bad. The positive condition is also our least precisely estimated arm, so we read its instability cautiously.

Political Outcomes

A core conjecture is that information regarding corruption prevalence in municipal governments will have consequences for public goods provision and incumbent vote. We proxy support for public goods provision with an incentivized contribution to a municipal public good. In both pre- and post-treatment, we give subjects the opportunity to make a 500 CLP (around 15% of their subject payments) contribution to their local municipality.²⁰ We treat this as an indicator of our subjects' willingness to contribute to the cost of providing local public goods (Beekman et al., 2014; Bonica, 2018). And we measure voting behavior with reported intention to vote for the incumbent government party. Both pre- and post-treatment we asked subjects to indicate whether they would vote for the incumbent mayor.

Table 4 presents ITT estimates regressing each post-treatment outcome on the randomly assigned treatment indicator, controlling for the pre-treatment value of the same outcome and, in the

¹⁹Figure A.18 corroborates this at the level of average *beliefs*.

²⁰500 CLP is around 50 cents.

Table 4: Regression Results for Donations and Support for the Incumbent

	Donations				Support for the Incumbent			
	(1)	(2)	(3)	(4)	(1)	(2)	(3)	(4)
(Intercept)	0.171*** (0.010)	0.073*** (0.009)	0.089*** (0.013)	0.070 (0.059)	0.282*** (0.012)	0.132*** (0.010)	0.118*** (0.013)	0.120+ (0.070)
Treat	-0.019 (0.013)	-0.017 (0.012)	-0.017 (0.012)	-0.012 (0.012)	-0.054*** (0.015)	-0.046*** (0.013)	-0.043*** (0.013)	-0.045*** (0.013)
Donation baseline		0.309*** (0.015)	0.310*** (0.015)	0.294*** (0.015)				
Partisanship center (4-6)			-0.030* (0.013)				0.006 (0.013)	
Partisanship Right (7-10)			0.004 (0.019)				0.075*** (0.022)	
Voting incumbent baseline						0.537*** (0.016)	0.532*** (0.017)	0.535*** (0.017)
Covariates	✓	✓	✓	✓	✓	✓	✓	✓
Num.Obs.	3592	3592	3592	3432	3592	3592	3592	3432
R2	0.001	0.154	0.156	0.173	0.004	0.308	0.311	0.316
R2 Adj.	0.000	0.154	0.155	0.169	0.003	0.308	0.310	0.312

Note: This table reports the intention-to-treat effects of the treatment on donations to the municipality and voting for the incumbent using linear regression. Both outcomes are binary. Baseline measures from the pre-treatment survey are included as covariates where indicated. Standard errors are clustered at the commune level. * $p < .05$, ** $p < .01$, *** $p < .001$.

third specification, partisanship. The pattern is asymmetric: the treatment has no detectable effect on donations to a municipal public good, but produces a statistically significant 4.5 percentage-point reduction in stated intention to vote for the incumbent (Column 4).

Donations. The null donation effect is informative once placed against the design. Studies that detect a corruption-information effect on contributions typically embed the manipulation in environments that activate strategic and reputational motives, group public goods games with direct interaction (Beekman et al., 2014; Beesley & Hawkins, 2022), monitoring or sanctioning (Campos-Vazquez & Mejia, 2016), or in-group reputational pressure (Solaz et al., 2019). Our donation environment is, by contrast, a purely voluntary contribution mechanism: anonymous, unmonitored, and free of any sanctioning. It isolates intrinsic prosocial motivation from those strategic and reputational channels. Read this way, our null is not a contradiction of the prior literature but a demarcation of where information about local malfeasance does and does not move behavior: it does not, on its own, depress unconditional intrinsic giving to the implicated institution.

Donation rates fall from baseline to post-treatment in both arms. Factors likely contributing

to this decline include “warm glow fatigue” (a second identical appeal generates less warm glow than the first (Merritt et al., 2010)) and a post-COVID wave of deteriorating household finances (reported income in our sample declines across waves (Table 1)). These forces are common to both arms: whatever they contribute to posteriors, they contribute to treated and placebo subjects alike, so the between-arm difference nets them out and remains the relevant causal test.

Incumbent vote. As with donations, stated incumbent support declines from baseline to post-treatment in both arms. This tracks a documented national swing rather than a survey artifact: our panel spans the post-2019 legitimacy crisis and the run-up to the May 2021 elections, in which incumbents and traditional parties suffered historic losses (Martínez & Olivares, 2022). Again, the common trend nets out in the between-arm contrast, which is where the causal interpretation lives.

The estimate in Table 4 is net of this common decline: relative to placebo subjects exposed to the same national trend, treated subjects report a further 4.5 percentage point reduction in support for the incumbent. This average effect masks meaningful partisan heterogeneity: support for the incumbent declines by roughly 7 points among copartisans and 2–3 points among non-copartisans. Partisan loyalty does not insulate incumbents from accountability and aligns with recent evidence that, conditional on accepting that wrongdoing occurred, copartisans sanction in-party incumbents at least as heavily as out-party voters do (Aarslew, 2023; Dulay & Lee, 2026). Two non-exclusive mechanisms are consistent with our data. First, copartisans begin from a substantially more optimistic baseline, about 55% rate the incumbent’s anti-corruption performance positively, compared with 8% of non-copartisans, so the audit message represents a larger gap between prior expectations and reported reality, and a correspondingly larger information shock. Second, work in the subjective group dynamics tradition (Abrams, 2008) holds that, once acknowledged, in-group wrongdoing is judged more severely because it violates the moral commitments that anchor partisan attachment. The two mechanisms are mutually reinforcing rather than competing.

Additional analyses examining alternative explanations, including the partisan composition of the negative benchmarking arm, issue-salience of corruption, partisan-attachment intensity, ideological asymmetry (left- vs. right-leaning partisans), and the local electoral-competition envi-

ronment (number of candidates, margins of victory), yield no evidence that any of these factors accounts for the differential effect. These are reported in Section A.8 in the appendix.

Discussion and Conclusions

Audit information causes belief updating. Receiving the WhatsApp video summarizing the CGO's audit of irregularities in their municipality leads Chileans to revise their malfeasance beliefs upward. The estimated treatment effect is large, precisely estimated, and consistent across all three outcomes. Most of these effects persist at one-month follow-up, most strongly for subjects who received the negative temporal benchmark, while effects under the positive benchmark attenuate and in places reverse sign. The shift is more pronounced among subjects whose priors initially understated corruption, but it does not require a low prior: average posteriors are more negative than priors in both improving and deteriorating municipalities.

This updating is heuristic rather than quantitatively precise. The video conveyed the magnitude of malfeasance: the absolute audited value and its equivalent in foregone municipal services. Yet belief updating did not vary with that magnitude. Magnitude was not randomly assigned, so the null magnitude correlation is descriptive rather than causal: municipalities with larger audit irregularities likely differ on unobserved dimensions (prior experience with corruption, media exposure, local political context) that independently shape responses to information, and a clean test of sensitivity to scale would require an experiment that randomized the magnitude conveyed while holding all other treatment content constant. There was other evidence on heuristic processing. Subjects' priors and posteriors do not converge on ground truth: The gap widens for the Resources metric and the narrowing for the Distribution metric is too weak to indicate genuine convergence. The prior-signal gap does not predict the direction or magnitude of updating. The only treatment component that produced differential updating was the temporal benchmark: subjects told their municipality's audit had worsened relative to the previous cycle revised more sharply than those told it had improved.

Our results diverge from the Bayesian pattern documented by Arias et al. (2022), whose voters update in line with their prior–signal gap. As developed in the Information Gap section, the two studies differ on dimensions any one of which could account for the divergence: how priors are measured, the delivery and cognitive load of the treatment, and unmeasured engagement with the information. Standalone trials cannot discriminate among these accounts. The design prescription follows: coordinated, multi-site studies in the Metaketa tradition (Dunning, Grossman, Humphreys, Hyde, et al., 2019) that hold delivery and content constant while varying engagement, treatment composition, and prior measurement, offer the most credible path to cumulative evidence on belief updating.

Partisanship does not undermine the causal effect of malfeasance information on belief updating. The treatment shifts beliefs upward for both copartisans and non-copartisans of the incumbent mayor, at roughly comparable magnitudes. We find no evidence that copartisans defend the incumbent by resisting unwelcome audit information: what partisanship governs is the level from which citizens start, not their responsiveness to evidence. This “parallel persuasion” pattern (Coppock, 2023, Figure 7) is encouraging in one respect: credible, concrete information about local government cuts through partisan filters at the stage of belief revision, much as a heuristic “corruption exists here” signal would be registered broadly across the electorate. It is sobering in another: the treatment does not close the partisan gap. Copartisans begin with less negative beliefs about local malfeasance than non-copartisans, and remain so after treatment. The persistence of this gap confirms that financial incentives for accuracy can mute, but not eliminate, partisan motivated reasoning on factual questions (Rathje et al., 2023). In our data the partisan baseline gap is in fact larger for the non-incentivised qualitative outcome than for the two incentivised quantitative measures, yet a gap remains for all three. Audits can move beliefs evenhandedly across partisan lines, but because they leave that gap intact, they bound the impact of transparency interventions.

The joint response of beliefs and political behavior to audit information is, we argue, a core finding of the paper rather than a peripheral one. The intent-to-treat estimates reveal an asymmetric pattern: a modest but statistically significant reduction in stated incumbent vote intention,

alongside no detectable effect on costly donations to municipal public goods. The electoral effect is consistent with an accountability mechanism: citizens who learned of malfeasance became less willing to support the incumbent mayor. The absence of any donation effect reveals the limits of that response: citizens revised their political evaluations of the incumbent but did not alter their material engagement with the institution implicated in wrongdoing. These are total effects of the randomized information treatment on each outcome; because belief updating is not itself randomly assigned, we do not estimate whether the electoral response is mediated by beliefs, and the parallel movement of beliefs and vote intention is consistent with, but does not demonstrate, mediation.

This asymmetry sits comfortably within a broader literature that struggles to identify a link between information and accountability behavior. The Dunning, Grossman, Humphreys, Hyde, et al. (2019) meta-analysis of coordinated field trials finds little evidence that voter information campaigns produce electoral accountability even when they shift stated beliefs. A. Chong et al. (2015) document that anti-corruption messaging in Mexico can demobilize voters, an accountability reversal rather than an improvement. Boas et al. (2019) and Enríquez et al. (2024) similarly find limited translation of malfeasance information into behavior. Our contribution is to isolate belief updating as a distinct causal stage but to note that the political behavior associated with the same information treatments responds differently: attitudinal evaluations of the incumbent (vote intention) respond as expected but a costly material commitment to the municipality (donation) does not.

Implications Why is updating heuristic rather than calibrated? Several non-exclusive mechanisms are consistent with our results. The first is cognitive: the quantitative content of an audit treatment imposes a non-trivial demand on numeracy. Recent evidence indicates that numeracy skills are correlated with correct engagement with quantitative policy data (Stagnaro et al., 2023) and that more accessible representations of such data substantially increase comprehension (Yu & Opfer, 2024). More broadly, low cognitive ability has been shown to undermine the stability and coherence of policy attitudes (Panish, 2025). We do not measure numeracy or cognitive ability di-

rectly, but our pattern of results, uniform upward updating that does not track magnitude, is exactly what such accounts would predict.

A second possibility is dosage. A one-time exposure, even when packaged in an accessible format, may simply be insufficient to overcome the cognitive demands of audit content. Bowles et al. (2025) report that sustained exposure to fact-checking information delivered via WhatsApp induced substantial internalization of fact-checked content; Grossman et al. (2024) similarly find that sustained transparency improves electoral accountability. Our short, single-treatment design may be near the lower bound of what audit transparency interventions can deliver in terms of belief calibration.

A third possibility is the transparency paradox itself: messages about corruption may shift attitudes primarily through priming rather than learning (Mistree & Dibley, 2018; Peiffer, 2020). Under this account, exposure to audit information increases perceived corruption, lowers efficacy, reduces willingness to engage in conventional corrective actions, and can even raise willingness to bribe (Cheeseman & Peiffer, 2020; A. Chong et al., 2015). For our setting, well-intentioned information campaigns may simply prime widely-shared beliefs that municipal corruption is pervasive, without triggering the cognitive effort that the policy-design rationale for transparency requires. Such priming may persuade citizens that “all politicians are corrupt”, fostering abstention or exit toward new, anti-system alternatives rather than accountability within the existing party system (A. Chong et al., 2015; De Vries & Solaz, 2017).

The agenda for next-generation studies is innovation in both content and delivery of anti-corruption information. One direction is to simplify and contextualize the key facts further, using vivid visual aids, analogies, or personalized impact statements. A second is repetition and reinforcement: citizens may need to encounter the information multiple times or through multiple channels before it sinks in (Bowles et al., 2025). A third is deliberation: encouraging public discussion of audit findings may prompt careful processing that passive listening does not. But the most important methodological lesson, we believe, is the need for progress on research design. The most direct test of our heuristic-processing interpretation would be an experiment that holds sub-

stantive audit content constant while randomly varying the cognitive load imposed on subjects, and an experiment that randomizes the scale of reported irregularities directly – ideally a coordinated, Metaketa-style program (Dunning, Grossman, Humphreys, Hyde, et al., 2019) that systematically varies a small number of design dimensions across sites and subject pools.

Our study suggests that the accountability gains from transparency are real but bounded. Audit information moves beliefs and shifts stated electoral preferences. It does not, on its own, calibrate beliefs to the scale of wrongdoing, elicit costly civic contribution, or close the partisan gap in baseline corruption perceptions. Making government audit data public is not the same as making it meaningful to the public. Strengthening democratic accountability requires closing the gap between information availability and citizens’ capacity (and willingness) to engage with detail.²¹

²¹**Data Availability Statement.** The replication data and code supporting the findings of this study will be deposited in a public repository; the link is omitted to preserve anonymity during peer review. The pre-registration is likewise available and the link withheld for review.

References

- Aarslew, L. F. (2023). Why don't partisans sanction electoral malpractice? *British Journal of Political Science*, 53(2), 407–423. <https://doi.org/10.1017/S0007123422000126>
- Anduiza, E., Gallego, A., & Muñoz, J. (2013). Turning a blind eye: Experimental evidence of partisan bias in attitudes toward corruption. *Comparative Political Studies*, 46(12), 1664–1692.
- Arel-Bundock, V., Blais, A., & Dassonneville, R. (2021). Do voters benchmark economic performance? *British Journal of Political Science*, 51(1), 437–449. <https://doi.org/10.1017/S0007123418000236>
- Arias, E., Balán, P., Larreguy, H., Marshall, J., & Querubín, P. (2019). Information provision, voter coordination, and electoral accountability: Evidence from Mexican social networks. *American Political Science Review*, 113(2), 475–498. <https://doi.org/10.1017/S0003055419000091>
- Arias, E., Larreguy, H., Marshall, J., & Querubín, P. (2022). Priors Rule: When Do Malfeasance Revelations Help Or Hurt Incumbent Parties? *Journal of the European Economic Association*, 20(4), 1433–1477.
- Arias, E., Larreguy, H. A., Marshall, J., & Querubín, P. (2025). Does the content and mode of delivery of information matter for electoral accountability? evidence from a field experiment in Mexico. *Latin American Economic Review*, 34, 1–41. <https://doi.org/10.60758/laer.v34i.335>
- Avenburg, A. (2019). Public costs versus private gain: Assessing the effect of different types of information about corruption incidents on electoral accountability. *Journal of Politics in Latin America*, 11(1), 71–108. <https://doi.org/10.1177/1866802X19840457>
- Aytaç, S. E. (2018). Relative economic performance and the incumbent vote: A reference point theory. *The Journal of Politics*, 80(1), 16–29. <https://doi.org/10.1086/693908>

- Baillon, A., Bleichrodt, H., & Granic, G. D. (2022). Incentives in surveys. *Journal of Economic Psychology*, 93, 102552.
- Bandiera, O., Prat, A., & Valletti, T. (2009). Active and passive waste in government spending: Evidence from a policy experiment. *American Economic Review*, 99(4), 1278–1308. <https://doi.org/10.1257/aer.99.4.1278>
- Banerjee, A., Green, D. P., McManus, J., & Pande, R. (2014). Are poor voters indifferent to whether elected leaders are criminal or corrupt? a vignette experiment in rural India. *Political Communication*, 31(3), 391–407. <https://doi.org/10.1080/10584609.2014.914615>
- Barro, R. (1973). The control of politicians: An economic model. *Public Choice*, 14(1), 19–42. <https://doi.org/10.1007/BF01718440>
- Barron, K. (2021). Belief updating: Does the ‘good-news, bad-news’ asymmetry extend to purely financial domains? *Experimental Economics*, 24(1), 31–58.
- Bastianello, F., & Imas, A. (2026). *Biases in belief updating within and across domains* (Working Paper No. 34638). National Bureau of Economic Research. <https://doi.org/10.3386/w34638>
- Becher, M., Brouard, S., & Stegmueller, D. (2023). Endogenous benchmarking and government accountability: Experimental evidence from the covid-19 pandemic. *British Journal of Political Science*, 1–18. <https://doi.org/10.1017/S0007123423000170>
- Beekman, G., Bulte, E., & Nillesen, E. (2014). Corruption, investments and contributions to public goods: Experimental evidence from rural Liberia. *Journal of Public Economics*, 115, 37–47.
- Beesley, C., & Hawkins, D. (2022). Corruption, institutional trust and political engagement in peru. *World Development*, 151, 105743.
- Bénabou, R., & Tirole, J. (2002). Self-confidence and personal motivation. *The Quarterly Journal of Economics*, 117(3), 871–915.

- Bénabou, R., & Tirole, J. (2016). Mindful economics: The production, consumption, and value of beliefs. *Journal of Economic Perspectives*, 30(3), 141–64. <https://doi.org/10.1257/jep.30.3.141>
- Berliner, D., & Wehner, J. (2022). Audits for accountability: Evidence from municipal by-elections in South Africa. *The Journal of Politics*, 84(3), 1581–1594. <https://doi.org/10.1086/716951>
- Bhandari, A., Larreguy, H., & Marshall, J. (2023). Able and mostly willing: An empirical anatomy of information’s effect on voter-driven accountability in Senegal. *American Journal of Political Science*, 67(4), 1040–1066. <https://doi.org/https://doi.org/10.1111/ajps.12591>
- Bisgaard, M. (2019). How getting the facts right can fuel partisan-motivated reasoning. *American Journal of Political Science*, 63(4), 824–839. <https://doi.org/10.1111/ajps.12432>
- Blinder, S., & Schaffner, B. F. (2019). Going with the flows: Information that changes Americans’ immigration preferences. *International Journal of Public Opinion Research*, 32(1), 153–164. <https://doi.org/10.1093/ijpor/edz007>
- Boas, T. C., Hidalgo, F. D., & Melo, M. A. (2019). Norms versus action: Why voters fail to sanction malfeasance in Brazil. *American Journal of Political Science*, 63(2), 385–400. <https://doi.org/10.1111/ajps.12413>
- Bobonis, G. J., Cámara Fuertes, L. R., & Schwabe, R. (2016). Monitoring corruptible politicians. *American Economic Review*, 106(8), 2371–2405. <https://doi.org/10.1257/aer.20130874>
- Bonica, A. (2018). Are donation-based measures of ideology valid predictors of individual-level policy preferences? *The Journal of Politics*, 81. <https://doi.org/10.1086/700722>
- Botero, S., Cornejo, R. C., Gamboa, L., Pavao, N., & Nickerson, D. W. (2015). Says who? an experiment on allegations of corruption and credibility of sources. *Political Research Quarterly*, 68(3), 493–504. <https://doi.org/10.1177/1065912915591607>
- Bowles, J., Croke, K., Larreguy, H., Liu, S., & Marshall, J. (2025). Sustaining exposure to fact-checks: Misinformation discernment, media consumption, and its political implications. *American Political Science Review*, 119(4), 1864–1887. <https://doi.org/10.1017/S0003055424001394>

- Breitenstein, S. (2019). Choosing the crook: A conjoint experiment on voting for corrupt politicians. *Research & Politics*, 6(1), 2053168019832230. <https://doi.org/10.1177/2053168019832230>
- Bullock, J. G., Gerber, A. S., Hill, S. J., & Huber, G. A. (2015). Partisan bias in factual beliefs about politics. *Quarterly Journal of Political Science*, 10(4), 519–578.
- Buntaine, M. T., Jablonski, R., Nielson, D. L., & Pickering, P. M. (2018). Sms texts on corruption help ugandan voters hold elected councillors accountable at the polls. *Proceedings of the National Academy of Sciences*, 115(26), 6668–6673. <https://doi.org/10.1073/pnas.1722306115>
- Camerer, C. F., & Hogarth, R. M. (1999). The effects of financial incentives in experiments: A review and capital-labor-production framework. *Journal of Risk and Uncertainty*, 19(1), 7–42.
- Campos-Vazquez, R. M., & Mejia, L. A. (2016). Does corruption affect cooperation? a laboratory experiment. *Latin American Economic Review*, 25(1), 5.
- Carey, J. M., Chun, E., Cook, A., Fogarty, B. J., Jacoby, L., Nyhan, B., Reifler, J., & Sweeney, L. (2025). The narrow reach of targeted corrections: No impact on broader beliefs about election integrity. *Political Behavior*, 47(2), 737–750. <https://doi.org/10.1007/s11109-024-09968-0>
- Cavallo, A., Cruces, G., & Perez-Truglia, R. (2017). Inflation expectations, learning, and supermarket prices: Evidence from survey experiments. *American Economic Journal: Macroeconomics*, 9(3), 1–35. <https://doi.org/10.1257/mac.20150147>
- Charbonneau, É., & Van Ryzin, G. G. (2015). Benchmarks and citizen judgments of local government performance: Findings from a survey experiment. *Public Management Review*, 17(2), 288–304. <https://doi.org/10.1080/14719037.2013.798027>
- Cheeseman, N., & Peiffer, C. (2020). *The unintended consequences of anti-corruption messaging in Nigeria: Why pessimists are always disappointed* (tech. rep.). Anti-Corruption Evidence SOAS Consortium.

- Cheeseman, N., & Peiffer, C. (2021). The curse of good intentions: Why anticorruption messaging can encourage bribery. *American Political Science Review*, 1–15. <https://doi.org/10.1017/S0003055421001398>
- Chong, A., De La O, A. L., Karlan, D., & Wantchekon, L. (2015). Does corruption information inspire the fight or quash the hope? a field experiment in Mexico on voter turnout, choice, and party identification. *The Journal of Politics*, 77(1), 55–71.
- Chong, D., & Druckman, J. N. (2010). Dynamic public opinion: Communication effects over time. *American Political Science Review*, 104(4), 663–680. <https://doi.org/10.1017/S0003055410000493>
- Coibion, O., Georgarakos, D., Gorodnichenko, Y., & van Rooij, M. (2023). How does consumption respond to news about inflation? field evidence from a randomized control trial. *American Economic Journal: Macroeconomics*, 15(3), 109–52. <https://doi.org/10.1257/mac.20200445>
- Coibion, O., & Gorodnichenko, Y. (2015). Information rigidity and the expectations formation process: A simple framework and new facts. *American Economic Review*, 105(8), 2644–78. <https://doi.org/10.1257/aer.20110306>
- Consejo Nacional para la Transparencia. (2019). Estudio nacional de transparencia.
- Consejo Nacional para la Transparencia. (2020). Estudio nacional de transparencia.
- Coppock, A. (2017, February). *The persistence of survey experimental treatment effects* [Working Paper. Department of Political Science, Yale University]. https://alexandercoppock.com/coppock_2017b.pdf
- Coppock, A. (2023). *Persuasion in parallel*. University of Chicago Press.
- Corbacho, A., Gingerich, D. W., Oliveros, V., & Ruiz-Vega, M. (2016). Corruption as a self-fulfilling prophecy: Evidence from a survey experiment in Costa Rica. *American Journal of Political Science*, 60(4), 1077–1092. <https://doi.org/10.1111/ajps.12244>
- Cornejo, R. C. (2022). Same scandal, different interpretations: Politics of corruption, anger, and partisan bias in Mexico. *Journal of Elections, Public Opinion and Parties*, 0(0), 1–22.

- Costello, T. H., Pennycook, G., & Rand, D. G. (2024). Durably reducing conspiracy beliefs through dialogues with AI. *Science*, 385(6714), eadq1814.
- Coutts, A. (2019). Good news and bad news are still news: Experimental evidence on belief updating. *Experimental Economics*, 22(2), 369–395. <https://doi.org/10.1007/s10683-018-9572-5>
- Cubel, M., Papadopoulou, A., & Sánchez-Pagés, S. (2024). Identity and political corruption: A laboratory experiment. *Economic Theory*.
- D’Acunto, F., & Weber, M. (2024). Why survey-based subjective expectations are meaningful and important. *Annual Review of Economics*, 16(Volume 16, 2024), 329–357. <https://doi.org/10.1146/annurev-economics-091523-043659>
- de Figueiredo, M. F., Hidalgo, F. D., & Kasahara, Y. (2022). When do voters punish corrupt politicians? experimental evidence from a field and survey experiment. *British Journal of Political Science*, 1–12. <https://doi.org/10.1017/S0007123421000727>
- de Quidt, J., Haushofer, J., & Roth, C. (2018). Measuring and bounding experimenter demand. *American Economic Review*, 108(11), 3266–3302. <https://doi.org/10.1257/aer.20171330>
- De Vries, C. E., & Solaz, H. (2017). The electoral consequences of corruption. *Annual Review of Political Science*, 20(Volume 20, 2017), 391–408.
- DellaVigna, S., & Gentzkow, M. (2010). Persuasion: Empirical evidence. *Annual Review of Economics*, 2(1), 643–669.
- Denly, M. (2022). Measuring corruption using governmental audits: A new framework and dataset [Accessed = 2022-11-01]. <https://mikedenly.com/research/audit-measurement>
- Duch, R., & Torres, F. (2020, May). Are voters updating when audit reports are informative? [Registered at the AEA RCT Registry].
- Duch, R. M., Loewen, P., Robinson, T. S., & Zakharov, A. (2025). Governing in the face of a global crisis: When do voters punish and reward incumbent governments? *Proceedings of the National Academy of Sciences*, 122(4), e2405021122. <https://doi.org/10.1073/pnas.2405021122>

- Dulay, D., & Lee, S. (2026). Partisanship versus principles: A study of voter response to intra-party corruption allegations. *Party Politics*, 32(1), 165–173.
- Dunning, T., Grossman, G., Humphreys, M., Hyde, S., McIntosh, C., & Nellis, G. (2019). *Information, accountability, and cumulative learning: Lessons from metaketa i*. Cambridge University Press.
- Dunning, T., Grossman, G., Humphreys, M., Hyde, S. D., McIntosh, C., Nellis, G., Adida, C. L., Arias, E., Bicalho, C., Boas, T. C., Buntaine, M. T., Chauchard, S., Chowdhury, A., Gottlieb, J., Hidalgo, F. D., Holmlund, M., Jablonski, R., Kramon, E., Larreguy, H., . . . Sircar, N. (2019). Voter information campaigns and political accountability: Cumulative findings from a preregistered meta-analysis of coordinated trials. *Science Advances*, 5(7).
- Eil, D., & Rao, J. M. (2011). The good news-bad news effect: Asymmetric processing of objective information about yourself. *American Economic Journal: Microeconomics*, 3(2), 114–38. <https://doi.org/10.1257/mic.3.2.114>
- Elia, E., & Schwindt-Bayer, L. A. (2022). Corruption perceptions, opposition parties, and reelecting incumbents in Latin America. *Electoral Studies*, 80, 102545.
- Enríquez, J. R., Larreguy, H., Marshall, J., & Simpser, A. (2024). Mass political information on social media: Facebook ads, electorate saturation, and electoral accountability in Mexico. *Journal of the European Economic Association*, 22(4), 1678–1722. <https://doi.org/10.1093/jeea/jvae011>
- Esberg, J., & Mummolo, J. (2018). *Explaining misperceptions of crime* [Unpublished manuscript].
- Fan, T. Q., Liang, Y., & Peng, C. (2026). *The inference-forecast gap in belief updating* [Forthcoming at *Econometrica*. Latest draft and experimental instructions available online]. <https://yuchengliang.com/ifgap/ifgap.pdf>
- Ferejohn, J. (1986). Incumbent performance and electoral control. *Public Choice*, 50, 5–25.
- Ferraz, C., & Finan, F. (2008). Exposing corrupt politicians: The effects of Brazil’s publicly released audits on electoral outcomes. *The Quarterly Journal of Economics*, 123(2), 703–745.

- Ferraz, C., & Finan, F. (2011). Electoral accountability and corruption: Evidence from the audits of local governments. *American Economic Review*, 101(4), 1274–1311. <https://doi.org/10.1257/aer.101.4.1274>
- Fiorina, M. P. (1981). *Retrospective voting in American national elections*. Yale University Press.
- Gagliarducci, S., & Manacorda, M. (2020). Politics in the family: Nepotism and the hiring decisions of Italian firms. *American Economic Journal: Applied Economics*, 12(2), 67–95. <https://doi.org/10.1257/app.20170778>
- Gerber, A. S., & Green, D. P. (2012). *Field Experiments: Design, Analysis, and Interpretation*. W.W. Norton & Company, Inc.
- Ghanem, D., Hirshleifer, S., & Ortiz-Beccera, K. (2023). Testing attrition bias in field experiments. *Journal of Human Resources*. <https://doi.org/10.3368/jhr.0920-11190R2>
- Gillen, B., Snowberg, E., & Yariv, L. (2019). Experimenting with measurement error: Techniques with applications to the caltech cohort study. *Journal of Political Economy*, 127(4), 1826–1863.
- Gonzalez-Ocantos, E., de Jonge, C. K., & Nickerson, D. W. (2014). The conditionality of vote buying norms: Experimental evidence from Latin America. *American Journal of Political Science*, 58(1), 197–211.
- Grossman, G., Michelitch, K., & Prato, C. (2024). The effect of sustained transparency on electoral accountability. *American Journal of Political Science*, 68(3), 1022–1040. <https://doi.org/10.1111/ajps.12787>
- Guess, A., & Coppock, A. (2020). Does counter-attitudinal information cause backlash? results from three large survey experiments. *British Journal of Political Science*, 50(4), 1497–1515. <https://doi.org/10.1017/S0007123418000327>
- Haaland, I., & Roth, C. (2021). Beliefs about racial discrimination and support for pro-black policies. *The Review of Economics and Statistics*, 1–38. https://doi.org/10.1162/rest_a_01036
- Haaland, I., Roth, C., & Wohlfart, J. (2023). Designing information provision experiments. *Journal of Economic Literature*, 61(1), 3–40. <https://doi.org/10.1257/jel.20211658>

- Healy, A., & Lenz, G. S. (2014). Substituting the end for the whole: Why voters respond primarily to the election-year economy. *American Journal of Political Science*, 58(1), 31–47. <https://doi.org/https://doi.org/10.1111/ajps.12053>
- Hicken, A., Leider, S., Ravanilla, N., & Yang, D. (2015). Measuring vote-selling: Field evidence from the Philippines. *American Economic Review*, 105(5), 352–56. <https://doi.org/10.1257/aer.p20151033>
- Hopkins, D. J., Sides, J., & Citrin, J. (2019). The muted consequences of correct information about immigration. *The Journal of Politics*, 81(1), 315–320.
- Jablonski, R. S., Buntaine, M. T., Nielson, D. L., & Pickering, P. M. (2021). Individualized text messages about public services fail to sway voters: Evidence from a field experiment on Ugandan elections. *Journal of Experimental Political Science*, 1–13.
- Jäger, S., Roth, C., Roussille, N., & Schoefer, B. (2024). Worker beliefs about outside options*. *The Quarterly Journal of Economics*, 139(3), 1505–1556. <https://doi.org/10.1093/qje/qjae001>
- Jahnke, B., & Weisser, R. A. (2019). How does petty corruption affect tax morale in Sub-Saharan Africa? *European Journal of Political Economy*, 60, 101751.
- Kalla, J. L., Broockman, D. E., & Sekhon, J. S. (2021). Field experiments with survey outcomes. In J. N. Druckman & D. P. Green (Eds.), *Advances in experimental political science* (pp. 56–78). Cambridge University Press.
- Kayser, M. A., & Peress, M. (2012). Benchmarking across borders: Electoral accountability and the necessity of comparison. *American Political Science Review*, 106(3), 661–684. <https://doi.org/10.1017/S0003055412000275>
- Kim, J. W. (2025). Evidence can change partisan minds but less so in hostile contexts. *British Journal of Political Science*, 55, e62. <https://doi.org/10.1017/S0007123424000899>
- Kovach, M. (2021). Conservative updating. <https://arxiv.org/abs/2102.00152>
- Krosnick, J. A. (1991). Response strategies for coping with the cognitive demands of attitude measures in surveys. *Applied Cognitive Psychology*, 5(3), 213–236.

- Kunda, Z. (1990). The case for motivated reasoning. *Psychological Bulletin*, 108(3), 480–498.
- Kuziemko, I., Norton, M. I., Saez, E., & Stantcheva, S. (2015). How elastic are preferences for redistribution? evidence from randomized survey experiments. *American Economic Review*, 105(4), 1478–1508. <https://doi.org/10.1257/aer.20130360>
- Lagunes, P. (2021, September). *The eye and the whip: Corruption control in the Americas*. Oxford University Press. <https://doi.org/10.1093/oso/9780197577622.001.0001>
- LAPOP. (2021). Chile: Americasbarometer 2021 [Vanderbilt University, Nashville, TN].
- Larreguy, H., Marshall, J., & Snyder, J. M., Jr. (2020). Publicising malfeasance: When the local media structure facilitates electoral accountability in Mexico. *The Economic Journal*, 130(631), 2291–2327. <https://doi.org/10.1093/ej/ueaa046>
- Larsen, M. V., & Olsen, A. L. (2020). Reducing bias in citizens’ perception of crime rates: Evidence from a field experiment on burglary prevalence. *The Journal of Politics*, 82(2), 747–752. <https://doi.org/10.1086/706595>
- Lee, D. S. (2009). Training, wages, and sample selection: Estimating sharp bounds on treatment effects. *The Review of Economic Studies*, 76(3), 1071–1102. <https://doi.org/10.1111/j.1467-937X.2009.00536.x>
- Letki, N., Górecki, M. A., & Gendźwił, A. (2023). ‘They accept bribes; we accept bribery’: Conditional effects of corrupt encounters on the evaluation of public institutions. *British Journal of Political Science*, 53(2), 690–697. <https://doi.org/10.1017/S0007123422000047>
- Liang, Y. (2025). Learning from unknown information sources. *Management Science*, 71(5), 3873–3890. <https://doi.org/10.1287/mnsc.2021.03551>
- Martínez, C. A., & Olivares, A. L. (2022). Chile 2021: entre un intenso calendario electoral y la acusación constitucional en contra de Sebastián Piñera. *Revista de Ciencia Política (Santiago)*, 42(2), 225–253. <https://doi.org/10.4067/S0718-090X2022005000119>
- Merritt, A. C., Effron, D. A., & Monin, B. (2010). Moral self-licensing: When being good frees us to be bad. *Social and Personality Psychology Compass*, 4(5), 344–357. <https://doi.org/10.1111/j.1751-9004.2010.00263.x>

- Mian, A., Sufi, A., & Khoshkhoh, N. (2021). Partisan bias, economic expectations, and household spending. *The Review of Economics and Statistics*, 1–46. https://doi.org/10.1162/rest_a_01056
- Mistree, D., & Dibley, A. (2018, April). *Corruption and the paradox of transparency* (Working Paper No. 03-2018) (Research Fellow and Lecturer (Mistree); JSD Candidate (Dibley), Stanford Law School). Stanford Law School. <https://law.stanford.edu/wp-content/uploads/2018/07/Mistree-Dibley-Corruption.4.18.18.pdf>
- Mummolo, J., & Peterson, E. (2019). Demand effects in survey experiments: An empirical assessment. *American Political Science Review*, 113(2), 517–529. <https://doi.org/10.1017/S0003055418000837>
- Mutz, D. C. (2011). *Population-based survey experiments* (STU - Student edition). Princeton University Press. Retrieved February 24, 2026, from <http://www.jstor.org/stable/j.ctt7sf3s>
- Mutz, D. C. (2021). *Winners and losers: The psychology of foreign trade* (Vol. 27). Princeton University Press. <https://doi.org/10.2307/j.ctv1fj85hw>
- Nyhan, B., & Reifler, J. (2010). When corrections fail: The persistence of political misperceptions. *Political Behavior*, 32(2), 303–330. <https://doi.org/10.1007/s11109-010-9112-2>
- OECD. (2016). *Supreme audit institutions and good governance: Oversight, insight and foresight* (OECD Public Governance Reviews). OECD Publishing. Paris. <https://doi.org/10.1787/9789264263871-en>
- O’Keefe, D. J. (2021). Persuasive message pretesting using non-behavioral outcomes: Differences in attitudinal and intention effects as diagnostic of differences in behavioral effects. *Journal of Communication*, 71(4), 623–645. <https://doi.org/10.1093/joc/jqab017>
- Olden, A., & Møen, J. (2022). The triple difference estimator. *The Econometrics Journal*, 25(3), 531–553.
- Pan, Y., Shu, Z., & Ye, Z. (2023). Exploring the dynamics of corruption perceptions in sustained anti-corruption campaigns: A survey experiment in China. *Humanities and Social Sciences Communications*, 10(1), 570.

- Panish, A. R. (2025). Informed or overwhelmed? disentangling the effects of cognitive ability and information on public opinion. *British Journal of Political Science*, 55, e58. <https://doi.org/10.1017/S0007123425000055>
- Peiffer, C. (2020). Message received? experimental findings on how messages about corruption shape perceptions. *British Journal of Political Science*, 50(3), 1207–1215. <https://doi.org/10.1017/S0007123418000108>
- Peterson, E., & Iyengar, S. (2021). Partisan gaps in political information and information-seeking behavior: Motivated reasoning or cheerleading? *American Journal of Political Science*, 65(1), 133–147.
- Peyton, K. (2020). Does trust in government increase support for redistribution? Evidence from randomized survey experiments. *American Political Science Review*, 114(2), 596–602. <https://doi.org/10.1017/S0003055420000076>
- Prior, M., Sood, G., & Khanna, K. (2015). You cannot be serious: The impact of accuracy incentives on partisan bias in reports of economic perceptions. *Quarterly Journal of Political Science*, 10(4), 489–518.
- Rathje, S., Roozenbeek, J., Bavel, J. J. V., & van der Linden, S. (2023). Accuracy and social motivations shape judgements of (mis)information. *Nature Human Behaviour*, 7(6), 892–903. <https://doi.org/10.1038/s41562-023-01540-w>
- Reinikka, R., & Svensson, J. (2006). Using micro-surveys to measure and explain corruption [Part Special Issue (pp. 324–404). Corruption and Development: Analysis and Measurement]. *World Development*, 34(2), 359–370. <https://doi.org/https://doi.org/10.1016/j.worlddev.2005.03.009>
- Rodríguez, I., Rodon, T., Unan, A., Herbig, L., Klüver, H., & Kuhn, T. (2025). Benchmarking pandemic response: How the UK’s COVID-19 vaccine rollout impacted diffuse and specific support for the EU. *British Journal of Political Science*, 55, e35. <https://doi.org/10.1017/S0007123424000802>

- Rodríguez Chatruc, M., Stein, E., & Vlaicu, R. (2021). How issue framing shapes trade attitudes: Evidence from a multi-country survey experiment. *Journal of International Economics*, 129, 103428. <https://doi.org/10.1016/j.jinteco.2021.103428>
- Roth, C., Settele, S., & Wohlfart, J. (2022). Beliefs about public debt and the demand for government spending [Annals Issue: Subjective Expectations & Probabilities in Economics]. *Journal of Econometrics*, 231(1), 165–187. <https://doi.org/10.1016/j.jeconom.2020.09.011>
- Roth, C., & Wohlfart, J. (2020). How do expectations about the macroeconomy affect personal expectations and behavior? *Review of Economics and Statistics*, 102(4), 731–748. https://doi.org/10.1162/rest_a_00867
- Sanna, G. A., & Lagnado, D. (2025). Belief updating in the face of misinformation: The role of source reliability. *Cognition*, 258, 106090. <https://doi.org/10.1016/j.cognition.2025.106090>
- Semenova, V. (2025). Generalized Lee bounds. *Journal of Econometrics*, 251, 106055.
- Solaz, H., Vries, C. E. D., & de Geus, R. A. (2019). In-group loyalty and the punishment of corruption. *Comparative Political Studies*, 52(6), 896–926.
- Spatharioti, S. E., Goldstein, D. G., & Hofman, J. M. (2024). Using open data to automatically generate localized analogies. *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. <https://doi.org/10.1145/3613904.3642638>
- Stagnaro, M. N., Tappin, B. M., & Rand, D. G. (2023). No association between numerical ability and politically motivated reasoning in a large US probability sample. *Proceedings of the National Academy of Sciences*, 120(32), e2301491120.
- Stuart, E. A., Lee, B. K., & Leacy, F. P. (2013). Prognostic score-based balance measures can be a useful diagnostic for propensity score methods in comparative effectiveness research. *Journal of Clinical Epidemiology*, 66(8), S84–S90.e1.
- Taber, C. S., & Lodge, M. (2006). Motivated skepticism in the evaluation of political beliefs. *American Journal of Political Science*, 50(3), 755–769. [https://doi.org/https://doi.org/10.1111/j.1540-5907.2006.00214.x](https://doi.org/10.1111/j.1540-5907.2006.00214.x)

- Tappin, B. M., Pennycook, G., & Rand, D. G. (2020). Bayesian or biased? analytic thinking and political belief updating [Epub 2020 Jun 24]. *Cognition*, 204, 104375. <https://doi.org/10.1016/j.cognition.2020.104375>
- Tappin, B. M., Berinsky, A. J., & Rand, D. G. (2023). Partisans' receptivity to persuasive messaging is undiminished by countervailing party leader cues [Epub 2023 Mar 2. Erratum in: *Nature Human Behaviour* 8(7):1431 (2024). DOI: 10.1038/s41562-024-01945-1]. *Nature Human Behaviour*, 7(4), 568–582. <https://doi.org/10.1038/s41562-023-01551-7>
- Thaler, M. (2024). The fake news effect: Experimentally identifying motivated reasoning using trust in news. *American Economic Journal: Microeconomics*, 16(2), 1–38. <https://doi.org/10.1257/mic.20220146>
- Thomas, R. J., Masthoff, J., & Oren, N. (2019). Can i influence you? development of a scale to measure perceived persuasiveness and two studies showing the use of the scale. *Frontiers in Artificial Intelligence*, 2. <https://doi.org/10.3389/frai.2019.00024>
- Transparency International. (2023). *Corruption perceptions index 2022* (tech. rep.) (Accessed: 14 May 2025). Transparency International. https://images.transparencycdn.org/images/Report_CPI2022_English.pdf
- World Bank. (2021, July). Supreme audit institutions independence index: 2021 global synthesis report.
- Yu, S., & Opfer, J. E. (2024). Cognitive support for political partisans' understanding of policy data. *PLOS ONE*, 19(10), 1–23.
- Zappalà, G. (2024). Adapting to climate change accounting for individual beliefs. *Journal of Development Economics*, 169, 103289. <https://doi.org/https://doi.org/10.1016/j.jdeveco.2024.103289>
- Zhou, F., & Oostendorp, R. (2014). Measuring true sales and underreporting with matched firm-level survey and tax office data. *The Review of Economics and Statistics*, 96(3), 563–576. https://doi.org/10.1162/REST_a_00408

- Zhu, J., Lu, J., & Shi, T. (2013). When grapevine news meets mass media: Different information sources and popular perceptions of government corruption in mainland China. *Comparative Political Studies*, 46(8), 920–946.
- Zimmermann, F. (2020). The dynamics of motivated beliefs. *American Economic Review*, 110(2), 337–61. <https://doi.org/10.1257/aer.20180728>

Government Audits of Municipal Corruption and Belief

Updating: Experimental Evidence.

Raymond Duch

Felipe Torres-Raposo

Nuffield College, University of Oxford

University College London

September 23, 2026

Contents

A.1	Placebo Screen Shots	1
A.2	Subject Recruitment	1
A.3	The Comptroller General Office in Chile	5
A.4	Hypotheses	9
A.5	Beliefs Formation	11
A.6	Treatment Random Assignment and Attrition	14
A.7	Political Outcomes Analysis	28
A.8	Further Discussion of Political and Donations Outcomes	29
A.8.1	Incumbent Performance Across Public Goods Domains	30
A.8.2	Contributions to Municipal Public Goods	30
A.8.3	Incumbent Vote Support and Partisan Heterogeneity	31
A.8.4	Electoral Context and Competition	32
A.9	Compliance	38

A.10 Heterogeneity	41
A.11 Corruption Audits in Chile	44
A.12 Additional Analyses, Subgroup Analyses and Persistence Ratios	46
A.13 Power calculations	51
A.14 Ethics Approvals	53
Appendix References	53

A.1 Placebo Screen Shots

Figure A.1 shows some screenshots of the placebo treatment. The video provided three pieces of information: 1) The number of people that live in the commune, 2) the percentage of people above the legal age, and the size of the commune in square kilometers.

Figure A.1: Screen Shots - Placebo Treatment



Note: This figure shows example screenshots of the placebo treatment arm. Subjects in the placebo condition received three pieces of information about their municipality. The first screenshot (left) reported the population that resides in the municipality. The second screenshot (middle) reported the population that is above the legal age (18 years old). The third (right) and final piece of information is the size of the commune in terms of square kilometers.

A.2 Subject Recruitment

The initial phase of the experiment ran from late January 2020 to January 2021, during which we recruited 46,723 Chilean residents aged 18 or older as potential participants across the 116 audited municipalities. We recruited subjects using Facebook, Instagram, and Twitter ads. We complemented our online recruitment efforts with traditional subject recruitment methods, including Flyers, posters, and newspaper advertisements. Study participants completed a brief form that requested information about the municipality where they reside, their gender, and contact details. Figure A.2 shows an example of the ads posted in different social media platforms and posters and Flyers distributed across the 116 municipalities. In December 2019, the CGO finalized the list of 116 municipalities to be audited in 2020. Our recruitment efforts focused on these communes.

The recruitment process formed part of Nuffield’s strategy to expand and enhance its subject panel in Chile. The recruitment advertisements did not specify the studies the centre would conduct, but they did include information about the centre’s objectives and the compensation offered

to respondents participating in Nuffield CESS's surveys.

Figure A.2: Flyers and Social Media Ads



Note: This figure shows the social media advertisements and flyers used in the recruitment process to invite individuals to join the Nuffield CESS recruitment panel. The materials provided information on how to sign up, the nature of the surveys in which participants would be invited to take part, and the range of compensation they would receive for their participation.

We recruited our sample from Chileans with a social media presence, focusing primarily on Facebook. Recruitment occurred in two phases. In the first phase, we identified 46,723 individuals as eligible for the experiment. We invited 22,755 eligible subjects to participate in the pre-treatment survey. The invitations did not specify that the survey concerned malfeasance; instead, they emphasized that respondents would be asked about issues within their local government.

We achieved a response rate of 24.29%, resulting in 5,528 subjects participating in the pre-treatment survey. Table A.1 compares the demographic features of the final sample with statistics from the Chilean national census and the National Socio-Economic Characterization Survey. Note

that the pool and the sample are very similar in terms of age, education, gender, and income. There are differences between our sample and Chile’s demographic population. Our sample is disproportionately younger and has higher educational attainment. In terms of gender, the study had a higher proportion of women, representing 65% of our sample. Our sample reports a slightly higher average income than the population’s average.

Table A.1: Profile of Final Sample

	Census	CASEN	Sample
N	17,574,003	2,164,439	5,528
Age	40	37	36
Education			
Secondary	40%	37%	67%
Tertiary	24%	22%	31%
Female	51%	52%	65%
Income (USD)	-	607	658


Note: In this table, we compare our sample to Chile’s population, examining key demographics such as age, income, educational attainment, and gender. The first column provides population estimates of these covariates using data from the 2017 census. The second column includes population estimates based on the 2020 National Socio-Economic Characterization Survey. Age is the age of respondents in 2020. Education is a categorical variable with levels: “Primary”, “Secondary”, and “Tertiary”. We only report “Secondary” and “Tertiary.” “Primary” is the reference category. Income is the respondents’ monthly income reported in USD dollars (Exchange rate 771 Chilean pesos equal to 1 USD).

Table A.2: Turnout by Group

	Negative	Positive	Audited	No Audit	All Country	Sample
Average turnout (%)	30.97 [30.96, 30.98]	34.93 [34.89, 34.97]	31.50 [31.48, 31.51]	41.73 [41.71, 41.75]	34.91 [34.89, 34.92]	64.67 [62.18, 67.17]
Right	37.806 [37.705, 37.906]	50.329 [50.12, 50.539]	39.749 [39.658, 39.84]	36.309 [36.202, 36.417]	38.266 [38.196, 38.335]	13.423 [10.97, 15.875]
Center-left	34.579 [34.476, 34.682]	24.78 [24.522, 25.038]	33.059 [32.963, 33.155]	39.275 [39.17, 39.379]	35.739 [35.669, 35.81]	10.999 [8.512, 13.486]
Left	7.706 [7.584, 7.829]	3.705 [3.414, 3.997]	7.086 [6.973, 7.198]	3.591 [3.459, 3.723]	5.579 [5.493, 5.664]	34.75 [32.621, 36.88]
No party affiliation	15.667 [15.55, 15.783]	18.202 [17.933, 18.471]	16.06 [15.953, 16.167]	17.638 [17.516, 17.76]	16.74 [16.66, 16.821]	38.278 [36.207, 40.349]
Other party	4.242 [4.117, 4.367]	2.984 [2.691, 3.276]	4.047 [3.932, 4.161]	3.187 [3.055, 3.319]	3.676 [3.589, 3.763]	2.551 [0, 5.153]
Female (%)	55.037 [54.954, 55.12]	55.65 [55.452, 55.848]	55.127 [55.051, 55.203]	54.156 [54.062, 54.251]	54.741 [54.681, 54.8]	64.734 [62.795, 66.672]
18-29	15.231 [15.118, 15.345]	15.385 [15.112, 15.658]	15.254 [15.149, 15.359]	15.219 [15.09, 15.348]	15.24 [15.159, 15.322]	43.326 [40.865, 45.787]
30-49	32.262 [32.16, 32.363]	33.087 [32.844, 33.331]	32.383 [32.29, 32.477]	33.901 [33.787, 34.014]	32.987 [32.915, 33.059]	40.684 [38.167, 43.201]
50-69	39.244 [39.148, 39.341]	38.013 [37.779, 38.247]	39.064 [38.975, 39.152]	38.334 [38.224, 38.444]	38.773 [38.704, 38.843]	15.573 [12.57, 18.576]
70-80+	13.262 [13.147, 13.377]	13.515 [13.238, 13.791]	13.299 [13.193, 13.405]	12.546 [12.415, 12.677]	12.999 [12.917, 13.082]	0.417 [0, 3.679]

Note: This table reports voter turnout (as percentage of eligible voters) and party preferences in the 2016 municipal election. *Source:* SERVEL (Servicio Electoral de Chile) electoral results. Estimates are reported in percentages and confidence intervals are reported in square brackets. CIs are calculated based on the normal approximation; lower bounds are truncated at zero where the interval would otherwise extend below zero. Results are disaggregated by benchmarking condition (Positive or Negative feedback), audit status (Audited vs. No Audit), and population level (full population average vs. study sample). Turnout for the sample is measured as the share of respondents who self-reported voting in the 2016 election, and who were eligible (above 18 years old) to vote in the 2016 election (numerator), divided by the total number of respondents who were eligible to vote in that election (denominator). Turnout is further broken down by demographic characteristics including age (at the time of the election). Party preferences and female percentages for the sample are measured only among respondents who self-reported voting in the 2016 election.

Figure A.3: Saliency of Corruption in Chile - AmericasBarometer Surveys




Note: This graph displays the proportion of respondents who identified the most pressing issue facing the country. The data come from the 2017, 2019, and 2021 waves of the AmericasBarometer, conducted by the Latin American Public Opinion Project (LAPOP, 2021). No survey was conducted in 2020. Respondents were asked the following question: *In your opinion, what is the most serious problem the country is currently facing?* Estimates are weighted using the sampling weights provided in the data to attain national representative estimates.

A.3 The Comptroller General Office in Chile

As noted in the manuscript, the Comptroller General's Office (CGO) is regarded as a trustworthy institution by the Chilean public. Although we do not have a direct measure of trust in the CGO in our surveys, we examine whether levels of trust differ by partisan affiliation using data from the National Transparency in Chile's annual representative survey. This survey asks respondents which institutions they trust to uncover corruption in Chile. Results indicate that a significant portion


of Chileans consider the CGO the fourth-most-trustworthy entity to tackle corruption. Analysis by partisanship reveals that levels of trust are broadly shared across political orientations, though left-leaning respondents report slightly higher levels of trust. Figure A.4 reports the proportion of respondents that trust the CGO as one of the main entities to tackle corruption in Chile. Figure A.5 reports levels of trust in the CGO broken down by partisanship and education. We observe no partisan differences in trust in the CGO, but we find that trust increases monotonically with education.

Figure A.4: Trust in the Comptroller General Office



Note: The top figure reports the proportion of respondents who identified each institution as among the three most trustworthy to fight corruption in Chile, out of a list of 12 institutions. Results are shown for two survey waves (2019 and 2020), distinguishing the CGO (Comptroller General Office) from other institutions.

Figure A.5: Trust in the Comptroller General Office by Partisanship



Note: This figure breaks down, by ideological group, the proportion of respondents who included the CGO among their top three choices. Ideology is measured on a 0 to 10 scale, where 0 represents left-leaning and 10 represents right leaning, grouped into Left (0–3), Center (4–6), and Right (7–10). Educational attainment is grouped into four categories: Primary or lower, Secondary, Technical, and Tertiary; each category incorporates both complete and incomplete levels of that educational tier. Points represent proportions with 95% confidence intervals; brackets display p-values for pairwise comparisons between groups.

A.4 Hypotheses

Table A.3: Pre-Registered and Exploratory Hypotheses

Hypothesis	Hypothesized		Findings			
	Dir.	Dir.	Qual.	Res.	Dist.	All
<i>Panel A: Pre-Registered Hypotheses</i>						
H1: Being informed of municipal government malfeasance increases beliefs about levels of corruption.	↑	↑	0.31	0.29	0.12	
H2: Beliefs regarding corruption levels will be higher if reported malfeasance has increased compared to the previous audit and will be lower if reported malfeasance has declined. (Negative/Positive)	↑ & ↓	↑ & ↑	0.32/0.23	0.32/0.12	0.43/0.08	
H4: Providing respondents with audit agency data about their municipality will reduce the gap between their subjective beliefs and objective corruption metrics.	$\underbrace{\Delta \downarrow}_{\text{True-Beliefs}}$	ϕ		-0.21	-0.34	
H7a: Reported malfeasance is negatively correlated with contributions to public goods.	↓	—				-1.7
H7b: Reported malfeasance is negatively correlated with vote for incumbent governing parties.	↓	↓				-4.5
<i>Panel B: Exploratory Hypotheses</i>						
H3: Belief updating will correlate positively with the reported magnitude of malfeasance levels in the respondent's municipality.	↑	—	0.07	0.01	0.39	
H5: Subjects with relatively high corruption priors will be unresponsive to negative malfeasance reports; subjects with relatively low priors will be unresponsive to positive reports. (High/Low).	$b_{High} \prec b_{Low}$	$b_{High} \preceq b_{Low}$	0.26/0.34	0.39/0.27	0.44/0.5	
H6a (<i>Motivated reasoning</i>): The effect of audit information on corruption beliefs is smaller for copartisans of the incumbent than for non-copartisans. (Partisan/Non-copartisan)	$\Delta b_P \prec \Delta b_{NP}$	$\Delta b_P \preceq \Delta b_{NP}$	0.30/0.30	0.33/0.31	0.39/0.46	

Continued on next page

Table A.3 continued

Hypothesis	Hypothesized		Findings			
	Dir.	Dir.	Qual.	Res.	Dist.	All
H6b (<i>Parallel persuasion</i>): Copartisans and non-copartisans differ in their priors but update by similar amounts, so the partisan gap in beliefs persists after treatment. (Partisan/Non-copartisan) (<i>Prior beliefs</i>)	Prior _P ≠ Prior _{NP}	Prior _P < Prior _{NP}	3.21/4.96	3.14/4.11	1.55/2.04	
H6b (<i>Parallel persuasion</i>): Copartisans and non-copartisans differ in their priors but update by similar amounts, so the partisan gap in beliefs persists after treatment. (Partisan/Non-copartisan) (Posterior beliefs)	Post _P ≠ Post _{NP}	Post _P < Post _{NP}	3.95/5.48	3.69/4.49	1.84/2.28	

Note: Treatment effects are standardized using the standard deviation of the control group. Panel A presents pre-registered hypotheses, while Panel B contains exploratory analyses. Dir. = direction; ↑/↓ = predicted increase or decrease; Δ = change; < = strictly less than; ≤ = less than or equal to; ≠ = not equal; — = not significant; ϕ = inconclusive. Bold indicates significance at the 5% level. Effect sizes for H1 correspond to estimates in Table 2. Effect sizes for H2 are derived from regressions of treatment on outcome beliefs, with data subset by negative and positive outcomes. Effect sizes for H3 and H4 are obtained from Table A.21. The *All* column reports pooled behavioral estimates in percentage points, corresponding to column 4 of Table 4. Effect sizes are calculated using regression coefficients from the *Magnitude* model in Table 2. Effect sizes for H5 are calculated using regression coefficients from Table A.15. Estimates in H6a (Prior beliefs) and H6b (Posterior beliefs) represent average beliefs. Prior belief comparisons use the full sample; posterior belief comparisons are restricted to treated respondents (Negative and Positive conditions), excluding the Control group. Bold indicates a statistically significant difference between the two groups, based on a Welch's two-sample *t*-test comparing group means.

A.5 Beliefs Formation

In this section, we examine the correlation between prior beliefs and respondents' media consumption behavior. Drawing on our baseline survey, we assess mainstream media, interpersonal networks, and social media usage, capturing both the frequency and content of media exposure. Table A.4 presents these correlations, with particular attention to the channels through which respondents acquire political information and corruption beliefs. Results from ordinary least squares regressions indicate that more frequent consumption of mainstream media is associated with lower, less pessimistic corruption beliefs across all three belief dimensions. In contrast, more frequent information-seeking from friends and family is associated with higher, more pessimistic beliefs, also consistent across all three dimensions. Social media consumption shows a similar positive association, but it is only statistically significant for the qualitative belief measure.

Table A.4: Correlation between priors and political information consumption


	Qualitative			Resources			Distribution		
	Friends & Family	Mainstream	Social Media	Friends & Family	Mainstream	Social Media	Friends & Family	Mainstream	Social Media
Media Consumption	0.112*** (0.024)	-0.078*** (0.022)	0.078** (0.025)	0.055* (0.026)	-0.096*** (0.023)	-0.026 (0.027)	0.034*** (0.008)	-0.037*** (0.008)	0.013 (0.009)
Covariates	✓	✓	✓	✓	✓	✓	✓	✓	✓
Num. obs.	5265	5273	5279	5265	5273	5279	5265	5273	5279
R2	0.030	0.029	0.028	0.016	0.019	0.016	0.026	0.027	0.023
R2 Adj.	0.026	0.025	0.024	0.012	0.014	0.011	0.022	0.023	0.019
RMSE	2.40	2.40	2.40	2.36	2.35	2.36	0.77	0.77	0.77

Note: This table reports the results of an ordinary least squares regression examining the relationship between political information consumption and corruption beliefs. The main outcome here is media consumption. Respondents were asked how frequently they gathered and discussed politics (From Never to Daily), and from which sources: (1) friends and family, (2) social media platforms, and (3) traditional media, including TV news, radio, and newspapers. Ordered factors were converted to a numeric scale (1 to 5) for use as a predictor in the regression models reported in this table.

Because WhatsApp served as the primary channel for intervention delivery, we conduct a more granular examination of respondents' information consumption behavior on this platform. Figure A.6 illustrates the types of content respondents consumed via WhatsApp. The bottom panel presents the average prior beliefs based on the frequency they receive political information on this platform. Although there are no significant differences between beliefs and the receipt of political information, beliefs tend to become more optimistic after individuals sometimes receive political information. However, these beliefs subsequently fluctuate across frequency categories, without

exhibiting a consistent upward or downward trend.

Figure A.6: Correlation between Consumption of Political Information via WhatsApp and Priors




Note: The top panel presents the proportion of respondents who received each type of information via WhatsApp. The bottom panel reports average beliefs (with 95% confidence intervals), broken down by whether respondents indicated the frequency of receiving political information. Red lines represent the average beliefs for the whole sample.

Audit Frequency, Audit Outcomes, and Corruption Beliefs As outlined in the paper, we examined whether priors are correlated with anti-corruption efforts. Figures A.7 and A.8 present correlations among audit frequency, the average number of irregularities identified, and the sever-


ity of these irregularities. Across these different metrics of malfeasance, we find no correlation.

Figure A.7: Correlation between Number of Audits and Prior Beliefs



Note: This figure displays the relationship between the total number of audits and special investigations conducted in a municipality (2015–2020) and average prior corruption beliefs, separately for each belief dimension (Qualitative, Resources, Distribution). Each point represents a municipality, sized by survey sample size (n). The fitted line is estimated using ordinary least squares weighted by municipality sample size. The reported slope (b), its standard error, and the associated p-value correspond to this weighted regression.

Figure A.8: Correlation between Type of Irregularities and Prior Beliefs




Note: This figure displays the relationship between average irregularities per audit and average prior corruption beliefs, separately for each belief dimension (Qualitative, Resources, Distribution) and each irregularity severity type (Low, Moderate, Medium, High). Each point represents a municipality, sized by the number of audits conducted. The fitted line in each panel is estimated using ordinary least squares weighted by the number of audits. The reported slope (b), its standard error, and the associated p -value correspond to this weighted regression.

A.6 Treatment Random Assignment and Attrition

Balance Respondents are observed at three different time points during the trial: baseline, when treatment is assigned and implemented; one week post-treatment; and one month post-treatment. Balance on covariates is assessed in Figure A.9 by comparing (in red) standardized mean dif-

ferences (raw differences in proportion for binary variables) for the treatment and placebo arms. The estimates are generated by the R package *cobalt* (Greifer, 2022). We also compare these unadjusted differences with those obtained when the sample is weighted using propensity score matching (in green). These comparisons are generated for the sample of 6,050 participants interviewed at baseline, the 3,592 participants interviewed one-week post-treatment, and the 2,255 participants contacted one-month post-treatment. We employ indicative balance tolerance levels of 0.1 (Stuart et al., 2013) in Figure A.9 (the vertical dotted lines). With a few exceptions among the age categories and the partisanship scale, the unadjusted standardized mean differences fall within this threshold; the adjusted differences all do.

Figure A.9: Balance on Unadjusted and Adjusted Standardized Mean Differences



Note: This figure shows unadjusted (blue) and adjusted (red) standardized mean differences for key covariates. Adjusted values are estimated using nearest neighbor matching based on propensity scores from a logistic regression predicting treatment. Covariates include age, gender, income, partisanship, education, civic, and political knowledge

Table A.5 reports the results of regressing subjects' treatment assignment (treatment versus placebo) on observed covariates, including both nested and full models. The initially assigned

sample of subjects is balanced – only one of the covariate predictor variables, partisanship, is weakly significant. Table A.5 also indicates that with rising attrition and declining compliance, the resulting samples remain balanced across the two treatment groups in most covariates. We find differences between the two groups for *Civic knowledge* and *Partisanship* covariates in the follow-up survey. This result indicates that respondents who received the treatment had lower levels of civic knowledge compared to the placebo group and lower partisanship.

Table A.5: Randomization Checks - Pre-treatment Variables

	Baseline	Post-Treatment	Follow-up
Intercept	0.661*** (0.022)	0.692*** (0.035)	0.628*** (0.054)
Gender	0.007 (0.013)	0.005 (0.017)	-0.018 (0.022)
Income	-0.005 (0.014)	-0.009 (0.018)	-0.009 (0.031)
Education	0.003 (0.015)	-0.007 (0.022)	-0.010 (0.029)
Partisanship	-0.006** (0.003)	-0.010** (0.004)	-0.015** (0.006)
Turnout	-0.014 (0.013)	-0.016 (0.019)	-0.001 (0.023)
Support incumbent	-0.011 (0.014)	-0.011 (0.021)	-0.010 (0.036)
Civic knowledge: Medium	0.038 (0.058)	-0.040 (0.073)	-0.160* (0.088)
Civic knowledge: High	0.027 (0.058)	-0.066 (0.075)	-0.184* (0.090)
Political information: Medium	0.003 (0.013)	0.009 (0.015)	0.019 (0.021)
Political information: High	0.016 (0.016)	0.032 (0.023)	0.046 (0.032)
Anova F-test Nested (P-value)	0.353	0.07	0.045*
Num.Obs.	6050	3592	2255
R2	0.002	0.004	0.008
R2 Adj.	0.000	0.001	0.003

Note: This table reports regression estimates of treatment assignment on pre-treatment variables. Treatment was assigned to 6,050 respondents, but 522 were excluded due to incomplete audits. Standard errors are clustered by municipality. The table also shows p-values from F-tests comparing models with and without covariates. * $p < .05$, ** $p < .01$, *** $p < .001$

Attrition As we point out in the main text, we observe attrition in both treatment and placebo conditions. In Table A.6, we report pre-treatment measures of respondents’ outcomes and covariates for the whole, observed, and attriters samples across all waves. This analysis shows near-zero differences for all three outcomes between attriters and the observed sample in post-treatment and follow-up waves. Similarly, we examined individual-level covariates, including gender, education, income, and copartisanship. We did not observe statistically significant differences across all covariates. There is evidence that higher-income subjects were more likely to drop out in the post-treatment and follow-up surveys.

We further examined differential attrition by treatment status. Table A.8 compares pre-treatment outcomes and covariates across treatment and placebo groups, distinguishing between attriters and observed respondents. We find no systematic baseline differences, except for the Qualitative outcome in the treatment group, where attriters reported higher prior corruption beliefs. Pre-treatment beliefs for the two quantitative outcomes remain balanced across all groups.

We find no differences in behavioral outcomes between attriters and non-attriters within treatment arms across both waves. At the individual level, groups are similar in terms of gender, education, copartisanship, and income. However, among treated respondents, attriters in the post-treatment survey tend to be older and more right-leaning. As noted, there is some evidence that attrition rates are higher among treated respondents living in high-corruption communities.

We assessed whether key covariates and treatment assignment were correlated with differential attrition. Table A.7 presents logistic regression results, where the outcome is a binary indicator for post-treatment survey dropout. Attrition is positively associated with treatment assignment and baseline donations but not with other covariates such as copartisanship, age, education, or political knowledge. Table A.9 reports complementary tests by block design.

A potential threat to internal validity is that attrition may be correlated with treatment assignment. To mitigate this, we block-randomized respondents based on prior beliefs about municipal malfeasance, forming “Low” and “High” blocks within each commune. Table A.9 shows no evidence of differential attrition across these blocks. We also examined whether attrition was as-

sociated with block size, which varied due to our randomization design. Results in columns 2 and 3 show that neither block type nor size predicts attrition. Finally, we tested whether attrition was related to recruitment intensity, measured by variation in sample sizes across the 116 communes. Column 4 shows no association, suggesting that attrition is not systematically related to recruitment procedures.

Table A.6: Pre-Treatment Descriptive Analysis for Malfeasance Belief Outcome Measures and Covariates for observed Samples and Attriters

	Baseline	Post-Treat			Follow-up		
	Whole sample	Observed sample	Attriters	Difference	Observed sample	Attriters	Difference
Qualitative Measures							
Subjective malfeasance scale 1-10	4.49 (2.44)	4.45 (2.40)	4.55 (2.51)	0.10	4.51 (2.40)	4.47 (2.47)	-0.05
Certainty Subjective malfeasance scale 1-10	5.40 (2.71)	5.38 (2.70)	5.45 (2.71)	0.07	5.40 (2.70)	5.41 (2.71)	0.00
Quantitative Measures - Incentivized							
Resources (1-10)	3.84 (2.38)	3.85 (2.40)	3.84 (2.35)	-0.01	3.85 (2.38)	3.84 (2.38)	-0.01
Distribution (1-4)	1.90 (0.78)	1.89 (0.79)	1.93 (0.78)	0.04	1.90 (0.78)	1.91 (0.79)	0.01
Behavioral measures							
500 CLP Municipal Donation	32% (0.63%)	31% (1.08%)	34% (0.77%)	3%	31% (0.83%)	33% (0.01%)	3%
Covariates							
Female	64.73 (0.64)	64.67 (1.09)	64.82 (0.80)	0.15	65.76 (1.00)	64.01 (0.84)	-1.75
Education (Yrs)	15.07 (0.04)	15.06 (0.07)	15.07 (0.05)	0.01	15.05 (0.05)	15.08 (0.06)	0.03
Age	35.88 (0.18)	35.76 (0.29)	36.11 (0.22)	0.35	35.77 (0.23)	35.96 (0.28)	0.19
Partisanship (1-10)	4.47 (0.03)	4.46 (0.05)	4.49 (0.03)	0.03	4.43 (0.04)	4.49 (0.04)	0.06
Support for the incumbent	27.13% (0.59%)	27.00% (1.01%)	27.38% (0.74%)	0.38%	26.87% (0.78%)	27.31% (0.93%)	0.44%
Income (USD)	658 (70)	896 (77)	1012 (140)	116	775 (18)	1046 (117)	271
Municipal-level covariates							
Share previous malfeasance	0.04 (0.14)	0.03 (0.13)	0.04 (0.15)	0.01	0.03 (0.11)	0.04 (0.16)	0.01
Historical irregularities	43.03 (16.51)	42.47 (16.64)	44.08 (17.31)	1.61	42.56 (16.51)	43.35 (17.15)	0.79
Sample	5528	3592	1936		2255	3273	

Note: This table reports pre-treatment values for all outcomes and relevant covariates broken down by attrition status. Along with mean values, we also calculate the difference in means for each variable. Standard deviations are reported in parenthesis.

Table A.7: Attrition Covariates

	(1)	(2)	(3)	(4)	(5)
(Intercept)	-1.010*** (0.053)	-1.329*** (0.349)	-1.994*** (0.501)	-1.784** (0.572)	-0.956 (0.844)
Treat	0.569*** (0.063)	0.552*** (0.064)	0.593*** (0.067)	0.593*** (0.067)	-0.755 (0.955)
Age (Years)		-0.002 (0.003)	-0.003 (0.003)	-0.002 (0.003)	-0.004 (0.005)
Log Income		0.066** (0.022)	0.072** (0.023)	0.072** (0.023)	0.109* (0.044)
Partisanship		0.005 (0.015)	-0.007 (0.016)	-0.005 (0.016)	-0.030 (0.028)
Education:Medium		-0.448 (0.258)	-0.396 (0.272)	-0.388 (0.272)	-0.810 (0.446)
Education:High		-0.421 (0.253)	-0.355 (0.267)	-0.347 (0.267)	-0.965* (0.437)
Copartisan		0.038 (0.065)	0.015 (0.073)	0.011 (0.073)	-0.001 (0.128)
Civic knowledge:Medium				-0.299 (0.272)	-0.912* (0.428)
Civic knowledge:High				-0.372 (0.273)	-1.053* (0.428)
Political knowledge:Medium				0.096 (0.075)	0.256 (0.142)
Political knowledge:High				0.112 (0.089)	0.327 (0.168)
Treat:Age(Years)					0.003 (0.006)
Treat:Log Income					-0.051 (0.052)
Treat:Partisanship					0.037 (0.034)
Treat:Education:Medium					0.671 (0.562)
Treat:Education:High					0.944 (0.552)
Treat:Partisan					0.022 (0.150)

Continued on next page

Table A.7: Attrition Covariates (Continued)

	(1)	(2)	(3)	(4)	(5)
Treat:Civic knowledge:Medium					0.968 (0.554)
Treat:Civic knowledge:High					1.060 (0.555)
Treat:Political Knowledge:Medium					-0.221 (0.168)
Treat:Political Knowledge:High					-0.301 (0.198)
Commune FE	×	×	✓	✓	✓
Num.Obs.	5528	5310	5310	5310	5310
AIC	7078.1	6816.8	6710.3	6713.7	6720.0
BIC	7091.3	6869.4	7519.3	7549.0	7621.1
McFadden's R squared	0.012	0.012	0.061	0.062	0.064

Note: This table reports logistic regressions where the dependent variable equals 1 if the respondent attrited in the post-treatment survey and 0 otherwise. Column 1 includes only *Treat*. Column 2 adds covariates (age, copartisanship, education). Column 3 adds commune fixed effects. Column 4 adds civic and political knowledge. Column 5 is a fully saturated model interacting *Treat* with all covariates. $*p < .05$, $**p < .01$, $***p < 0.001$.

Table A.8: Descriptive for Malfeasance Belief Outcome Measures and Covariates for Samples and Attriters by Treatment Status

	Post-Treat													
	Baseline		Observed Sample				Differences		Follow-up				Differences	
	Placebo	Treatment	Placebo	Treatment	Placebo	Treatment	Placebo	Treat	Placebo	Treat	Placebo	Treatment	Placebo	Treat
Qualitative Measures														
Subjective malfeasance scale (1-10)	4.47 (2.44)	4.49 (2.44)	4.50 (2.39)	4.42 (2.41)	4.42 (2.56)	4.60 (2.49)	0.08	-0.18	4.53 (2.43)	4.50 (2.37)	4.40 (2.45)	4.49 (2.48)	0.13	0.01
Certainty of Subjective malfeasance scale (1-10)	5.44 (2.72)	5.39 (2.70)	5.45 (2.73)	5.34 (2.69)	5.41 (2.70)	5.46 (2.72)	0.05	-0.12	5.52 (2.73)	5.30 (2.66)	5.33 (2.70)	5.43 (2.72)	0.19	-0.13
Quantitative Measures - Incentivized														
Resources (1-10)	3.83 (2.36)	3.85 (2.39)	3.84 (2.35)	3.85 (2.42)	3.83 (2.39)	3.85 (2.33)	0.01	0.01	3.82 (2.33)	3.88 (2.42)	3.86 (2.40)	3.83 (2.37)	-0.04	0.05
Distribution (1-4)	1.90 (0.78)	1.91 (0.79)	1.90 (0.78)	1.89 (0.79)	1.91 (0.79)	1.94 (0.78)	-0.01	-0.05	1.91 (0.78)	1.90 (0.77)	1.90 (0.78)	1.91 (0.79)	0.01	-0.01
Behavioral measures														
500 CLP Municipal Donation	32% (1.09%)	33% (0.77%)	32% (2.12%)	31% (1.26%)	33% (1.27%)	35% (0.98%)	-1%	-4%	31% (1.69%)	31% (0.94%)	33% (1.42%)	33% (1.34%)	-2%	-2%
Covariates														
Female	64.79 (1.12%)	64.69 (0.79%)	64.78 (2.16%)	64.61 (1.26%)	64.83 (1.30%)	64.82 (1.01%)	-0.05	-0.22	64.81 (1.72%)	66.61 (0.96%)	64.77 (1.47%)	63.77 (1.36%)	0.03	2.84
Education (Yrs)	15.09 (0.07)	15.05 (0.05)	15.10 (0.13)	15.04 (0.08)	15.06 (0.08)	15.07 (0.06)	0.04	-0.03	15.06 (0.1)	15.03 (0.06)	15.13 (0.08)	15.06 (0.08)	-0.06	-0.03
Age	36.42 (0.07)	35.61 (0.05)	36.41 (0.57)	35.37 (0.34)	36.46 (0.37)	35.99 (0.28)	-0.05	-0.62	36.45 (0.46)	35.17 (0.26)	36.38 (0.43)	35.83 (0.38)	0.07	-0.66
Partisanship	4.55 (0.05)	4.43 (0.03)	4.58 (0.09)	4.38 (0.05)	4.47 (0.06)	4.50 (0.04)	0.11	-0.12	4.58 (0.07)	4.31 (0.04)	4.51 (0.06)	4.49 (0.06)	0.07	-0.18
Copartisanship	27.78% (1.05%)	26.81% (0.73%)	27.92% (0.09%)	26.46% (0.05%)	27.40% (0.06%)	27.37% (0.04%)	0.52%	-0.91%	27.63% (0.07%)	26.21% (0.04%)	28.00% (0.06%)	27.10% (0.06%)	-0.37%	-0.89%
Income (USD)	1003 (149)	904 (74)	853 (544)	921 (25)	1404 (28)	878 (121)	551	-43	793 (346)	973 (108)	1282 (26)	973 (24)	489	214
Municipal-level covariates														
Share previous malfeasance	0.04 (0.15)	0.03 (0.13)	0.04 (0.14)	0.03 (0.12)	0.05 (0.19)	0.04 (0.14)	0.01	0.01	0.04 (0.14)	0.02 (0.04)	0.04 (0.16)	0.04 (0.15)	0.00	0.02
Historical irregularities	42.53 (17.1)	43.28 (16.8)	42.27 (17.04)	42.58 (16.41)	43.23 (17.2)	44.37 (17.34)	0.26	1.79	42.14 (16.82)	42.92 (16.23)	42.92 (17.43)	43.45 (17.06)	0.78	0.53
Sample	1832	3696	1343	2249	489	2249			1057	1198	775	1198		

Note: This table reports pre-treatment values for all outcomes and relevant covariates broken down by attrition status and treatment status. Along with mean values, we report standard deviations in parentheses. We also calculate the difference in means for each variable, grouped by their attrition and treatment status.

Table A.9: Attrition Test - Block type and Block Size and Commune Sample

	Post-treatment				Follow-up			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
(Intercept)	0.345*** (0.013)	0.386*** (0.025)	0.383*** (0.036)	0.398*** (0.039)	0.591*** (0.025)	0.615*** (0.037)	0.628*** (0.057)	0.626*** (0.060)
Block - High	0.019 (0.017)		0.038 (0.042)		0.002 (0.024)		-0.034 (0.071)	
Block size		-0.001 (0.001)	-0.001 (0.001)			0.000 (0.001)	-0.001 (0.001)	
Block size:Block-High			-0.002 (0.001)				0.001 (0.003)	
Sample commune				-0.001 (0.001)				0.000 (0.001)
Num.Obs.	5528	5528	5528	5528	5528	5528	5528	5528
R2	0.000	0.002	0.003	0.002	0.000	0.001	0.001	0.001
R2 Adj.	0.000	0.002	0.002	0.002	0.000	0.001	0.000	0.001

Note: This table reports the results of three models using OLS. The dependent variable is a binary variable, equal to 1 if the respondent is an attriter and 0 if not, in the post-treatment survey. The independent variable is the block randomization variable used to conduct complete randomization. Each commune had two blocks: *Low* and *High*, where high corresponded to the block that grouped respondents with low prior beliefs about malfeasance. At the same time, the high block grouped respondents with high prior beliefs. We also evaluated whether the block size predicts differential attrition. The independent variable in this model is the sample within each commune and each block. Finally, in model *Both*, we interacted with the block type (High or Low) and block size variables. Standard errors are clustered at the commune level. $*p < .05$, $**p < .01$, $***p < 0.001$.

Inverse Probability Weighting With the complete set of demographic measures collected in the pre-treatment survey we can model attrition quite robustly. Inverse probability weighting is one of the estimation strategies we can implement. It has the value of being very straightforward – essentially modeling the attrition process as a function of observable covariates Anderson et al. (2021). The weights are based on the predicted values from a logistic regression of a binary variable indicating whether the observation is missing on the available covariates.¹ The weight is simply 1 over 1 minus these predicted probabilities.² We re-estimate treatment effects on the subset of the data where outcomes are observed and weight that estimate using these weights. The estimate from this regression is a consistent estimate for the treatment effect, assuming the censoring process is observable. Table A.10 compares the unweighted baseline treatment model coefficients with those that incorporate IPW for the three outcome variables. The coefficients for

¹An alternative here that we would consider using is a propensity score estimation algorithm to fully model any possible nonlinearities – the *twang* package, for example (Ridgeway et al., 2021).

²These weights are often characterized as “Weighted” – a slightly modified estimation strategy can generate more “stable” weights.

the weighted estimates are very similar to those for the unweighted estimates. In particular, the estimated treatment effects are similar in both weighted and unweighted estimations.

Table A.10: Stabilizing Weighting Model Regressions and Unweighted Models

Unweighted	Qualitative				Resources				Distribution			
	Information	Magnitudes	Negative	Positive	Information	Magnitudes	Negative	Positive	Information	Magnitudes	Negative	Positive
Intercept	2.734*** (0.401)	2.717*** (0.402)	2.826*** (0.407)	2.527+ (1.341)	2.302*** (0.662)	2.274*** (0.648)	2.333*** (0.686)	2.481 (1.969)	1.122*** (0.153)	1.125*** (0.152)	1.191*** (0.161)	0.861+ (0.459)
Prior	0.467*** (0.020)	0.464*** (0.020)	0.473*** (0.022)	0.436*** (0.055)	0.446*** (0.018)	0.444*** (0.018)	0.442*** (0.019)	0.458*** (0.054)	0.421*** (0.024)	0.415*** (0.024)	0.406*** (0.024)	0.453*** (0.093)
Treat	0.749*** (0.079)	0.736*** (0.091)	0.773*** (0.103)	0.566** (0.203)	0.740*** (0.088)	0.739*** (0.094)	0.834*** (0.103)	0.316 (0.200)	0.347*** (0.028)	0.327*** (0.035)	0.375*** (0.036)	0.073 (0.081)
Stand - Malfeasance		0.490 (0.489)	0.334 (0.476)	-1.877 (22.915)		0.590 (0.432)	0.559 (0.448)	2.175 (16.192)		0.217 (0.233)	0.220 (0.240)	-9.230+ (5.108)
Treat:Stand - Malfeasance		0.179 (0.882)	0.065 (0.904)	4.383 (14.757)		0.012 (0.597)	-0.259 (0.518)	11.969 (16.308)		0.302 (0.353)	0.158 (0.294)	8.216 (5.262)
Covariates	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Num.Obs.	3439	3439	2941	498	3439	3439	2941	498	3439	3439	2941	498
R2	0.263	0.264	0.270	0.268	0.220	0.221	0.217	0.304	0.203	0.207	0.204	0.244
R2 Adj.	0.257	0.258	0.262	0.226	0.214	0.214	0.209	0.264	0.196	0.200	0.196	0.201
Weighted												
Intercept	1.929*** (0.548)	1.916*** (0.553)	2.460*** (0.135)	2.280*** (0.229)	0.812*** (0.184)	2.011*** (0.099)	2.060*** (0.115)	1.726*** (0.227)	0.812*** (0.184)	1.132*** (0.053)	1.152*** (0.056)	1.085*** (0.159)
Treat	0.751*** (0.081)	0.743*** (0.093)	0.795*** (0.109)	0.620** (0.216)	0.343*** (0.028)	0.729*** (0.095)	0.812*** (0.104)	0.369+ (0.222)	0.343*** (0.028)	0.326*** (0.035)	0.372*** (0.037)	0.095 (0.082)
Stand - Malfeasance		0.532 (0.547)	0.287 (0.625)	8.205 (21.666)		0.525 (0.516)	0.425 (0.546)	13.680 (19.139)		0.183 (0.253)	0.169 (0.267)	-6.502 (5.090)
Treat:Stand - Malfeasance		0.116 (0.906)	-0.050 (0.962)	-3.552 (11.227)		-0.115 (0.646)	-0.357 (0.589)	2.506 (17.272)		0.317 (0.357)	0.174 (0.306)	6.608 (4.107)
Prior	0.465*** (0.020)	0.463*** (0.019)	0.480*** (0.021)	0.440*** (0.047)	0.417*** (0.024)	0.452*** (0.018)	0.448*** (0.019)	0.474*** (0.058)	0.417*** (0.024)	0.418*** (0.024)	0.411*** (0.025)	0.443*** (0.084)
Covariates	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Num.Obs.	3439	3439	2941	498	3439	3439	2941	498	3439	3439	2941	498
R2	0.262	0.263	0.260	0.216	0.199	0.203	0.198	0.241	0.199	0.190	0.189	0.197
R2 Adj.	0.256	0.256	0.259	0.209	0.192	0.202	0.197	0.235	0.192	0.189	0.188	0.191


Note: Regression estimates are reported using different specifications across all three outcomes. To reduce the influence of extreme weight values and improve the stability of our estimates, we stabilized the IPW weights by dividing each weight by the marginal probability of treatment. The variable *Stand-Malfeasance* is a min-max standardized measure of the amount of malfeasance reported in the video, ranging from 0 to 1. The *Negative* columns report the *Magnitude* specification for respondents who received negative information (increasing malfeasance). In contrast, the *Positive* columns report the specifications for those who received positive information (i.e., declining malfeasance). Standard errors are clustered at the commune level. * $p < .05$, ** $p < .01$, *** $p < .001$.

Figure A.10 reports confusion matrices comparing predicted and actual attritor status across four model specifications. Focusing on Model 4, the fullest specification, the model correctly classifies 91.9% of actual attriters but only 22.4% of actual non-attriters, yielding an overall accuracy of 67.4%, only modestly above the base attrition rate of 64.8% within Model 4's estimation sample.³ This modest classification performance is consistent with the attrition model being only weakly predictive, which is why the resulting IPW weights are themselves weakly informative and

³Accuracy is calculated as the share of respondents correctly classified (true positives plus true negatives) out of the total sample: $(3,159 + 420)/5,310 = 67.4\%$. Sensitivity is the share of actual attriters correctly classified as attriters: $3,159/(3,159 + 280) = 91.9\%$. Specificity is the share of actual non-attriters correctly classified as non-attriters: $420/(420 + 1,451) = 22.4\%$. The base attrition rate is the share of Model 4's estimation sample that attrited, regardless of model predictions: $3,439/5,310 = 64.8\%$.

the unweighted and weighted estimates coincide.⁴

Figure A.10: Confusion Matrix



Note: This figure displays the confusion matrix summarizing the predictive performance of Models 2 through 4 reported in Table A.7. The classification task distinguishes between attriters and non-attriters, where attrition is defined as respondents who did not complete the follow-up survey. Each cell reports the number of observations and their corresponding percentage classified into the actual and predicted attrition status. Correctly classified cases appear along the main diagonal, while misclassified cases appear in the off-diagonal cells.

⁴Sample sizes differ slightly across models due to listwise deletion of missing covariates: Model 1 (treatment only) retains the full sample of 5,528 respondents, while Model 4's fuller specification, which includes additional demographic, political covariates, and commune fixed effects, retains 5,310 respondents.

Lee, Generalized Lee Bounds Given the presence of differential attrition, we estimate Lee, *Tight* Lee, and Generalized Lee bounds. We compute standard Lee bounds without covariates, while we narrow *Tight* Lee bounds by including covariates selected through a LASSO procedure that identifies variables predicting attrition. Both *Income* and *Civic Knowledge* exhibit strong correlations with attrition status.

Table A.11: Lee and Tight Lee Bounds

	Standard Lee Bounds			Tight - Income			Tight - Civic Knowledge		
	Qualitative	Resources	Distribution	Qualitative	Resources	Distribution	Qualitative	Resources	Distribution
Lower	0.0446 (0.108)	0.261** (0.105)	0.160*** (0.0394)	0.0471 (0.100)	0.252** (0.102)	0.165*** (0.0385)	0.0536 (0.107)	0.259** (0.102)	0.165*** (0.0385)
Upper	1.431*** (0.115)	1.597*** (0.104)	0.591*** (0.0412)	1.466*** (0.104)	1.596*** (0.102)	0.611*** (0.0395)	1.478*** (0.103)	1.620*** (0.104)	0.612*** (0.0395)
Num.Obs.	5,528	5,528	5,528	5,528	5,528	5,528	5,528	5,528	5,528

Note: This table presents estimated Lee bounds and their corresponding confidence intervals. The bounds are tightened using *Income* and *Civic Knowledge*, which are the two covariates that most strongly predict attrition. * $p < .05$, ** $p < .01$, *** $p < .001$.

Table A.12 presents Generalized Lee bounds for all covariate specifications. The lower bounds are positive, and the confidence regions exclude zero in every column except the Qualitative outcome, indicating that the treatment effect remains positive and statistically significant even under worst-case attrition assumptions. The results are consistent across both the full covariate set and the LASSO-selected specification. The bounds are most precise for the *Distribution* outcome, indicating that this estimate is least sensitive to differential attrition.

Table A.12: Generalized Lee Bounds

	All Covariates			Selection (8)		
	Qualitative	Resources	Distribution	Qualitative	Resources	Distribution
Bounds	[0.338, 1.361]	[0.364, 1.566]	[0.341, 0.423]	[0.316, 1.308]	[0.321, 1.465]	[0.345, 0.408]
95% CR	(0.000, 1.618)	(0.001, 1.807)	(0.309, 0.469)	(0.040, 1.492)	(0.027, 1.651)	(0.330, 0.420)

Note: This table presents generalized Lee bounds in square brackets along with their 95% confidence region set in parenthesis. The estimates were obtained using all covariates. The same analysis was also performed using Lasso-selected covariates (eight covariates).

Table A.13: Summary Statistics Negative/Positive Respondents and Audit Activity

	Positive (N=866)		Negative (N=4668)		Diff. in Means	P-value
	Mean	Std. Dev.	Mean	Std. Dev.		
Prior beliefs						
Qualitative	4.339	2.574	4.512	2.416	0.173	0.067
Resources	3.789	2.547	3.854	2.347	0.066	0.481
Distribution	1.859	0.847	1.912	0.772	0.053	0.086
Covariates						
Gender (Female)	0.658	0.475	0.645	0.478	−0.013	0.460
Civic knowledge (High)	0.510	0.500	0.521	0.500	0.011	0.558
Political knowledge (High)	0.197	0.398	0.245	0.430	0.047	<0.05
Copartisan	0.323	0.468	0.262	0.440	−0.062	<0.05
Municipal benefits (Received)	0.203	0.403	0.247	0.431	0.044	0.004
Log Income	12.203	1.556	12.176	1.553	−0.027	0.646
Age (Yrs)	35.644	13.290	35.926	13.160	0.281	0.567
Education (Yrs)	15.009	2.848	15.077	2.826	0.067	0.524
Partisanship (1-10)	4.673	2.091	4.431	1.995	−0.242	<0.05

Note: The table presents summary statistics for outcomes and individual- and audit-level covariates, comparing respondents/municipalities for which malfeasance was increasing (Negative) with those for which malfeasance was reported as declining (Positive). Differences in means across outcomes and covariates are reported, along with their corresponding p-values.

Table A.14: Municipal Electoral Context by Benchmarking Condition (Negative vs. Positive)

	Negative (N=94)		Positive (N=22)		Diff. in Means	P-value
	Mean	Std. Dev.	Mean	Std. Dev.		
Municipal-level malfeasance						
Total Audits	3.990	2.328	3.318	2.033	−0.671	0.183
Number irregularities	17.598	16.457	17.812	15.798	0.215	0.854
Highly severe irregularities	6.229	4.124	5.864	3.533	−0.364	0.676
Political and voting variables						
Turnout	0.379	0.114	0.443	0.093	0.064	<0.05
Number of electoral coalitions	4.319	1.736	3.500	1.012	−0.819	<0.05
Minimum percentage of votes for an electoral coalitions	9.772	12.957	11.630	10.876	1.858	0.492
Maximum percentage of votes for an electoral coalitions	51.749	11.587	54.467	14.146	2.718	0.409
Total candidates	4.319	1.736	3.500	1.012	−0.819	<0.05

Continued on next page

Table A.14: Municipal Electoral Context by Benchmarking Condition (Negative vs. Positive)
(Continued)

Total independents	1.021	1.136	0.818	0.733	−0.203	0.303
% independent candidates	20.407	19.630	20.833	18.350	0.426	0.923
Effective Number of Parties (Votes)	2.615	0.774	2.441	0.669	−0.173	0.296
Conditional reelection rates	0.545	–	0.667	–	0.122	0.673
Unconditional reelection rates	0.128	–	0.182	–	0.054	0.5014
Ideological Polarization Index (Dalton, 2008)	2.775	0.816	2.531	0.577	−0.244	0.108
Ethnic Fractionalization (Alesina et al., 2003)	0.215	0.127	0.221	0.128	0.007	0.824
Ethnic Number	5.291	10.296	3.541	3.775	−1.750	0.194
Ethnic Presence	13.739	12.279	13.293	8.863	−0.446	0.845
Average Ideological (Bunker, 2025)	5.039	0.426	5.145	0.470	0.106	0.341
Disproportionality Index (Rae, 1967)	0.249	0.087	0.265	0.069	0.016	0.352
Disproportionality Index (Loosemore & Hanby)	0.483	0.115	0.455	0.142	−0.027	0.411
Disproportionality Index (Lijphart I, 1990)	0.483	0.115	0.455	0.142	−0.027	0.411
Disproportionality Index (Lijphart II, 1994)	0.374	0.049	0.370	0.050	−0.005	0.678
Disproportionality Index (Gallagher)	0.689	0.086	0.665	0.114	−0.024	0.364
Disproportionality Index (Cox & Shugart)	0.334	0.142	0.345	0.159	0.011	0.762

Note: This table presents average electoral participation and competition, ethnic composition, electoral system outcomes, and party system characteristics for municipalities assigned to the negative condition (where respondents received information that malfeasance was increasing) and the positive condition (where respondents received information that malfeasance was declining). Law 21.238, enacted in Chile in July 2020, established term limits for municipal mayors, barring approximately 96 incumbents from seeking reelection in 2021; thus, unconditional reelection rates should be interpreted with caution. We estimate differences in means using a two-sample t-test and report the corresponding p-values.

A.7 Political Outcomes Analysis

Table A.15: Regression Results of Malfeasance Beliefs - Copartisan and Non-Copartisans

	Qualitative		Resources		Distribution	
	Copartisan	Non-copartisan	Copartisan	Non-copartisan	Copartisan	Non-copartisan
Panel A: Information						
Intercept	2.061+	3.144***	2.877	2.125***	1.058**	1.166***
	(1.115)	(0.393)	(1.885)	(0.533)	(0.401)	(0.153)
Treat	0.742***	0.733***	0.788***	0.723***	0.303***	0.357***
	(0.152)	(0.083)	(0.135)	(0.107)	(0.056)	(0.034)
Prior	0.424***	0.415***	0.414***	0.435***	0.382***	0.381***
	(0.037)	(0.023)	(0.038)	(0.019)	(0.050)	(0.026)
Covariates	✓	✓	✓	✓	✓	✓
Num.Obs.	929	2510	929	2510	929	2510
R2	0.195	0.233	0.219	0.215	0.158	0.193
R2 Adj.	0.172	0.224	0.197	0.206	0.134	0.185
Panel B: Benchmarking						
Intercept	1.483	3.121***	2.471	2.236***	0.926*	1.223***
	(1.151)	(0.417)	(1.864)	(0.517)	(0.388)	(0.178)
Treat	0.908	0.362**	0.471*	0.213	0.143	0.080
	(0.498)	(0.139)	(0.186)	(0.251)	(0.100)	(0.101)
Negative	0.710*	0.036	0.494*	-0.106	0.169	-0.055
	(0.292)	(0.234)	(0.207)	(0.246)	(0.106)	(0.113)
Treat:Negative	-0.217	0.430*	0.371	0.590*	0.189	0.320**
	(0.528)	(0.170)	(0.247)	(0.275)	(0.120)	(0.107)
Prior	0.424***	0.416***	0.414***	0.435***	0.377***	0.382***
	(0.036)	(0.022)	(0.038)	(0.020)	(0.048)	(0.028)
Covariates	✓	✓	✓	✓	✓	✓
Num.Obs.	929	2510	929	2510	929	2510
R2	0.204	0.236	0.231	0.217	0.175	0.201
R2 Adj.	0.179	0.227	0.207	0.208	0.149	0.192
Panel C: Magnitudes						
Intercept	1.932+	3.142***	2.651	2.126***	1.011**	1.172***
	(1.084)	(0.395)	(1.789)	(0.527)	(0.363)	(0.155)
Treat	0.783***	0.713***	0.884***	0.698***	0.314***	0.330***
	(0.180)	(0.092)	(0.160)	(0.115)	(0.068)	(0.041)
Stand - Malfeasance	1.844	0.144	2.536	0.121	0.819+	0.081
	(1.991)	(0.394)	(1.903)	(0.294)	(0.486)	(0.274)
Treat:Stand - Malfeasance	-0.907	0.303	-1.846	0.375	-0.277	0.394
	(2.135)	(0.673)	(2.450)	(0.465)	(0.664)	(0.320)
Prior	0.417***	0.415***	0.412***	0.434***	0.368***	0.378***
	(0.038)	(0.022)	(0.038)	(0.019)	(0.049)	(0.027)
Covariates	✓	✓	✓	✓	✓	✓
Num.Obs.	929	2510	929	2510	929	2510
R2	0.199	0.233	0.225	0.215	0.166	0.197
R2 Adj.	0.174	0.224	0.200	0.206	0.140	0.188
Panel D: Information Gap						
Intercept			2.713	2.015***	0.766*	1.071***
			(1.846)	(0.530)	(0.351)	(0.162)
Treat			0.739***	0.741***	0.332***	0.333***
			(0.146)	(0.120)	(0.076)	(0.034)
Corrup diff			0.122*	0.034	0.129*	-0.006
			(0.049)	(0.033)	(0.053)	(0.035)
Treat:Corrup diff			-0.038	0.012	-0.045	0.085**
			(0.053)	(0.033)	(0.058)	(0.026)
Prior			0.507***	0.476***	0.468***	0.428***
			(0.041)	(0.027)	(0.064)	(0.035)
Covariates			✓	✓	✓	✓
Num.Obs.			929	2510	929	2510
R2			0.226	0.216	0.167	0.199
R2 Adj.			0.202	0.207	0.141	0.190

Note: This table replicates the models from Tables 2 and A.21, disaggregated by partisanship. Copartisans intend to vote for the incumbent; non-copartisans are undecided or intend to vote against. In the Information Gap model, *Corrup Diff* is the gap between respondents' priors and true malfeasance: $(\text{Corrupt}_{ti}^{\text{True}} - \text{Corrupt}_{ti})$. Standard errors are clustered at the commune level. * $p < .05$, ** $p < .01$, *** $p < .001$.

Table A.16: Triple Interaction - Negative

	Qualitative	Resources	Distribution
Intercept	2.399*** (0.506)	3.094*** (0.753)	1.090*** (0.197)
Prior	0.422*** (0.019)	0.435*** (0.018)	0.386*** (0.025)
Treat	0.365* (0.142)	0.195 (0.258)	0.077 (0.100)
Negative	0.070 (0.236)	-0.098 (0.256)	-0.047 (0.118)
Copartisan	-1.248*** (0.372)	-0.897** (0.323)	-0.394* (0.160)
Treat:Negative	0.426* (0.170)	0.611* (0.282)	0.322** (0.106)
Treat:Copartisan	0.501 (0.532)	0.249 (0.353)	0.054 (0.153)
Treat:Copartisan:Negative	-0.613 (0.558)	-0.243 (0.405)	-0.116 (0.170)
Num.Obs.	3414	3414	3414
R2	0.286	0.232	0.230
R2 Adj.	0.279	0.224	0.222

Note: This table presents models including a three-way interaction term between the treatment condition, the negative condition, and copartisan status. The negative condition indicates that respondents were informed that malfeasance was increasing in their municipality; the reference category is respondents informed it was declining. Copartisans are respondents who reported an intention to vote for the incumbent in the upcoming municipal election; non-copartisans include those who were undecided or intended to vote for a different candidate. Standard errors are clustered at the commune level. * $p < .05$, ** $p < .01$, *** $p < .001$.

A.8 Further Discussion of Political and Donations Outcomes

Our understanding of voter responses to corruption information remains incomplete unless we consider the broader political context in which such information is received. In this section, we elaborate on our main findings by examining three dimensions of this context: the influence of corruption information on evaluations of incumbent performance across key public goods domains, the reasons it does not lead to reduced voluntary contributions to municipal public goods, and the

ways in which partisan attachments and electoral conditions affect voters' willingness to withdraw support from a corrupt incumbent.

A.8.1 Incumbent Performance Across Public Goods Domains

We investigate this using pre- and post-treatment measures of respondents' evaluations of the mayoral incumbent's performance across four policy areas where municipal governments exercise substantial authority: crime control, healthcare services, sanitation services, and corruption control. As shown in Figure A.14, after respondents are informed about malfeasance, their assessments of the incumbent decline significantly across all four areas, with the largest decline concentrated in perceptions of corruption control. These results suggest that although treated respondents broadly revise their evaluations of the incumbent's competence, this does not translate into reduced willingness to contribute to a shared municipal public good, which is consistent with our null finding on public goods donations.

A.8.2 Contributions to Municipal Public Goods

Contributions to municipal public goods decline by approximately 14.8 percentage points from baseline to post-treatment in both the treatment and placebo groups, consistent with well-documented declines in voluntary contributions over repeated rounds of public goods games (Chaudhuri, 2011). Notably, there is no significant difference in this decline between treated and placebo respondents. This null between-arm effect contrasts with prior research indicating that corruption information reduces public goods contributions (Beekman et al., 2014; Cagala et al., 2018; Campos-Vazquez & Mejia, 2016; Solaz et al., 2019), but aligns with evidence that such information does not affect charitable donations or redistributive preferences (Beesley & Hawkins, 2022; Peyton, 2020). This null result is likely attributable to the architecture of the contribution environment: donations were anonymous, unmonitored, and free of any sanctioning. This design isolates intrinsic prosocial preferences from the strategic and reputational motives through which information about corruption most plausibly operates, which are typically present in designs that incorporate group

interaction, monitoring, or reputational costs (Beekman et al., 2014; Beesley & Hawkins, 2022; Campos-Vazquez & Mejia, 2016; Solaz et al., 2019).

We probe two potential explanations for this null effect. First, Figure A.11 shows no differences in pre- and post-treatment contributions between treated and placebo groups among either copartisans or non-copartisans, ruling out differential partisan response as an explanation. Second, donation rates move in opposite directions across benchmarking conditions: in the positive condition, the treatment group donates at a higher rate than the placebo group post-treatment, whereas in the negative condition, the treatment group donates at a lower rate than the placebo group. These offsetting effects likely cancel out in the pooled estimate, consistent with Solaz et al., 2019, who similarly find that malfeasance information reduces contributions more among out-group than in-group members; however, in our study, the estimates do not reach statistical significance.

A.8.3 Incumbent Vote Support and Partisan Heterogeneity

We contextualize our findings on support for the incumbent by contrasting our results to those synthesized by Incerti (2020). We observe a 4.5 percentage-point decline in incumbent support following malfeasance information, a relatively large effect compared to similar field experiments, and this decline is more pronounced among partisans (7 percentage points) than non-partisans (2 to 3 percentage points).

We begin by examining whether the extensive margin of belief updating accounts for the gap. As shown in Table A.15, both groups update in the same direction and with broadly comparable magnitude, but the share of copartisans updating negatively consistently exceeds that of non-copartisans across outcomes (35%–52% versus 36%–46%). This greater share of copartisans revising beliefs upward mechanically produces a larger predicted decline in incumbent support, indicating the gap is driven primarily by the extensive rather than the intensive margin of updating.

We next consider whether over-representation of partisans in the negative treatment condition drives the differential effect. In our sample, 83% of partisan and 86% of non-copartisan respon-

dents received the negative condition.⁵ This near-identical distribution across groups indicates that treatment-arm composition does not account for the differential effect.

A related possibility is that differential issue salience explains the gap. Drawing on our baseline survey, 33% of partisan respondents identified corruption as a key policy dimension for quality of life in their commune, compared to 41% of non-copartisans. Table A.18 shows that although treatment effects are larger among partisans who prioritize corruption, this difference does not reach statistical significance relative to non-copartisans who also prioritize corruption, indicating issue salience does not explain the gap.

We further examine whether the intensity of partisan attachment accounts for the gap, distinguishing strong, moderate, and weak partisans from politically disenchanted copartisans (incumbent voters reporting no party sympathy). Figure A.12 shows the most pronounced belief revisions and the largest drops in incumbent support concentrated among politically disenchanted and moderate copartisans, a gradient that does not map onto a simple attachment-intensity account of the copartisan gap.

Finally, motivated by evidence that liberals may place greater emphasis on harm and fairness in evaluating moral violations (Hull et al., 2024; Walter & Redlawsk, 2019), we examine whether ideology accounts for the gap. Table A.17 shows that both left-leaning partisans and independents exhibit a statistically significant decline in support for the incumbent, of 4.5 and 8.9 percentage points, respectively, while right-leaning partisans show a smaller, statistically insignificant decline of 2.5 percentage points. However, an interaction model directly comparing the left-leaning and right-leaning groups' treatment effects finds no statistically significant difference between them.

A.8.4 Electoral Context and Competition


An important contextual factor is whether voters reside in constituencies where electoral competition is sufficiently rich to enable defection toward other electoral alternatives: the probability of voting against incumbents declines when ideologically viable alternatives are present (Agerberg,

⁵Percentages are calculated only including respondents who took part in the post-treatment survey.

2020; Elia & Schwindt-Bayer, 2022), while weak or absent challengers limit the electoral pressure incumbents face. Following Boas et al., 2019, we assess competitive conditions using the number of candidates and margins of victory in the 2016 and 2021 elections. Figure A.13 shows that, among the 116 municipalities in our sample, the median number of candidates per constituency is five and the median incumbent margin of victory is approximately 13%. Although coalition parties agreed to field a single candidate in each municipality, independents contested widely: on average, at least one independent candidate with prior party affiliation ran in each constituency, giving voters a realistic alternative to the coalition incumbent in essentially every municipality. This indicates municipal elections in our sample exhibit moderate rather than intense electoral competition.⁶


⁶Chile presents a puzzle for accounts of increasingly competence-based, valence voting and low local turnover: mayoral incumbents enjoy a high reelection rate of 68%. This sits in some tension with Klačnja and Titiunik (2017), who, extending their regression discontinuity design to the Chilean case, found no statistically discernible incumbency advantage in mayoral elections, though they note this null result may reflect a small sample rather than a true absence of advantage. Argote (2020) offers a more fine-grained, Chile-specific account: incumbency advantage was positive and substantial before 2012, disappeared in the first election under the newly introduced voluntary voting regime, and reemerged at an even higher level by 2016.

Figure A.11: Donations by Benchmark and Partisanship




Note: This figure reports pre- and post-treatment raw means of donation rates for the full sample (top), by copartisan status (bottom left), and by treatment benchmark (bottom right). The positive benchmark indicates that corruption is declining, while the negative benchmark indicates that corruption is increasing. Vertical lines represent 95% confidence intervals.

Figure A.12: Belief Updating and Incumbent Support by Partisan Attachment Strength




Note: This figure presents the relationship between partisan attachment, beliefs about malfeasance, and support for the incumbent. Panel A displays mean beliefs with corresponding 95% confidence intervals, disaggregated by partisan attachment. Panel B shows the proportion of respondents indicating an intention to vote for the incumbent, also disaggregated by partisan attachment. Respondents who report no attachment to any party are classified as *No partisan attachment*.

Figure A.13: Number of Candidates and Margins of Victory



Note: This figure reports the distribution of margin of victory (left column) and number of candidates (right column) in the 2016 and 2021 municipal elections for the municipalities included in the study. Dashed vertical lines indicate the median of each distribution.

Figure A.14: Evaluations of incumbent's performance on multiple dimensions



Note: This figure displays the mean change in respondents' evaluations of incumbent performance, from baseline to post-treatment, among treated respondents only, for health, crime, corruption, and cleaning. Performance is measured on a 5-point scale (1 = Very bad to 5 = Very good); values below zero indicate a decline in evaluations after treatment. Thick lines show 90% confidence intervals; thin lines show 95% confidence intervals. Standard errors are clustered at the commune level.

Table A.17: Effect of Treatment on Incumbent Voting by Coalition

	Left-leaning	Right-leaning	Independents	Interaction
Treat	-0.045* (0.019)	-0.025 (0.027)	-0.089*** (0.017)	-0.028 (0.027)
Voting incumbent baseline	0.557*** (0.040)	0.488*** (0.028)	0.535*** (0.059)	0.534*** (0.026)
Intercept	0.247* (0.100)	-0.160** (0.055)	0.546* (0.244)	0.035 (0.067)
Liberal				0.028 (0.028)
Treat:Liberal				-0.016 (0.033)
Covariates	✓	✓	✓	✓
Num.Obs.	1413	1542	449	2955
R2	0.336	0.327	0.340	0.319
R2 Adj.	0.324	0.315	0.304	0.312

Note: The regression table presents the effects of the treatment on support for the incumbent, disaggregated by partisanship groups. These groups are defined according to respondents' party preferences as reported in the baseline survey. *Independents* are those who indicated no party attachments. Standard errors are clustered at the municipality level. * $p < .05$, ** $p < .01$, *** $p < .001$.

Table A.18: Treatment Effects Among Voters Who Prioritize Corruption

	All	Other issues	Corruption issues
Treat	-0.012 (0.017)	-0.040* (0.018)	-0.053** (0.019)
Support incumbent baseline	0.598*** (0.032)	0.542*** (0.024)	0.541*** (0.030)
Corruption priority	0.011 (0.019)		
Treat:Support incumbent baseline	-0.090* (0.041)		
Treat:Corruption priority	-0.046+ (0.025)		
Support incumbent baseline:Corruption priority	-0.070 (0.054)		
Treat:Support incumbent baseline:Corruption priority	0.111 (0.071)		
Intercept	0.119*** (0.014)	0.137*** (0.015)	0.127*** (0.015)
Num.Obs.	3526	2167	1359
R2	0.315	0.313	0.307
R2 Adj.	0.314	0.312	0.306

Note: The regression table reports the treatment effects on support for the incumbent, disaggregated by issue-priority subgroups. These subgroups are defined by respondents' rankings of the most important issues facing their community. The list of options comprises corruption, education, public transport, the quality of local health centers, and access to arts and cultural activities. The *Corruption priority* variable is dichotomized: 1 for respondents who selected corruption as their main priority, and 0 otherwise. The *Support incumbent baseline* variable indicates whether the respondent intended to vote for the incumbent in the upcoming election prior to the treatment. Standard errors are clustered at the municipality level. * $p < .05$, ** $p < .01$, *** $p < .001$.

A.9 Compliance

We assessed whether results hold under treatment compliance. Table A.19 shows that 1,831 subjects received the placebo video and 3,694 the treatment video via WhatsApp. While we cannot confirm whether recipients viewed the video due to WhatsApp's privacy rules, compliance is proxied by correct responses to a ChatBot attention check immediately afterward. Response rates were 59% in the placebo group and 56% in the treatment group.

The ChatBot survey included two attention checks: (1) “*What was the topic of the video?*”. Around 98% respondents answered correctly in the placebo group and 95% in the treatment; and (2) “*What was the information reported in the video?*” answered correctly by 97% and 99%, re-

spectively. Most respondents passed both checks. As noted in the main text, we define compliance based on correct answers to the second question: 59% in the placebo group and 56% in the treatment group.

A post-treatment survey was conducted 3–6 days after subjects received the videos. Among the 3,592 respondents, 2,249 were from the treatment group (61% response rate) and 1,343 from the placebo group (73%). Respondents were asked several attention checks. When asked whether they remembered receiving a video from Chile Transparente, 63% of placebo and 74% of treatment respondents said yes. Among those, 80% in the placebo group and 84% in the treatment group correctly recalled the video’s content.

We defined “Post-Treatment Compliers” as the approximately 38% of subjects who answered the post-treatment survey and reported watching the treatment videos. We defined “Chatbot Compliers” as subjects who answered attention check questions in both the Chatbot and Post-treatment surveys. We used two definitions of compliance (Chatbot Compliers and Post-Treatment Compliers) to estimate the Complier Average Causal Effects (CACE).

Table A.20 reports pre-treatment differences between compliers and non-compliers. Compliers scored lower on the Resource Outcome measure and were more likely to donate compared to non-compliers. They are also more likely to be male and younger. In terms of political and socioeconomic characteristics, compliers tend to be more left-leaning and slightly wealthier than non-compliers. However, compliers and non-compliers are similar in their prior beliefs for the Qualitative and Distribution outcome and also similar in years of education and partisanship.

Table A.19: Treatment Compliance

Treatment status	Placebo	%	Treatment	%	Total	% T. sample
Videos sent	1832		3696	-	5528	100
A: Chatbot survey						
Clicked video	No information					
Respondent Chatbot survey	1077	100	2054	100	3131	57
First compliance question - correct	1055	98	1961	95	3016	
First compliance question - incorrect	22	2	112	5	134	
Second compliance question - correct	1042	97	2039	99	3130	
Second compliance question - incorrect	34	3	15	1	49	
B: Post-treatment survey						
Total post-treatment sample	1832	100	3696	100	5528	
Responded	1343	73	2249	61	3592	65
No response	489	27	1447	39	1936	
C: Post-treatment - compliance						
Total post-treatment sample	1343	100	2249	100	3592	
Reported receiving the video	846	63	1670	74	2516	46
Attention check question - correct	680	51	1401	62	2081	
Attention check question - incorrect	166	12	269	12	435	
D: Chatbot and Post-Treatment survey						
Total post-treatment sample	1343	100	2249	100	3592	
Responded both surveys	987	73	2055	91	3042	55
Attention check Chatbot correct	967	72	1962	87	2929	
Attention check Chatbot incorrect	20	1	93	4	111	
F: Follow up survey - compliance						
Total follow-up sample	1057	100	1198	100	2255	
Reported receiving the video	534	50	729	60	1263	23
Attention check question - correct	371	35	548	46	919	
Attention check question - incorrect	163	15	181	15	344	
G: Chatbot and Follow up survey						
Total follow-up sample	1057	100	1198	100	2255	
Reported receiving the video	534	44	729	60	1263	22
Attention check Chatbot correct	411	39	559	45	970	
Attention check Chatbot incorrect	3	0.3	25	0.02	28	

Note: This table reports compliance and attention check results across three survey waves. Panel A displays the correct responses to attention checks in the Chatbot survey, which were asked immediately after the video. Panel B reports the completion of the post-treatment survey one week later. Panel C shows respondents who completed the post-treatment survey, watched the video, and answered the attention check. Panel D includes those who completed both surveys and passed both attention checks. Panel F summarizes follow-up survey compliance, and Panel G presents respondents who passed both Chatbot and post-treatment checks.

Table A.19 reports “Post-Treatment Compliers” in Panel C, showing counts by treatment status, the percentage who reported receiving the videos, and those who answered the attention check correctly as a share of the post-treatment sample. Panel D presents analogous data for “Chatbot

Compliers” in both groups. Panels F and G present the same metrics for the follow-up survey. Panel F displays the proportion of post-treatment compliers, excluding follow-up respondents, while Panel G summarizes chatbot compliance among follow-up survey participants by treatment condition.

Table A.20: Pre-treatment Measures - Compliers Versus Non-Compliers

	Baseline	Post-Treat			Follow-up		
	Whole sample	Compliers	Non-compliers	Difference	Compliers	Non-compliers	Difference
Qualitative Measures							
Subjective malfeasance scale 1-10	4.49 (2.44)	4.44 (2.42)	4.46 (2.38)	-0.02	4.45 (2.39)	4.38 (2.34)	0.08
Certainty Subjective malfeasance scale 1-10	5.40 (2.71)	5.43 (2.67)	5.31 (2.75)	0.13	5.25 (2.66)	5.49 (2.68)	-0.24
Quantitative Measures - Incentivized							
Resources	3.84 (2.38)	3.75 (2.33)	3.97 (2.48)	-0.22	3.75 (2.35)	3.72 (2.18)	0.03
Distribution	1.90 (0.78)	1.88 (0.79)	1.91 (0.78)	-0.03	1.88 (0.76)	1.89 (0.73)	-0.01
Behavioral measures							
500 CLP Municipal Donation	32% (0.63%)	34% (0.93%)	28% (1.38%)	6%	33% (1.33%)	35% (1.43%)	-3%
Covariates							
Female	64.73% (0.01%)	63.77% (0.96%)	65.92% (1.45%)	-2.15%	67.34% (1.33%)	66.41% (1.51%)	0.93%
Education (Yrs)	15.07 (0.04)	15.21 (0.06)	14.86 (0.09)	0.35	15.36 (0.07)	15.27 (0.09)	0.08
Age	35.88 (0.18)	35.05 (0.26)	36.73 (0.43)	-1.68	32.80 (0.34)	33.00 (0.45)	-0.20
Partisanship	4.47 (0.03)	4.34 (0.04)	4.61 (0.06)	-0.27	4.21 (0.06)	4.39 (0.06)	-0.18
Copartisanship	27.13 (0.60%)	25.32 (0.88%)	29.32 (1.38%)	-3.99	25.38 (1.25%)	28.54 (1.41%)	-3.15
Income (USD)	658 (16)	646 (25)	606 (23)	40	521 (19)	652 (32)	-131
N	5528	2081	1511		987	1268	

Note: This table reports pre-treatment values for all outcomes and relevant covariates broken down by compliance status. Along with mean values, we report standard deviations in parentheses. We also calculate the difference in means for each variable. Compliance status was determined by whether respondents reported receiving the video treatment in the post-treatment and follow surveys.

A.10 Heterogeneity

Education Our *Information Treatment* model of the intent to treat subjects with the malfeasance information video results in a very robust and significant estimated treatment effect. This model implies a minimal engagement with the quantitative content of the message – beliefs update when subjects are simply reminded that there is malfeasance in local government. We hypothesize that individuals will update their beliefs based on the quantitative content of the malfeasance informa-

tion (i.e., the amount of malfeasance reported), and that this updating is conditional on cognitive ability, since processing quantitative treatment content is more cognitively demanding. We use education as a proxy for cognitive ability. We have two expectations: We expect the treatment effect for the information model, with a dichotomous treatment variable, to be similar across education segments in the subject population. On the other hand, we conjecture that the higher educated are more responsive to the quantitative details of the malfeasance information treatment.

We exploit machine learning techniques in order to evaluate 1) whether education is a source of heterogeneous treatment effects and 2) how education-related heterogeneity varies by treatment. Our non-parametric modeling strategy uses Bayesian Additive Regression Trees (BART).

We recover ATE estimates by first using the BART models to generate predicted outcomes for the observed data and then for a set of counterfactual observations. To train the BART model, we provide the observed outcome of interest and a training data matrix consisting of the information treatment variable and the covariates of interest. We then simulate a second set of outcomes based on a separate test data matrix that contains “synthetic” observations that are identical to the training data except that the treatment assignment is inverted (“treated” becomes “placebo” and vice versa). The test data are only used post-estimation to predict counterfactual outcomes based on the results of the BART model, which were trained on the observed training data. Estimating outcomes for *both* the observed and synthetic observations ensures that for each observation there is a corresponding counterfactual “placebo” case, where the only difference between the two is the value for treatment assignment. The individual-level effect of being treated is simply the difference between these two predictions.⁷ Conveniently, with individual-level predictions, we can assess the extent to which education, taking into consideration other covariates, is a principal source of heterogeneity in the malfeasance video treatment effect.⁸

Figure A.15 plots the distribution of the estimated individual-level treatment effects by mag-


⁷Our BART model of heterogeneous effects is a simple specification generated using the BART R package with inputs described above. All other options within BART are left at their default value.

⁸Although there are clear data constraints here – we do not claim to have measured and estimated the effect of all possible covariates. Instead, our claim pertains to the various “pre-treatment” variables collected as part of the surveys.

nitude, along with the histogram of the corresponding education covariate values beneath. For the Qualitative outcome, the top part of Panel (A) shows that the Information treatment effect is heterogeneous in the sample population: the overall estimated ATE (on a measure ranging from 1 to 10) is approximately 0.742 (identical to the estimated coefficient in Table 2). The bottom part of Panel (A) shows the density of the three education covariate values over this range of estimated CATEs; the distribution suggests that education is not correlated with the magnitude of the treatment effect, and a one-way ANOVA finds no statistically significant difference in mean individual treatment effects by education level ($F(2, 3436) = 0.86, p = 0.424$). Panel (C) reports the analogous Information treatment effect for the Distribution outcome: the average individual treatment effect is 0.34, and a one-way ANOVA again finds no statistically significant difference in mean effects by education level ($F(2, 3436) = 0.91, p = 0.402$). Taken together, these results broadly confirm our expectation that simple belief updating in response to information about malfeasance is not meaningfully conditioned on educational level.

We next examine whether responsiveness to the quantitative magnitude of reported malfeasance, rather than to the mere presence of an information treatment, is conditioned on education, consistent with our conjecture that engaging with this more demanding quantitative content requires greater cognitive effort. Panel (B) reports the malfeasance-magnitude effect for the Qualitative outcome: a one-way ANOVA finds a statistically significant difference in mean effects across education groups ($F(2, 3436) = 14.57, p < 0.001$), a difference that persists when comparing only High- and Medium-education respondents ($F(1, 3400) = 28.67, p < 0.001$). A Wilcoxon rank-sum test confirms that High-education respondents' CATEs (median = 0.273) are skewed toward the upper end of the distribution relative to Medium-education respondents (median = 0.263; $p < 0.001$). Panel (D) reports the analogous result for the Distribution outcome, showing the same pattern: a significant difference in mean effects across education groups ($F(2, 3436) = 78.51, p < 0.001$) and a significant Wilcoxon rank-sum test confirming that High-education respondents' CATEs (median = 0.112) are skewed above those of Medium-education respondents (median = 0.109; $p < 0.001$).

Figure A.15: Heterogeneous Conditional Average Treatment Effects: Qualitative and Distribution Corruption Beliefs



Note: This figure displays individual-level treatment effect heterogeneity from a Bayesian Additive Regression Trees (BART) model. Malfeasance is split at the median into high and low categories. Holding treatment status fixed, the model generates counterfactual predictions of post-treatment corruption beliefs under high versus low malfeasance for each respondent; the plotted effect is the difference between these two predictions, sorted from most negative to most positive. The dashed red line marks zero; the solid blue line marks the average effect.

In summary, when the treatment is modeled as a simple binary variable, the effect remains highly significant and consistent across education categories. This result aligns with expectations, as the binary treatment effect does not require substantial cognitive engagement with the video’s quantitative content. In contrast, our hypothesis was that belief updating would correlate with the magnitude of malfeasance, and that this correlation would depend on education level, since processing quantitative details demands greater cognitive engagement. The analysis indicates that, for the Qualitative and Distribution outcomes, the malfeasance-magnitude treatment effect has a larger effect for respondents with higher education. However, these differences are small in absolute terms across all outcomes.

A.11 Corruption Audits in Chile

Information about the audit process and malfeasance found in audits The source of the information treatment in this experiment is the CGO annual audit program. The CGO typically focuses

on government activities such as procurement, hiring processes, or municipal finances. Audited municipalities are selected based on a scoring system incorporating factors such as budget size, transfers to the private sector, and results from previous audits. This scoring system is neither publicly available nor accessible to municipalities. The factors and weights of each element change from one audit program to another.

The CGO audits reveal an average of 31 irregularities per municipality, amounting to USD 1,119 to 6,695,744, or roughly 4% of the average municipal budget in Chile. Irregularities are classified by severity, with about 58% categorized accordingly. Highly complex irregularities involve significant restitution requests or potential civil or criminal liabilities and affect critical processes, often triggering disciplinary actions by audit or government agencies. Complex irregularities indicate severe control weaknesses that may lead to investigations by other institutions. Moderately complex irregularities indicate procedural failures that do not impact critical processes and may not necessitate further action. Mild irregularities are minor and have no impact on critical processes or follow-up. We argue that highly complex irregularities correspond to outright corruption, while the others represent mismanagement, given the absence of criminal prosecution or restitution of funds.

The CGO (and other independent governmental institutions) can prosecute civil servants and mayors for the irregularities found in these audits. Civil servants or mayors found guilty of violating statutory rules or engaging in criminal activities are subject to disciplinary action, which may include reparation orders, unpaid leave, and more severe penalties, such as dismissals, demotions, fines, and non-custodial penalties. On average, 5 out of 31 of the irregularities found in each audit involve some form of disciplinary action.

As noted earlier, our study takes place in a context where corruption is highly salient. Overall perceptions of corruption are high, with 88% of respondents in the *Consejo para la Transparencia*'s annual representative survey believing that more than half of Chilean politicians are involved in corrupt practices (Consejo Nacional para la Transparencia, 2020). However, when asked which government institutions exhibit high corruption levels, municipal governments rank among the

lowest; fewer than 4% of respondents perceive corruption as widespread at the local level (Consejo Nacional para la Transparencia, 2019). In terms of direct experience, less than 3% of respondents reported being solicited for bribes by public officials (Consejo Nacional para la Transparencia, 2018, 2019, 2020).

A.12 Additional Analyses, Subgroup Analyses and Persistence Ratios

Table A.21: Regression Results of Malfeasance Beliefs - Information Gap

	Resources	Distribution
Intercept	2.171*** (0.654)	0.990*** (0.156)
Prior	0.502*** (0.025)	0.481*** (0.034)
Treat	0.735*** (0.096)	0.327*** (0.030)
Corrup diff	0.059* (0.030)	0.030 (0.034)
Treat:Corrup diff	0.000 (0.028)	0.051* (0.024)
Num.Obs.	3439	3439
R2	0.223	0.207
R2 Adj.	0.216	0.200


Note: This table reports regression estimates for the *Information Gap* model. Standard errors are clustered at the commune level. The Corrup diff variable captures the difference between respondents beliefs pre-treatment versus the ground true $*p < .05$, $**p < .01$, $***p < .001$.

Table A.22: Regression Results of Malfeasance Beliefs - Priors Group

	Qualitative			Resources			Distribution		
	Pooled	High	Low	Pooled	High	Low	Pooled	High	Low
Panel A: Information									
Intercept	2.734*** (0.401)	2.330*** (0.480)	2.174*** (0.611)	2.302*** (0.662)	2.513** (0.825)	2.061* (0.865)	1.122*** (0.153)	0.267 (0.567)	1.059*** (0.145)
Prior	0.467*** (0.020)	0.561*** (0.043)	0.516*** (0.030)	0.446*** (0.018)	0.355*** (0.035)	0.668*** (0.054)	0.421*** (0.024)	0.512*** (0.058)	0.499*** (0.034)
Treat	0.749*** (0.079)	0.641*** (0.129)	0.824*** (0.096)	0.740*** (0.088)	0.912*** (0.151)	0.643*** (0.098)	0.347*** (0.028)	0.340*** (0.060)	0.351*** (0.032)
Covariates	✓	✓	✓	✓	✓	✓	✓	✓	✓
Num.Obs.	3439	1336	2103	3439	1384	2055	3439	770	2669
R2	0.263	0.206	0.209	0.220	0.136	0.139	0.203	0.235	0.176
R2 Adj.	0.257	0.191	0.199	0.214	0.120	0.128	0.196	0.208	0.167
Panel B: Benchmarking									
Intercept	2.471*** (0.446)	1.762** (0.575)	2.107** (0.647)	2.198*** (0.642)	2.121** (0.821)	2.120* (0.867)	1.101*** (0.166)	0.161 (0.599)	1.109*** (0.163)
Prior	0.466*** (0.020)	0.562*** (0.043)	0.515*** (0.029)	0.445*** (0.019)	0.358*** (0.035)	0.665*** (0.063)	0.420*** (0.026)	0.502*** (0.057)	0.498*** (0.039)
Treat	0.575** (0.208)	0.665 (0.369)	0.539 (0.293)	0.321 (0.167)	1.014*** (0.287)	-0.116 (0.260)	0.110 (0.067)	0.336 (0.193)	0.024 (0.076)
Negative	0.311 (0.212)	0.618 (0.354)	0.109 (0.236)	0.132 (0.174)	0.443 (0.320)	-0.129 (0.184)	0.031 (0.084)	0.249 (0.159)	-0.067 (0.087)
Treat:Negative	0.202 (0.227)	-0.027 (0.401)	0.334 (0.310)	0.489* (0.190)	-0.114 (0.327)	0.887** (0.279)	0.278*** (0.073)	0.003 (0.205)	0.380*** (0.082)
Covariates	✓	✓	✓	✓	✓	✓	✓	✓	✓
Num.Obs.	3439	1336	2103	3439	1384	2055	3439	770	2669
R2	0.267	0.214	0.212	0.225	0.138	0.148	0.213	0.247	0.186
R2 Adj.	0.261	0.198	0.201	0.218	0.121	0.136	0.206	0.218	0.177
Panel C: Magnitudes									
Intercept	2.717*** (0.402)	3.098*** (0.506)	2.498*** (0.569)	2.274*** (0.648)	2.731*** (0.788)	1.737* (0.868)	1.125*** (0.152)	0.980*** (0.227)	0.969*** (0.194)
Prior	0.464*** (0.020)	0.436*** (0.033)	0.426*** (0.024)	0.444*** (0.018)	0.342*** (0.035)	0.671*** (0.053)	0.415*** (0.024)	0.414*** (0.041)	0.578*** (0.035)
Treat	0.736*** (0.091)	0.812*** (0.128)	0.698*** (0.127)	0.739*** (0.094)	0.899*** (0.165)	0.676*** (0.123)	0.327*** (0.035)	0.335*** (0.053)	0.337*** (0.040)
Stand - Malfeasance	0.490 (0.489)	0.404 (0.664)	0.666 (0.894)	0.590 (0.432)	0.636 (0.617)	0.628 (0.842)	0.217 (0.233)	0.181 (0.291)	0.251 (0.264)
Treat:Stand - Malfeasance	0.179 (0.882)	-0.649 (0.502)	0.636 (1.263)	0.012 (0.597)	-0.103 (0.985)	-0.227 (1.414)	0.302 (0.353)	0.364 (0.530)	0.162 (0.354)
Covariates	✓	✓	✓	✓	✓	✓	✓	✓	✓
Num.Obs.	3439	1366	2073	3439	1366	2073	3439	1595	1844
R2	0.264	0.223	0.226	0.221	0.129	0.138	0.207	0.180	0.222
R2 Adj.	0.258	0.207	0.215	0.214	0.111	0.126	0.200	0.165	0.210
Panel D: Information Gap									
Intercept				2.171*** (0.654)	2.366** (0.835)	1.952* (0.862)	0.990*** (0.156)	0.121 (0.558)	0.990*** (0.159)
Prior				0.502*** (0.025)	0.418*** (0.042)	0.701*** (0.055)	0.481*** (0.034)	0.594*** (0.066)	0.543*** (0.046)
Treat				0.735*** (0.096)	0.992*** (0.248)	0.639*** (0.099)	0.327*** (0.030)	0.349*** (0.069)	0.288*** (0.041)
Corrup diff				0.059* (0.030)	0.054 (0.041)	0.024 (0.036)	0.030 (0.034)	0.085 (0.048)	-0.009 (0.036)
Treat:Corrup diff				0.0003 (0.028)	0.022 (0.052)	0.033 (0.044)	0.051* (0.024)	0.008 (0.055)	0.089* (0.035)
Covariates				✓	✓	✓	✓	✓	✓
Num.Obs.				3439	1384	2055	3439	770	2669
R2				0.223	0.140	0.142	0.207	0.244	0.180
R2 Adj.				0.216	0.122	0.129	0.200	0.216	0.171


Note: This table reports regression estimates by prior beliefs about malfeasance for three outcomes and belief groups (“Pooled”, “High”, “Low”). We use same models reported in Tables 2 and A.21. Standard errors clustered at the commune level. * $p < .05$, ** $p < .01$, *** $p < .001$.

Figure A.16: Corruption Priors and Posteriors Copartisans and Non-Copartisans.




Note: This figure shows prior and posterior beliefs by treatment and partisanship across all three outcome measures. The sample is split into copartisans (respondents intending to vote for the incumbent) and Non-copartisans (respondents not supporting the incumbent). Estimates are broken down by treatment condition: Negative (increasing malfeasance) and Positive (declining malfeasance). No Information corresponds to respondents' prior beliefs measured in the baseline survey.

Figure A.17: Persistence ratios



Note: This figure reports persistence ratios of average updating estimates for all three outcomes. Persistence ratios are calculated by dividing one-month estimates by estimates obtained 3-6 days after the treatment was delivered. A ratio of 1 would be equivalent to no treatment decay a month after the intervention. 3-6 days treatment effects (Post-treatment) are reported in the x-axis, while one-month estimates are reported in the y-axis. Persistence ratios are reported next to each crosshair point. Negative persistence ratios correspond to cases where there is a sign reversal between the post-treatment and follow-up estimates. We broke down persistence ratios by treatment status: “Information Treatment”, “Negative”, and “Positive”, and different levels of standardized measures of malfeasance. The “Negative” treatment condition comprises all respondents who received unfavorable information, whereas “Positive” includes subjects who received a favorable information frame.

Figure A.18: Average Beliefs Pre, Post, and Follow-up




Note: This table reports average beliefs and 95% confidence intervals by benchmarking status (Positive or Negative) across survey waves. Sample sizes vary across waves due to attrition at follow-up. P-values correspond to two-sample independent t-tests comparing means between the different waves.

A.13 Power calculations

We conducted power calculations using simulations to test the probability of rejecting the null hypothesis of no effect for different treatment effects. Based on the results of this simulation, we are well-powered for both the whole sample and the “negative” frame sample. However, we are underpowered for the “positive” frame. Figure A.19 report the results of our power calculations.

Figure A.19: Power Calculations



Note: This figure displays the probability of rejecting the null hypothesis of no treatment effect across a range of effect sizes. Power calculations are based on the sample sizes used in the study.

Table A.23: Audit-Information Experiments and the Present Study

Study	Design	Information treatment & delivery	Prior beliefs	Belief outcomes	Behavioral outcomes
Ferraz & Finan (2008), Brazil	Natural experiment (randomized audit lottery)	Public release of municipal audit findings; diffusion via radio and media; exposure uncontrolled	Not measured	Not directly measured	Electoral punishment: reelection falls from 53% to 20% where three irregularities uncovered
Chong et al. (2015), Mexico	Field experiment (precinct-level)	Flyers summarizing audit results, delivered before municipal elections	Not measured	No significant change in corruption perceptions (non-incentivized)	Reduced turnout (demobilization); no significant change in incumbent support
Arias et al. (2022), Mexico	Field experiment (precinct-level)	Flyers and coordinated group meetings; share of funds irregular; tied to upcoming election	Inferred from control-group posteriors	Updating tracks the prior–signal gap (“priors rule”; Bayesian pattern)	Incumbent-party support responds conditional on priors
Dunning et al. (2019), Metaketa I (6 countries)	Coordinated field experiments, harmonized design	Varied politician-performance information; Flyers, door-to-door, and video	Baseline vote intentions and beliefs in most sites	Limited average movement in beliefs	Pooled: no average effect on incumbent vote choice or turnout
Enríquez et al. (2024), Mexico	Field experiment (saturation design)	Facebook advertising of municipal spending irregularities, incl. high-saturation arms	Not directly measured	Not the focus	Electoral effects on incumbent party under sustained/saturated exposure
This study, Chile	Field survey experiment (individual-level randomization)	Customized 50-second WhatsApp video: audited amount, foregone flu vaccines, temporal benchmark; temporally separated from surveys	Directly measured pre-treatment; three measures, two financially incentivized	Large, persistent upward updating; heuristic rather than calibrated (insensitive to magnitude; negative benchmark dominates); parallel across partisan groups	Incumbent vote intention -4.5 pp; no effect on incentivized donations to municipal public good

Note: The table positions the present study against the audit-information studies most closely related to our design. “Prior beliefs” indicates whether subjects’ pre-treatment corruption beliefs are directly elicited. Our design is, to our knowledge, the only one combining directly measured priors, financially incentivized belief outcomes, and individually randomized private delivery of real audit content.

A.14 Ethics Approvals

This research project received ethics approval from the relevant institutions; details have been withheld to preserve anonymity during the review process. We adhered to the [Country] Data Protection Acts of 1998 and 2018, as well as Chile's Data Protection Act (Law No. 19.628). All participants provided informed consent, and we ensured that their data was handled securely and in compliance with relevant data protection regulations. Participants were compensated approximately \$3.70 per survey, and no deception was involved in the study. The University of — served as the primary data controller

Appendix References

- Agerberg, M. (2020). The lesser evil? corruption voting and the importance of clean alternatives. *Comparative Political Studies*, 53(2), 253–287.
- Alesina, A., Devleeschauwer, A., Easterly, W., Kurlat, S., & Wacziarg, R. (2003). Fractionalization. *Journal of Economic Growth*, 8(2), 155–194.
- Anderson, K., Zamarro, G., Steele, J., & Miller, T. (2021). Comparing performance of methods to deal with differential attrition in randomized experimental evaluations. *Evaluation Review*, 45(1-2), 70–104.
- Argote, P. (2020). Incumbency advantage and shocks in the electorate: The adoption of voluntary voting. *European Political Science Review*, 12(2), 173–197. <https://doi.org/10.1017/S1755773920000041>
- Beekman, G., Bulte, E., & Nillesen, E. (2014). Corruption, investments and contributions to public goods: Experimental evidence from rural Liberia. *Journal of Public Economics*, 115, 37–47.
- Beesley, C., & Hawkins, D. (2022). Corruption, institutional trust and political engagement in peru. *World Development*, 151, 105743.

- Boas, T. C., Hidalgo, F. D., & Melo, M. A. (2019). Norms versus action: Why voters fail to sanction malfeasance in Brazil. *American Journal of Political Science*, 63(2), 385–400. <https://doi.org/10.1111/ajps.12413>
- Cagala, T., Glogowsky, U., Grimm, V., & Rincke, J. (2018). Public goods provision with rent-extracting administrators. *The Economic Journal*, 129(620), 1593–1617.
- Campos-Vazquez, R. M., & Mejia, L. A. (2016). Does corruption affect cooperation? a laboratory experiment. *Latin American Economic Review*, 25(1), 5.
- Chaudhuri, A. (2011). Sustaining cooperation in laboratory public goods experiments: A selective survey of the literature. *Experimental Economics*, 14(1), 47–83.
- Consejo Nacional para la Transparencia. (2018). Estudio nacional de transparencia.
- Consejo Nacional para la Transparencia. (2019). Estudio nacional de transparencia.
- Consejo Nacional para la Transparencia. (2020). Estudio nacional de transparencia.
- Dalton, R. J. (2008). The quantity and the quality of party systems: Party system polarization, its measurement, and its consequences. *Comparative Political Studies*, 41(7), 899–920.
- Elia, E., & Schwindt-Bayer, L. A. (2022). Corruption perceptions, opposition parties, and reelecting incumbents in Latin America. *Electoral Studies*, 80, 102545.
- Greifer, N. (2022). *Cobalt: Covariate balance tables and plots* [R package version 4.4.1]. <https://cran.r-project.org/package=cobalt>
- Hull, K., Warren, C., & Smith, K. (2024). Politics makes bastards of us all: Why moral judgment is politically situational. *Political Psychology*, 45(6), 1013–1029. <https://doi.org/https://doi.org/10.1111/pops.12954>
- Incerti, T. (2020). Corruption information and vote share: A meta-analysis and lessons for experimental design. *American Political Science Review*, 114(3), 761–774. <https://doi.org/10.1017/S000305542000012X>
- Klašnja, M., & Titunik, R. (2017). The incumbency curse: Weak parties, term limits, and unfulfilled accountability. *American Political Science Review*, 111(1), 129–148. <https://doi.org/10.1017/S0003055416000575>

- LAPOP. (2021). Chile: Americasbarometer 2021 [Vanderbilt University, Nashville, TN].
- Peyton, K. (2020). Does trust in government increase support for redistribution? Evidence from randomized survey experiments. *American Political Science Review*, 114(2), 596–602. <https://doi.org/10.1017/S0003055420000076>
- Ridgeway, G., McCaffrey, D., Morral, A., Cefalu, M., Burgette, L., Pane, J., & Griffin, B. A. (2021, October). *Toolkit for weighting and analysis of nonequivalent groups: A guide to the twang package* (R Package). RAND. <https://cran.r-project.org/web/packages/twang/vignettes/twang.pdf>
- Solaz, H., Vries, C. E. D., & de Geus, R. A. (2019). In-group loyalty and the punishment of corruption. *Comparative Political Studies*, 52(6), 896–926.
- Stuart, E. A., Lee, B. K., & Leacy, F. P. (2013). Prognostic score-based balance measures can be a useful diagnostic for propensity score methods in comparative effectiveness research. *Journal of Clinical Epidemiology*, 66(8), S84–S90.e1.
- Walter, A. S., & Redlawsk, D. P. (2019). Voters’ partisan responses to politicians’ immoral behavior. *Political Psychology*, 40(5), 1075–1097. <https://doi.org/https://doi.org/10.1111/pops.12582>